



An Introduction to Hibernate 6

Version 6.2.45.Final

Table of Contents

Preface	1
1. Introduction	2
1.1. Hibernate and JPA	2
1.2. Writing Java code with Hibernate	4
1.3. Hello, Hibernate	4
1.4. Hello, JPA	6
1.5. Organizing persistence logic	8
1.6. Testing persistence logic	9
1.7. Architecture and the persistence layer	11
1.8. Overview	12
2. Configuration and bootstrap	13
2.1. Including Hibernate in your project build	13
2.2. Optional dependencies	14
2.3. Configuration using JPA XML	14
2.4. Configuration using Hibernate API	15
2.5. Configuration using Hibernate properties file	16
2.6. Basic configuration settings	16
2.7. Automatic schema export	17
2.8. Logging the generated SQL	17
2.9. Minimizing repetitive mapping information	18
2.10. Nationalized character data in SQL Server	18
3. Entities	20
3.1. Entity classes	20
3.2. Access types	21
3.3. Entity class inheritance	21
3.4. Identifier attributes	22
3.5. Generated identifiers	22
3.6. Natural keys as identifiers	23
3.7. Composite identifiers	24
3.8. Version attributes	25
3.9. Natural id attributes	25
3.10. Basic attributes	26
3.11. Enumerated types	27
3.12. Converters	28
3.13. Compositional basic types	29
3.14. Embeddable objects	30
3.15. Associations	31
3.16. Many-to-one	32
3.17. One-to-one (first way)	33
3.18. One-to-one (second way)	34
3.19. Many-to-many	35
3.20. Collections of basic values and embeddable objects	36
3.21. Collections mapped to SQL arrays	37
3.22. Collections mapped to a separate table	37
3.23. Summary of annotations	38
3.24. equals() and hashCode()	39
4. Object/relational mapping	41
4.1. Mapping entity inheritance hierarchies	41
4.2. Mapping to tables	43
4.3. Mapping entities to tables	43
4.4. Mapping associations to tables	44
4.5. Mapping to columns	45
4.6. Mapping basic attributes to columns	45
4.7. Mapping associations to foreign key columns	46
4.8. Mapping primary key joins between tables	48
4.9. Column lengths and adaptive column types	49
4.10. LOBs	49
4.11. Mapping embeddable types to UDTs or to JSON	50
4.12. Summary of SQL column type mappings	51
4.13. Mapping to formulas	52

5. Interacting with the database	53
5.1. Persistence Contexts	53
5.2. Creating a session	54
5.3. Managing transactions	55
5.4. Operations on the persistence context	55
5.5. Cascading persistence operations	56
5.6. Proxies and lazy fetching	56
5.7. Entity graphs and eager fetching	58
5.8. Flushing the session	58
5.9. Queries	59
5.10. HQL queries	59
5.11. Criteria queries	60
5.12. Native SQL queries	62
5.13. Limits, pagination, and ordering	63
5.14. Representing projection lists	63
5.15. Named queries	64
5.16. Controlling lookup by id	65
5.17. Interacting directly with JDBC	66
5.18. What to do when things go wrong	66
6. Tuning and performance	68
6.1. Tuning the connection pool	68
6.2. Enabling statement batching	69
6.3. Association fetching	69
6.4. Batch fetching and subselect fetching	70
6.5. Join fetching	71
6.6. The second-level cache	72
6.7. Specifying which data is cached	73
6.8. Caching by natural id	74
6.9. Caching and association fetching	75
6.10. Configuring the second-level cache provider	75
6.11. Caching query result sets	76
6.12. Second-level cache management	77
6.13. Session cache management	78
6.14. Stateless sessions	78
6.15. Optimistic and pessimistic locking	79
6.16. Collecting statistics	80
6.17. Tracking down slow queries	80
6.18. Reactive programming with Hibernate	81
7. Advanced Topics	82
7.1. Filters	82
7.2. Multi-tenancy	84
7.3. Using custom-written SQL	85
7.4. Handling database-generated columns	85
7.5. User-defined generators	86
7.6. Naming strategies	87
7.7. Spatial datatypes	88
7.8. Ordered and sorted collections and map keys	89
7.9. Any mappings	90
7.10. Selective column lists in inserts and updates	91
7.11. Using the bytecode enhancer	92
7.12. Named fetch profiles	93

Preface

Hibernate 6 is a major redesign of the world's most popular and feature-rich ORM solution. The redesign has touched almost every subsystem of Hibernate, including the APIs, mapping annotations, and the query language. This new Hibernate is more powerful, more robust, and more typesafe.

With so many improvements, it's very difficult to summarize the significance of this work. But the following general themes stand out. Hibernate 6:

- finally takes advantage of the advances in relational databases over the past decade, updating the query language to support a raft of new constructs in modern dialects of SQL,
- exhibits much more consistent behavior across different databases, greatly improving portability, and generates much higher-quality DDL from dialect-independent code,
- improves error reporting by more scrupulous validation of queries *before* access to the database,
- improves the type-safety of O/R mapping annotations, clarifies the separation of API, SPI, and internal implementation, and fixes some long-standing architectural flaws,
- removes or deprecates legacy APIs, laying the foundation for future evolution, and
- makes far better use of Javadoc, putting much more information at the fingertips of developers.

Hibernate 6 and Hibernate Reactive are now core components of Quarkus 3, the most exciting new environment for cloud-native development in Java, and Hibernate remains the persistence solution of choice for almost every major Java framework or server.

Unfortunately, the changes in Hibernate 6 have obsoleted much of the information about Hibernate that's available in books, in blog posts, and on stackoverflow.

This guide is an up-to-date, high-level discussion of the current set of Hibernate features. It does not attempt to cover every feature and should be used in conjunction with other documentation -

- Hibernate's extensive [Javadoc](#)
- The [Hibernate Query Language Guide](#)
- The [Hibernate User Guide](#).



The Hibernate User Guide provides detailed discussion about most aspects of Hibernate. But with so much information to cover, readability is difficult to achieve - it should be considered more of a reference guide.

Where applicable, this guide will provide links to detailed User Guide content within its covered topics.

Chapter 1. Introduction

Hibernate is usually described as a library that makes it easy to map Java classes to relational database tables. But this formulation does not justice to the central role played by the relational data itself. So a better description might be:

Hibernate makes **relational data** visible to a program written in Java, in a **natural** and **typesafe** form,

1. making it easy to write complex queries and work with their results,
2. letting the program easily synchronize changes made in memory with the database, respecting the ACID properties of transactions, and
3. allowing performance optimizations to be made after the basic persistence logic has already been written.

Here the relational data is the focus, along with the importance of typesafety. The goal of *Object/relational mapping* (ORM) is to eliminate fragile and untypesafe code, and make large programs easier to maintain in the long run.

ORM takes the pain out of persistence by relieving the developer of the need to hand-write tedious, repetitive, and fragile code for flattening graphs of objects to database tables and rebuilding graphs of objects from flat SQL query result sets. Even better, ORM makes it much easier to tune performance later, after the basic persistence logic has already been written.



A perennial question is: should I use ORM, or plain SQL? The answer is usually: *use both*. JPA and Hibernate were designed to work *in conjunction with* handwritten SQL. You see, most programs with nontrivial data access logic will benefit from the use of ORM at least *somewhere*. But if Hibernate is making things more difficult, for some particularly tricky piece of data access logic, the only sensible thing to do is to use something better suited to the problem! Just because you're using Hibernate for persistence doesn't mean you have to use it for *everything*.

Developers often ask about the relationship between Hibernate and JPA, so let's take a short detour into some history.

1.1. Hibernate and JPA

Hibernate was the inspiration behind the *Java* (now *Jakarta*) *Persistence API*, or JPA, and includes a complete implementation of the latest revision of this specification.

The early history of Hibernate and JPA

The Hibernate project began in 2001, when Gavin King's frustration with Entity Beans in EJB 2 boiled over. It quickly overtook other open source and commercial contenders to become the most popular persistence solution for Java, and the book *Hibernate in Action*, written with Christian Bauer, was an influential bestseller.

In 2004, Gavin and Christian joined a tiny startup called JBoss, and other early Hibernate contributors soon followed: Max Rydahl Andersen, Emmanuel Bernard, Steve Ebersole, and Sanne Grinovero.


Soon after, Gavin joined the EJB 3 expert group and convinced the group to deprecate Entity Beans in favor of a brand-new persistence API modelled after Hibernate. Later, members of the TopLink team got involved, and the Java Persistence API evolved as a collaboration between—primarily—Sun, JBoss, Oracle, and Sybase, under the leadership of Linda Demichiel.

Over the intervening two decades, *many* talented people have contributed to the development of Hibernate. We're all especially grateful to Steve, who has led the project for many years, since Gavin stepped back to focus in other work.

We can think of the API of Hibernate in terms of three basic elements:

- an implementation of the JPA-defined APIs, most importantly, of the interfaces `EntityManagerFactory` and `EntityManager`, and of the JPA-defined O/R mapping annotations,
- a *native API* exposing the full set of available functionality, centered around the interfaces `SessionFactory`, which extends `EntityManagerFactory`, and `Session`, which extends `EntityManager`, and
- a set of *mapping annotations* which augment the O/R mapping annotations defined by JPA, and which may be used with the JPA-defined interfaces, or with the native API.

Hibernate also offers a range of SPIs for frameworks and libraries which extend or integrate with Hibernate, but we're not interested in any of that stuff here.



As an application developer, you must decide whether to:

- write your program in terms of `Session` and `SessionFactory`, or
- maximize portability to other implementations of JPA by, wherever reasonable, writing code in terms of `EntityManager` and `EntityManagerFactory`, falling back to the native APIs only where necessary.

Whichever path you take, you will use the JPA-defined mapping annotations most of the time, and the Hibernate-defined annotations for more advanced mapping problems.



You might wonder if it's possible to develop an application using *only* JPA-defined APIs, and, indeed, that's possible in principle. JPA is a great baseline that really nails the basics of the object/relational mapping problem. But without the native APIs, and extended mapping annotations, you miss out on much of the power of Hibernate.

Since Hibernate existed before JPA, and since JPA was modelled on Hibernate, we unfortunately have some competition and duplication in naming between the standard and native APIs. For example:

Table 1. Examples of competing APIs with similar naming

Hibernate	JPA
<code>org.hibernate.annotations.CascadeType</code>	<code>javax.persistence.CascadeType</code>
<code>org.hibernate.FlushMode</code>	<code>javax.persistence.FlushModeType</code>
<code>org.hibernate.annotations.FetchMode</code>	<code>javax.persistence.FetchType</code>
<code>org.hibernate.query.Query</code>	<code>javax.persistence.Query</code>
<code>org.hibernate.Cache</code>	<code>javax.persistence.Cache</code>
<code>@org.hibernate.annotations.NamedQuery</code>	<code>@javax.persistence.NamedQuery</code>
<code>@org.hibernate.annotations.Cache</code>	<code>@javax.persistence.Cacheable</code>

Typically, the Hibernate-native APIs offer something a little extra that's missing in JPA, so this isn't exactly a *flaw*. But it's something to watch out for.

1.2. Writing Java code with Hibernate

If you're completely new to Hibernate and JPA, you might already be wondering how the persistence-related code is structured.

Well, typically, our persistence-related code comes in two layers:

1. a representation of our data model in Java, which takes the form of a set of annotated entity classes, and
2. a larger number of functions which interact with Hibernate's APIs to perform the persistence operations associated with your various transactions.

The first part, the data or "domain" model, is usually easier to write, but doing a great and very clean job of it will strongly affect your success in the second part.

Most people implement the domain model as a set of what we used to call "Plain Old Java Objects", that is, as simple Java classes with no direct dependencies on technical infrastructure, nor on application logic which deals with request processing, transaction management, communications, or interaction with the database.



Take your time with this code, and try to produce a Java model that's as close as reasonable to the relational data model. Avoid using exotic or advanced mapping features when they're not really needed. When in the slightest doubt, map a foreign key relationship using @ManyToOne with @OneToMany(mappedBy=...) in preference to more complicated association mappings.

The second part of the code is much trickier to get right. This code must:

- manage transactions and sessions,
- interact with the database via the Hibernate session,
- fetch and prepare data needed by the UI, and
- handle failures.



Responsibility for transaction and session management, and for recovery from certain kinds of failure, is best handled in some sort of framework code.

We're going to [come back soon](#) to the thorny question of how this persistence logic should be organized, and how it should fit into the rest of the system.

1.3. Hello, Hibernate

Before we get deeper into the weeds, we'll quickly present a basic example program that will help you get started if you don't already have Hibernate integrated into your project.

We begin with a simple gradle build file:

build.gradle

```
plugins {
    id 'java'
}

group = 'org.example'
version = '1.0-SNAPSHOT'

repositories {
    mavenCentral()
}

dependencies {
    // the GOAT ORM
    implementation 'org.hibernate.orm:hibernate-core:6.2.2.Final'

    // Hibernate Validator
    implementation 'org.hibernate.validator:hibernate-validator:8.0.0.Final'
    implementation 'org.glassfish:jakarta.el:4.0.2'

    // Agroal connection pool
    implementation 'org.hibernate.orm:hibernate-agroal:6.2.2.Final'
    implementation 'io.agroal:agroal-pool:2.1'

    // logging via Log4j
    implementation 'org.apache.logging.log4j:log4j-core:2.20.0'
```

```

// JPA Metamodel Generator
annotationProcessor 'org.hibernate.orm:hibernate-jpamodelgen:6.2.2.Final'

// Compile-time checking for HQL
//implementation 'org.hibernate:query-validator:2.0-SNAPSHOT'
//annotationProcessor 'org.hibernate:query-validator:2.0-SNAPSHOT'

// H2 database
runtimeOnly 'com.h2database:h2:2.1.214'
}

```

Only the first of these dependencies is absolutely *required* to run Hibernate.

Next, we'll add a logging configuration file for log4j:

log4j2.properties

```

rootLogger.level = info
rootLogger.appenderRefs = console
rootLogger.appenderRef.console.ref = console

logger.hibernate.name = org.hibernate.SQL
logger.hibernate.level = info

appender.console.name = console
appender.console.type = Console
appender.console.layout.type = PatternLayout
appender.console.layout.pattern = %highlight{[%p]} %m%n

```

Now we need some Java code. We begin with our *entity class*:

Book.java

```

package org.hibernate.example;

import jakarta.persistence.Entity;
import jakarta.persistence.Id;
import jakarta.validation.constraints.NotNull;

@Entity
class Book {
    @Id
    String isbn;

    @NotNull
    String title;

    Book() {}

    Book(String isbn, String title) {
        this.isbn = isbn;
        this.title = title;
    }
}

```

Finally, let's see code which configures and instantiates Hibernate and asks it to persist and query the entity. Don't worry if this makes no sense at all right now. It's the job of this Introduction to make all this crystal clear.

Main.java

```

package org.hibernate.example;

import org.hibernate.cfg.Configuration;

import static java.lang.Boolean.TRUE;
import static java.lang.System.out;
import static org.hibernate.cfg.AvailableSettings.*;

public class Main {
    public static void main(String[] args) {
        var sessionFactory = new Configuration()
            .addAnnotatedClass(Book.class)
            // use H2 in-memory database

```



```

        .setProperty(URL, "jdbc:h2:mem:db1")
        .setProperty(USER, "sa")
        .setProperty(PASS, "")
        // use Agroal connection pool
        .setProperty("hibernate.agroal.maxSize", "20")
        // display SQL in console
        .setProperty(SHOW_SQL, TRUE.toString())
        .setProperty(FORMAT_SQL, TRUE.toString())
        .setProperty(HIGHLIGHT_SQL, TRUE.toString())
        .buildSessionFactory();

// export the inferred database schema
sessionFactory.getSchemaManager().exportMappedObjects(true);

// persist an entity
sessionFactory.inTransaction(session -> {
    session.persist(new Book("9781932394153", "Hibernate in Action"));
});

// query data using HQL
sessionFactory.inSession(session -> {
    out.println(session.createQuery("select isbn||': '||title from Book").getSingleResult());
});

// query data using criteria API
sessionFactory.inSession(session -> {
    var builder = sessionFactory.getCriteriaBuilder();
    var query = builder.createQuery(String.class);
    var book = query.from(Book.class);
    query.select(builder.concat(builder.concat(book.get(Book_.isbn), builder.literal(": ")),
        book.get(Book_.title)));
    out.println(session.createQuery(query).getSingleResult());
});
}
}

```

Here we've used Hibernate's native APIs. We could have used JPA-standard APIs to achieve the same thing.

1.4. Hello, JPA

If we limit ourselves to the use of JPA-standard APIs, we need to use XML to configure Hibernate.

META-INF/persistence.xml

```

<persistence xmlns="https://jakarta.ee/xml/ns/persistence"
    xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
    xsi:schemaLocation="https://jakarta.ee/xml/ns/persistence
https://jakarta.ee/xml/ns/persistence/persistence_3_0.xsd"
    version="3.0">

    <persistence-unit name="example">

        <class>org.hibernate.example.Book</class>

        <properties>

            <!-- H2 in-memory database -->
            <property name="jakarta.persistence.jdbc.url"
                value="jdbc:h2:mem:db1"/>

            <!-- Credentials -->
            <property name="jakarta.persistence.jdbc.user"
                value="sa"/>
            <property name="jakarta.persistence.jdbc.password"
                value=""/>

            <!-- Agroal connection pool -->
            <property name="hibernate.agroal.maxSize"
                value="20"/>

            <!-- display SQL in console -->
            <property name="hibernate.show_sql" value="true"/>

```

```

        <property name="hibernate.format_sql" value="true"/>
        <property name="hibernate.highlight_sql" value="true"/>

    </properties>

</persistence-unit>
</persistence>

```

Note that our `build.gradle` and `log4j2.properties` files are unchanged.

Our entity class is also unchanged from what we had before.

Unfortunately, JPA doesn't offer an `inSession()` method, so we'll have to implement session and transaction management ourselves. We can put that logic in our own `inSession()` function, so that we don't have to repeat it for every transaction. Again, you don't need to understand any of this code right now.

`Main.java` (JPA version)

```

package org.hibernate.example;

import jakarta.persistence.EntityManager;
import jakarta.persistence.EntityManagerFactory;

import java.util.Map;
import java.util.function.Consumer;

import static jakarta.persistence.Persistence.createEntityManagerFactory;
import static java.lang.System.out;
import static org.hibernate.cfg.AvailableSettings.JAKARTA_HBM2DDL_DATABASE_ACTION;
import static org.hibernate.tool.schema.Action.CREATE;

public class Main {
    public static void main(String[] args) {
        var factory = createEntityManagerFactory("example",
            // export the inferred database schema
            Map.of(JAKARTA_HBM2DDL_DATABASE_ACTION, CREATE));

        // persist an entity
        inSession(factory, entityManager -> {
            entityManager.persist(new Book("9781932394153", "Hibernate in Action"));
        });

        // query data using HQL
        inSession(factory, entityManager -> {
            out.println(entityManager.createQuery("select isbn||': '||title from Book").getSingleResult());
        });

        // query data using criteria API
        inSession(factory, entityManager -> {
            var builder = factory.getCriteriaBuilder();
            var query = builder.createQuery(String.class);
            var book = query.from(Book.class);
            query.select(builder.concat(builder.literal(":"),
                book.get(Book_.title)));
            out.println(entityManager.createQuery(query).getSingleResult());
        });
    }

    // do some work in a session, performing correct transaction management
    static void inSession(EntityManagerFactory factory, Consumer<EntityManager> work) {
        var entityManager = factory.createEntityManager();
        var transaction = entityManager.getTransaction();
        try {
            transaction.begin();
            work.accept(entityManager);
            transaction.commit();
        }
        catch (Exception e) {
            if (transaction.isActive()) transaction.rollback();
            throw e;
        }
        finally {
            entityManager.close();
        }
    }
}

```

```

    }
}

```

In practice, we never access the database directly from a `main()` method. So now let's talk about how to organize persistence logic in a real system. The rest of this chapter is not compulsory. If you're itching for more details about Hibernate itself, you're quite welcome to skip straight to the [next chapter](#), and come back later.

1.5. Organizing persistence logic

In a real program, persistence logic like the code shown above is usually interleaved with other sorts of code, including logic:

- implementing the rules of the business domain, or
- for interacting with the user.

Therefore, many developers quickly—even *too quickly*, in our opinion—reach for ways to isolate the persistence logic into some sort of separate architectural layer. We're going to ask you to suppress this urge for now.



The *easiest* way to use Hibernate is to call the `Session` or `EntityManager` directly. If you're new to Hibernate, frameworks which wrap JPA are only going to make your life more difficult.

We prefer a *bottom-up* approach to organizing our code. We like to start thinking about methods and functions, not about architectural layers and container-managed objects. To illustrate the sort of approach to code organization that we advocate, let's consider a service which queries the database using HQL or SQL.

We might start with something like this, a mix of UI and persistence logic:

```

@Path("/") @Produces("application/json")
public class BookResource {
    @GET @Path("book/{isbn}")
    public Book getBook(String isbn) {
        var book = sessionFactory.fromTransaction(session -> session.find(Book.class, isbn));
        return book == null ? Response.status(404).build() : book;
    }
}

```

Indeed, we might also *finish* with something like that—it's quite hard to identify anything concretely wrong with the code above, and for such a simple case it seems really difficult to justify making this code more complicated by introducing additional objects.

One very nice aspect of this code, which we wish to draw your attention to, is that session and transaction management is handled by generic "framework" code, just as we already recommended above. In this case, we're using the `fromTransaction()` method, which happens to come built in to Hibernate. But you might prefer to use something else, for example:

- in a container environment like Jakarta EE or Quarkus, *container-managed transactions* and *container-managed persistence contexts*, or
- something you write yourself.

The important thing is that calls like `createEntityManager()` and `getTransaction().begin()` don't belong in regular program logic, because it's tricky and tedious to get the error handling correct.

Let's now consider a slightly more complicated case.

```

@Path("/") @Produces("application/json")
public class BookResource {
    private static final RESULTS_PER_PAGE = 20;

    @GET @Path("books/{titlePattern}/{page:\\d+}")
    public List<Book> findBooks(String titlePattern, int page) {
        var books = sessionFactory.fromTransaction(session -> {
            return entityManager.createQuery("from Book where title like ?1 order by title", Book.class)
                .setParameter(1, titlePattern)
                .setMaxResults(RESULTS_PER_PAGE)           // return at most 20 results
                .setFirstResult(page*RESULTS_PER_PAGE)     // start from the given page of results
                .getResultList();
        });
        return books.isEmpty() ? Response.status(404).build() : books;
    }
}

```

This is fine, and we won't complain if you prefer to leave the code exactly as it appears above. But there's one thing we could perhaps improve. We love super-short methods with single responsibilities, and there looks to be an opportunity to introduce one here. Let's hit the code with our favorite thing, the Extract Method refactoring. We obtain:

```
static List<Book> findBooksByTitleWithPagination(EntityManager entityManager,
                                                String titlePattern, int max, int start) {
    return entityManager.createQuery("from Book where title like ?1 order by title", Book.class)
        .setParameter(1, titlePattern)
        .setMaxResults(max)
        .setFirstResult(start)
        .getResultList();
}
```

This is an example of a *query method*, a function which accepts arguments to the parameters of a HQL or SQL query, and executes the query, returning its results to the caller. And that's all it does; it doesn't orchestrate additional program logic, and it doesn't perform transaction or session management.

It's even better to specify the query string using the `@NamedQuery` annotation, so that Hibernate can validate the query it at startup time, that is, when the `SessionFactory` is created, instead of when the query is first executed.

We need a place to put the annotation, so let's move our query method to a new class:

```
@NamedQuery(name="findBooksByTitle",
            query="from Book where title like :title order by title")
class Queries {

    static List<Book> findBooksByTitleWithPagination(EntityManager entityManager,
                                                    String titlePattern, int max, int start) {
        return entityManager.createNamedQuery("findBooksByTitle", Book.class)
            .setParameter("title", titlePattern)
            .setMaxResults(max)
            .setFirstResult(start)
            .getResultList();
    }
}
```

Notice that our query method doesn't attempt to hide the `EntityManager` from its clients. Indeed, the client code is responsible for providing the `EntityManager` or `Session` to the query method. This is a quite distinctive feature of our whole approach.

The client code may:

- obtain an `EntityManager` or `Session` by calling `inTransaction()` or `fromTransaction()`, as we saw above, or,
- in an environment with container-managed transactions, it might obtain it via dependency injection.

Whatever the case, the code which orchestrates a unit of work usually just calls the `Session` or `EntityManager` directly, passing it along to helper methods like our query method if necessary.

```
@GET
@Path("books/{titlePattern}")
public List<Book> findBooks(String titlePattern) {
    var books = sessionFactory.fromTransaction(session ->
        Queries.findBooksByTitleWithPagination(session, titlePattern,
            RESULTS_PER_PAGE, RESULTS_PER_PAGE*page));
    return books.isEmpty() ? Response.status(404).build() : books;
}
```

You might be thinking that our query method looks a bit boilerplatey. That's true, perhaps, but we're much more concerned that it's not very typesafe. Indeed, for many years, the lack of compile-time checking for HQL queries and code which binds arguments to query parameters was our number one source of discomfort with Hibernate.

Fortunately, there's now a solution to both problems: as an incubating feature of Hibernate 6.3, we now offer the possibility to have the Metamodel Generator fill in the implementation of such query methods for you.

Now that we have a rough picture of what our persistence logic might look like, it's natural to ask how we should test this code.

1.6. Testing persistence logic

When we write tests for our persistence logic, we're going to need:

1. a database, with

2. an instance of the schema mapped by our persistent entities, and
3. a set of test data, in a well-defined state at the beginning of each test.

It might seem obvious that we should test against the same database system that we're going to use in production, and, indeed, we should certainly have at least *some* tests for this configuration. But on the other hand, tests which perform I/O are much slower than tests which don't, and most databases can't be set up to run in-process.

So, since most persistence logic written using Hibernate 6 is *extremely* portable between databases, it often makes good sense to test against an in-memory Java database. (H2 is the one we recommend.)



We do need to be careful here if our persistence code uses native SQL, or if it uses concurrency-management features like pessimistic locks.

Whether we're testing against our real database, or against an in-memory Java database, we'll need to export the schema at the beginning of a test suite. We *usually* do this when we create the Hibernate `SessionFactory` or JPA `EntityManager`, and so traditionally we've used a [configuration property](#) for this.

The JPA-standard property is `jakarta.persistence.schema-generation.database.action`. For example, if we're using `Configuration` to configure Hibernate, we could write:

```
configuration.setProperty(AvailableSettings.JAKARTA_HBM2DDL_DATABASE_ACTION,  
                           Action.SPEC_ACTION_DROP_AND_CREATE);
```

Alternatively, in Hibernate 6, we may use the new `SchemaManager` API to export the schema, just as we did [above](#).

```
sessionFactory.getSchemaManager().exportMappedObjects(true);
```

Since executing DDL statements is very slow on many databases, we don't want to do this before every test. Instead, to ensure that each test begins with the test data in a well-defined state, we need to do two things before each test:

1. clean up any mess left behind by the previous test, and then
2. reinitialize the test data.

We may truncate all the tables, leaving an empty database schema, using the `SchemaManager`.

```
sessionFactory.getSchemaManager().truncateMappedObjects();
```

After truncating tables, we might need to initialize our test data. We may specify test data in a SQL script, for example:

`/import.sql`

```
insert into Books (isbn, title) values ('9781932394153', 'Hibernate in Action')  
insert into Books (isbn, title) values ('9781932394887', 'Java Persistence with Hibernate')  
insert into Books (isbn, title) values ('9781617290459', 'Java Persistence with Hibernate, Second Edition')
```

If we name this file `import.sql`, and place it in the root classpath, that's all we need to do.

Otherwise, we need to specify the file in the [configuration property](#) `jakarta.persistence.schema-generation.create-script-source`.

This SQL script will be executed every time `exportMappedObjects()` or `truncateMappedObjects()` is called.



Another important test we'll need is one which validates our [O/R mapping annotations](#) against the actual database schema. This is again the job of the schema management tooling, either:

```
configuration.setProperty(AvailableSettings.JAKARTA_HBM2DDL_DATABASE_ACTION,  
                           Action.ACTION_VALIDATE);
```

Or:

```
sessionFactory.getSchemaManager().validateMappedObjects();
```

This "test" is one which many people like to run even in production, when the system starts up.

1.7. Architecture and the persistence layer

Let's now consider a different approach to code organization, one we treat with suspicion.



In this section, we're going to give you our *opinion*. If you're only interested in facts, or if you prefer not to read things that might undermine the opinion you currently hold, please feel free to skip straight to the [next chapter](#).

Hibernate is an architecture-agnostic library, not a framework, and therefore integrates comfortably with a wide range of Java frameworks and containers. Consistent with our place within the ecosystem, we've historically avoided giving out much advice on architecture. This is a practice we're now perhaps inclined to regret, since the resulting vacuum has come to be filled with advice from people advocating architectures, design patterns, and extra frameworks which we suspect make Hibernate a bit less pleasant to use than it should be.

In particular, frameworks which wrap JPA seem to add bloat while subtracting some of the fine-grained control over data access that Hibernate works so hard to provide. These frameworks don't expose the full feature set of Hibernate, and so the program is forced to work with a less powerful abstraction.

The stodgy, dogmatic, *conventional* wisdom, which we hesitate to challenge for simple fear of pricking ourselves on the erect hackles that inevitably accompany such dogma-baiting is:

Code which interacts with the database belongs in a separate persistence layer.

We lack the courage—perhaps even the conviction—to tell you categorically to *not* follow this recommendation. But we do ask you to consider the cost in boilerplate of any architectural layer, and whether the benefits this cost buys are really worth it in the context of your system.

To add a little background texture to this discussion, and at the risk of our Introduction degenerating into a rant at such an early stage, we're going to ask you to humor us while talk a little more about ancient history.

An epic tale of DAOs and Repositories

Back in the dark days of Java EE 4, before the standardization of Hibernate, and subsequent ascendance of JPA in Java enterprise development, it was common to hand-code the messy JDBC interactions that Hibernate takes care of today. In those terrible times, a pattern arose that we used to call *Data Access Objects* (DAOs). A DAO gave you a place to put all that nasty JDBC code, leaving the important program logic cleaner.

When Hibernate arrived suddenly on the scene in 2001, developers loved it. But Hibernate implemented no specification, and many wished to reduce or at least *localize* the dependence of their project logic on Hibernate. An obvious solution was to keep the DAOs around, but to replace the JDBC code inside them with calls to the Hibernate *Session*.

We partly blame ourselves for what happened next.

Back in 2002 and 2003 this really seemed like a pretty reasonable thing to do. In fact, we contributed to the popularity of this approach by recommending—or at least not discouraging—the use of DAOs in *Hibernate in Action*. We hereby apologize for this mistake, and for taking much too long to recognize it.


Eventually, some folks came to believe that their DAOs shielded their program from depending in a hard way on ORM, allowing them to "swap out" Hibernate, and replace it with JDBC, or with something else. In fact, this was never really true—there's quite a deep difference between the programming model of JDBC, where every interaction with the database is explicit and synchronous, and the programming model of stateful sessions in Hibernate, where updates are implicit, and SQL statements are executed asynchronously.

But then the whole landscape for persistence in Java changed in April 2006, when the final draft of JPA 1.0 was approved. Java now had a standard way to do ORM, with multiple high-quality implementations of the standard API. This was the end of the line for the DAOs, right?

Well, no. It wasn't. DAOs were rebranded "repositories", and continue to enjoy a sort-of zombie afterlife as a front-end to JPA. But are they really pulling their weight, or are they just unnecessary extra complexity and bloat? An extra layer of indirection that makes stack traces harder to read and code harder to debug?

Our considered view is that they're mostly just bloat. The JPA *EntityManager* is a "repository", and it's a standard repository with a well-defined specification written by people who spend all day thinking about persistence. If these repository frameworks offered anything actually *useful*—and not obviously foot-shooty—over and above what *EntityManager* provides, we would have already added it to *EntityManager* decades ago.

Ultimately, we're not sure you need a separate persistence layer at all. At least *consider* the possibility that it might be OK to call the *EntityManager* directly from your business logic.



We can already hear you hissing at our heresy. But before slamming shut the lid of your laptop and heading off to fetch garlic and a pitchfork, take a couple of hours to weigh what we're proposing.

OK, so, look, if it makes you feel better, one way to view `EntityManager` is to think of it as a single *generic* "repository" that works for every entity in your system. From this point of view, JPA is your persistence layer. And there's few good reasons to wrap this abstraction in a second abstraction that's *less* generic.

Even where a distinct persistence layer *is* appropriate, DAO-style repositories aren't the unambiguously most-correct way to factorize the equation:

- most nontrivial queries touch multiple entities, and so it's often quite ambiguous which repository such a query belongs to,
- most queries are extremely specific to a particular fragment of program logic, and aren't reused in different places across the system, and
- the various operations of a repository rarely interact or share common internal implementation details.

Indeed, repositories, by nature, exhibit very low *cohesion*. A layer of repository objects might make sense if you have multiple implementations of each repository, but in practice almost nobody ever does. That's because they're also extremely highly *coupled* to their clients, with a very large API surface. And, on the contrary, a layer is only easily replaceable if it has a very *narrow* API.



Some people do indeed use mock repositories for testing, but we really struggle to see any value in this. If you don't want to run your tests against our real database, it's usually very easy to "mock" the database itself by running tests against an in-memory Java database like H2. This works even better in Hibernate 6 than in older versions of Hibernate, since HQL is now *much* more portable between platforms.

Phew, let's move on.

1.8. Overview

It's now time to begin our journey toward actually *understanding* the code we saw earlier.

This introduction will guide you through the basic tasks involved in developing a program that uses Hibernate for persistence:

1. configuring and bootstrapping Hibernate, and obtaining an instance of `SessionFactory` or `EntityManagerFactory`,
2. writing a *domain model*, that is, a set of *entity classes* which represent the persistent types in your program, and which map to tables of your database,
3. customizing these mappings when the model maps to a pre-existing relational schema,
4. using the `Session` or `EntityManager` to perform operations which query the database and return entity instances, or which update the data held in the database,
5. writing complex queries using the Hibernate Query Language (HQL) or native SQL, and, finally
6. tuning performance of the data access logic.

Naturally, we'll start at the top of this list, with the least-interesting topic: *configuration*.

Chapter 2. Configuration and bootstrap

We would love to make this section short. Unfortunately, there's several distinct ways to configure and bootstrap Hibernate, and we're going to have to describe at least two of them in detail.

The four basic ways to obtain an instance of Hibernate are shown in the following table:

Using the standard JPA-defined XML, and the operation <code>Persistence.createEntityManagerFactory()</code>	Usually chosen when portability between JPA implementations is important.
Using the <code>Configuration</code> class to construct a <code>SessionFactory</code>	When portability between JPA implementations is not important, this option is quicker, adds some flexibility and saves a typecast.
Using the more complex APIs defined in <code>org.hibernate.boot</code>	Used primarily by framework integrators, this option is outside the scope of this document.
By letting the container take care of the bootstrap process and of injecting the <code>SessionFactory</code> or <code>EntityManagerFactory</code>	Used in a container environment like WildFly or Quarkus.

Here we'll focus on the first two options.

Hibernate in containers

Actually, the last option is extremely popular, since every major Java application server and microservice framework comes with built-in support for Hibernate. Such container environments typically also feature facilities to automatically manage the lifecycle of an `EntityManager` or `Session` and its association with container-managed transactions.

To learn how to configure Hibernate in such a container environment, you'll need to refer to the documentation of your chosen container. For Quarkus, here's the [relevant documentation](#).

If you're using Hibernate outside of a container environment, you'll need to:

- include Hibernate ORM itself, along with the appropriate JDBC driver, as dependencies of your project, and
- configure Hibernate with information about your database, by specifying configuration properties.

2.1. Including Hibernate in your project build

First, add the following dependency to your project:

```
org.hibernate.orm:hibernate-core:{version}
```

Where `{version}` is the version of Hibernate you're using.

You'll also need to add a dependency for the JDBC driver for your database.

Table 2. JDBC driver dependencies

Database	Driver dependency
PostgreSQL or CockroachDB	<code>org.postgresql:postgresql:{version}</code>
MySQL or TiDB	<code>com.mysql:mysql-connector-j:{version}</code>
MariaDB	<code>org.mariadb.jdbc:mariadb-java-client:{version}</code>
DB2	<code>com.ibm.db2:jcc:{version}</code>
SQL Server	<code>com.microsoft.sqlserver:mssql-jdbc:\${version}</code>
Oracle	<code>com.oracle.database.jdbc:ojdbc11:\${version}</code>
H2	<code>com.h2database:h2:{version}</code>
HSQLDB	<code>org.hsqldb:hsqldb:{version}</code>

Where {version} is the latest version of the JDBC driver for your database.

2.2. Optional dependencies

Optionally, you might also add any of the following additional features:

Table 3. Optional dependencies

Optional feature	Dependencies
An SLF4J logging implementation	org.apache.logging.log4j:log4j-core or org.slf4j:slf4j-jdk14
A JDBC connection pool, for example, Agroal	org.hibernate.orm:hibernate-agroal and io.agroal:agroal-pool
The Hibernate Metamodel Generator , especially if you're using the JPA criteria query API	org.hibernate.orm:hibernate-jpamodelgen
The Query Validator , for compile-time checking of HQL	org.hibernate:query-validator
Hibernate Validator , an implementation of Bean Validation	org.hibernate.validator:hibernate-validator and org.glassfish:jakarta.el
Local second-level cache support via JCache and EHCache	org.hibernate.orm:hibernate-jcache and org.ehcache:ehcache
Local second-level cache support via JCache and Caffeine	org.hibernate.orm:hibernate-jcache and com.github.ben-manes.caffeine:jcache
Distributed second-level cache support via Infinispan	org.infinispan:infinispan-hibernate-cache-v60
A JSON serialization library for working with JSON datatypes, for example, Jackson or Yasson	com.fasterxml.jackson.core:jackson-databind or org.eclipse:yasson
Hibernate Spatial	org.hibernate.orm:hibernate-spatial
Envers , for auditing historical data	org.hibernate.orm:hibernate-envers

You might also add the [Hibernate bytecode enhancer](#) to your Gradle build if you want to use [field-level lazy fetching](#).

2.3. Configuration using JPA XML

Sticking to the JPA-standard approach, we would provide a file named `persistence.xml`, which we usually place in the `META-INF` directory of a *persistence archive*, that is, of the `.jar` file or directory which contains our entity classes.

```
<persistence xmlns="http://java.sun.com/xml/ns/persistence"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://java.sun.com/xml/ns/persistence
https://jakarta.ee/xml/ns/persistence/persistence_3_0.xsd"
  version="2.0">

  <persistence-unit name="org.hibernate.example">

    <class>org.hibernate.example.Book</class>
    <class>org.hibernate.example.Author</class>

    <properties>
      <!-- PostgreSQL -->
      <property name="jakarta.persistence.jdbc.url"
        value="jdbc:postgresql://localhost/example"/>

      <!-- Credentials -->
      <property name="jakarta.persistence.jdbc.user"
        value="gavin"/>
      <property name="jakarta.persistence.jdbc.password"
        value="hibernate"/>

      <!-- Automatic schema export -->
```

```

        <property name="jakarta.persistence.schema-generation.database.action"
            value="drop-and-create"/>

        <!-- SQL statement logging -->
        <property name="hibernate.show_sql" value="true"/>
        <property name="hibernate.format_sql" value="true"/>
        <property name="hibernate.highlight_sql" value="true"/>

    </properties>

</persistence-unit>

</persistence>

```

The `<persistence-unit>` element defines a named *persistence unit*, that is:

- a collection of associated entity types, along with
- a set of default configuration settings, which may be augmented or overridden at runtime.

Each `<class>` element specifies the fully-qualified name of an entity class.

Scanning for entity classes

In some container environments, for example, in any EE container, the `<class>` elements are unnecessary, since the container will scan the archive for annotated classes, and automatically recognize any class annotated `@Entity`.

Each `<property>` element specifies a *configuration property* and its value. Note that:

- the configuration properties in the `jakarta.persistence` namespace are standard properties defined by the JPA spec, and
- properties in the `hibernate` namespace are specific to Hibernate.

We may obtain an `EntityManagerFactory` by calling `Persistence.createEntityManagerFactory()`:

```

EntityManagerFactory entityManagerFactory =
    Persistence.createEntityManagerFactory("org.hibernate.example");

```

If necessary, we may override configuration properties specified in `persistence.xml`:

```

EntityManagerFactory entityManagerFactory =
    Persistence.createEntityManagerFactory("org.hibernate.example",
        Map.of(AvailableSettings.JAKARTA_JDBC_PASSWORD, password));

```

2.4. Configuration using Hibernate API

Alternatively, the venerable class `org.hibernate.cfg.Configuration` allows an instance of Hibernate to be configured in Java code:

```

SessionFactory sessionFactory =
    new Configuration()
        .addAnnotatedClass(Book.class)
        .addAnnotatedClass(Author.class)
        // PostgreSQL
        .setProperty(AvailableSettings.JAKARTA_JDBC_URL, "jdbc:postgresql://localhost/example")
        // Credentials
        .setProperty(AvailableSettings.JAKARTA_JDBC_USER, user)
        .setProperty(AvailableSettings.JAKARTA_JDBC_PASSWORD, password)
        // Automatic schema export
        .setProperty(AvailableSettings.JAKARTA_HBM2DDL_DATABASE_ACTION,
            Action.SPEC_ACTION_DROP_AND_CREATE)
        // SQL statement logging
        .setProperty(AvailableSettings.SHOW_SQL, TRUE.toString())
        .setProperty(AvailableSettings.FORMAT_SQL, TRUE.toString())
        .setProperty(AvailableSettings.HIGHLIGHT_SQL, TRUE.toString())
        // Create a new SessionFactory
        .buildSessionFactory();

```

The `Configuration` class has survived almost unchanged since the very earliest (pre-1.0) versions of Hibernate, and so it doesn't look particularly modern. On the other hand, it's very easy to use, and exposes some options that `persistence.xml` doesn't support.

Advanced configuration options

Actually, the `Configuration` class is just a very simple facade for the more modern, much more powerful—but more complex—API defined in the package `org.hibernate.boot`. This API is useful if you have very advanced requirements, for example, if you’re writing a framework or implementing a container. You’ll find more information in the [User Guide](#), and in the `{versionDocBase}/javadocs/org/hibernate/boot/package-summary.html` [package-level documentation] of `org.hibernate.boot`.

2.5. Configuration using Hibernate properties file

If we’re using the Hibernate `Configuration` API, but we don’t want to put certain configuration properties directly in the Java code, we can specify them in a file named `hibernate.properties`, and place the file in the root classpath.

```
# PostgreSQL
jakarta.persistence.jdbc.url=jdbc:postgresql://localhost/example
# Credentials
jakarta.persistence.jdbc.user=hibernate
jakarta.persistence.jdbc.password=zAh7mY$2MNshzAQ5

# SQL statement logging
hibernate.show_sql=true
hibernate.format_sql=true
hibernate.highlight_sql=true
```

2.6. Basic configuration settings

The class `org.hibernate.cfg.AvailableSettings` enumerates all the configuration properties understood by Hibernate.

Of course, we’re not going to cover every useful configuration setting in this chapter. Instead, we’ll mention the ones you need to get started, and come back to some other important settings later, especially when we talk about performance tuning.



Hibernate has many—too many—switches and toggles. Please don’t go crazy messing about with these settings; most of them are rarely needed, and many only exist to provide backward compatibility with older versions of Hibernate. With rare exception, the default behavior of every one of these settings was carefully chosen to be *the behavior we recommend*.

The properties you really do need to get started are these three:

Table 4. JDBC connection settings

Configuration property name	Purpose
<code>jakarta.persistence.jdbc.url</code>	JDBC URL of your database
<code>jakarta.persistence.jdbc.user</code> and <code>jakarta.persistence.jdbc.password</code>	Your database credentials



In Hibernate 6, you don’t need to specify `hibernate.dialect`. The correct Hibernate SQL Dialect will be determined for you automatically. The only reason to specify this property is if you’re using a custom user-written Dialect class.

Similarly, neither `hibernate.connection.driver_class` nor `jakarta.persistence.jdbc.driver` is needed when working with one of the supported databases.

Pooling JDBC connections is an extremely important performance optimization. You can set the size of Hibernate’s built-in connection pool using this property:

Table 5. Built-in connection pool size

Configuration property name	Purpose
<code>hibernate.connection.pool_size</code>	The size of the built-in connection pool



By default, Hibernate uses a simplistic built-in connection pool. This pool is not meant for use in production, and

later, when we discuss performance, we'll see how to [select a more robust implementation](#).

Alternatively, in a container environment, you'll need at least one of these properties:

Table 6. Transaction management settings

Configuration property name	Purpose
<code>jakarta.persistence.transactionType</code>	(Optional, defaults to JTA) Determines if transaction management is via JTA or resource-local transactions. Specify <code>RESOURCE_LOCAL</code> if JTA should not be used.
<code>jakarta.persistence.jtaDataSource</code>	JNDI name of a JTA datasource
<code>jakarta.persistence.nonJtaDataSource</code>	JNDI name of a non-JTA datasource

In this case, Hibernate obtains pooled JDBC database connections from a container-managed `DataSource`.

2.7. Automatic schema export

You can have Hibernate infer your database schema from the mapping annotations you've specified in your Java code, and export the schema at initialization time by specifying one or more of the following configuration properties:

Table 7. Schema management settings

Configuration property name	Purpose
<code>jakarta.persistence.schema-generation.database.action</code>	• If <code>drop-and-create</code>, first drop the schema and then export tables, sequences, and constraints • If <code>create</code>, export tables, sequences, and constraints, without attempting to drop them first • If <code>create-drop</code>, drop the schema and recreate it on <code>SessionFactory</code> startup. Additionally, drop the schema on <code>SessionFactory</code> shutdown • If <code>drop</code>, drop the schema on <code>SessionFactory</code> shutdown • If <code>validate</code>, validate the database schema without changing it • If <code>update</code>, only export what's missing in the schema
<code>jakarta.persistence.create-database-schemas</code>	(Optional) If <code>true</code> , automatically create schemas and catalogs
<code>jakarta.persistence.schema-generation.create-source</code>	(Optional) If <code>metadata-then-script</code> or <code>script-then-metadata</code> , execute an additional SQL script when exported tables and sequences
<code>jakarta.persistence.schema-generation.create-script-source</code>	(Optional) The name of the SQL script to be executed

This feature is extremely useful for testing.



The easiest way to pre-initialize a database with test or "reference" data is to place a list of SQL insert statements in a file named, for example, `import.sql`, and specify the path to this file using the property `jakarta.persistence.schema-generation.create-script-source`.

This approach is cleaner than writing Java code to instantiate entity instances and calling `persist()` on each of them.



Alternatively, the `SchemaManager` API allow you to control schema export programmatically.

```
sessionFactory.getSchemaManager().exportMappedObjects(true);
```

2.8. Logging the generated SQL

To see the generated SQL as it's sent to the database, you have two options.

One way is to set the property `hibernate.show_sql` to `true`, and Hibernate will log SQL direct to the console. You can make the output much more readable by enabling formatting or highlighting. These settings really help when troubleshooting the generated

SQL statements.

Table 8. Settings for SQL logging to the console

Configuration property name	Purpose
<code>hibernate.show_sql</code>	If true, log SQL directly to the console
<code>hibernate.format_sql</code>	If true, log SQL in a multiline, indented format
<code>hibernate.highlight_sql</code>	If true, log SQL with syntax highlighting via ANSI escape codes

Alternatively, you can enable debug-level logging for the category `org.hibernate.SQL` using your preferred SLF4J logging implementation.

For example, if you're using Log4J 2 (as above in [Optional dependencies](#)), add these lines to your `log4j2.properties` file:

```
# SQL execution
logger.hibernate.name = org.hibernate.SQL
logger.hibernate.level = debug

# JDBC parameter binding
logger.jdbc-bind.name=org.hibernate.orm.jdbc.bind
logger.jdbc-bind.level=trace
# JDBC result set extraction
logger.jdbc-extract.name=org.hibernate.orm.jdbc.extract
logger.jdbc-extract.level=trace
```

But with this approach we miss out on the pretty highlighting.

2.9. Minimizing repetitive mapping information

The following properties are very useful for minimizing the amount of information you'll need to explicitly specify in `@Table` and `@Column` annotations, which we'll discuss below in [Object/relational mapping](#):

Table 9. Settings for minimizing explicit mapping information

Configuration property name	Purpose
<code>hibernate.default_schema</code>	A default schema name for entities which do not explicitly declare one
<code>hibernate.default_catalog</code>	A default catalog name for entities which do not explicitly declare one
<code>hibernate.physical_naming_strategy</code>	A <code>PhysicalNamingStrategy</code> implementing your database naming standards
<code>hibernate.implicit_naming_strategy</code>	An <code>ImplicitNamingStrategy</code> which specifies how "logical" names of relational objects should be inferred when no name is specified in annotations



Writing your own `PhysicalNamingStrategy` and/or `ImplicitNamingStrategy` is an especially good way to reduce the clutter of annotations on your entity classes, and to implement your database naming conventions, and so we think you should do it for any nontrivial data model. We'll have more to say about them in [Naming strategies](#).

2.10. Nationalized character data in SQL Server

By default, SQL Server's `char` and `varchar` types don't accommodate Unicode data. But a Java string may contain any Unicode character. So, if you're working with SQL Server, you might need to force Hibernate to use the `nchar` and `nvarchar` column types.

Table 10. Setting the use of nationalized character data

Configuration property name	Purpose
<code>hibernate.use_nationalized_character_data</code>	Use <code>nchar</code> and <code>nvarchar</code> instead of <code>char</code> and <code>varchar</code>

On the other hand, if only *some* columns store nationalized data, use the `@Nationalized` annotation to indicate fields of your entities which map these columns.



Alternatively, you can configure SQL Server to use the UTF-8 enabled collation `_UTF8`.

Chapter 3. Entities

An *entity* is a Java class which represents data in a relational database table. We say that the entity *maps* or *maps to* the table. Much less commonly, an entity might aggregate data from multiple tables, but we'll get to that [later](#).

An entity has *attributes*—properties or fields—which map to columns of the table. In particular, every entity must have an *identifier* or *id*, which maps to the primary key of the table. The id allows us to uniquely associate a row of the table with an instance of the Java class, at least within a given *persistence context*.

We'll explore the idea of a persistence context [later](#). For now, think of it as a one-to-one mapping between ids and entity instances.

An instance of a Java class cannot outlive the virtual machine to which it belongs. But we may think of an entity instance having a lifecycle which transcends a particular instantiation in memory. By providing its id to Hibernate, we may re-materialize the instance in a new persistence context, as long as the associated row is present in the database. Therefore, the operations `persist()` and `remove()` may be thought of as demarcating the beginning and end of the lifecycle of an entity, at least with respect to persistence.

Thus, an id represents the *persistent identity* of an entity, an identity that outlives a particular instantiation in memory. And this is an important difference between entity class itself and the values of its attributes—the entity has a persistent identity, and a well-defined lifecycle with respect to persistence, whereas a `String` or `List` representing one of its attribute values doesn't.

An entity usually has associations to other entities. Typically, an association between two entities maps to a foreign key in one of the database tables. A group of mutually associated entities is often called a *domain model*, though *data model* is also a perfectly good term.

3.1. Entity classes

An entity must:

- be a non-final class,
- with a non-private constructor with no parameters.

On the other hand, the entity class may be either concrete or abstract, and it may have any number of additional constructors.



An entity class may be a static inner class.

Every entity class must be annotated `@Entity`.

```
@Entity
class Book {
    Book() {}
    ...
}
```

Alternatively, the class may be identified as an entity type by providing an XML-based mapping for the class.

Mapping entities using XML

When XML-based mappings are used, the `<entity>` element is used to declare an entity class:

```
<entity-mappings>
  <package>org.hibernate.example</package>

  <entity class="Book">
    <attributes> ... </attributes>
  </entity>

  ...
</entity-mappings>
```

Since the `orm.xml` mapping file format defined by the JPA specification was modelled closely on the annotation-based mappings, it's usually easy to go back and forth between the two options.

We won't have much more to say about XML-based mappings in this Introduction, since it's not our preferred way to do things.

"Dynamic" models

We love representing entities as classes because the classes give us a *type-safe* model of our data. But Hibernate also has the ability to represent entities as detyped instances of `java.util.Map`. There's information in the [User Guide](#), if you're curious.

This must sound like a weird feature for a project that places importance on type-safety. Actually, it's a useful capability for a very particular sort of generic code. For example, [Hibernate Envers](#) is a great auditing/versioning system for Hibernate entities. Envers makes use of maps to represent its *versioned model* of the data.

3.2. Access types

Each entity class has a default *access type*, either:

- *direct field access*, or
- *property access*.

Hibernate automatically determines the access type from the location of attribute-level annotations. Concretely:

- if a field is annotated `@Id`, field access is used, or
- if a getter method is annotated `@Id`, property access is used.

Back when Hibernate was just a baby, property access was quite popular in the Hibernate community. Today, however, field access is *much* more common.



The default access type may be specified explicitly using the `@Access` annotation, but we strongly discourage this, since it's ugly and never necessary.



Mapping annotations should be placed consistently:

- if `@Id` annotates a field, the other mapping annotations should also be applied to fields, or,
- if `@Id` annotates a getter, the other mapping annotations should be applied to getters.

It is in principle possible to mix field and property access using explicit `@Access` annotations at the attribute level. We don't recommend doing this.

An entity class like `Book`, which does not extend any other entity class is called a *root entity*. Every root entity must declare an identifier attribute.

3.3. Entity class inheritance

An entity class may extend another entity class.

```
@Entity
class AudioBook extends Book {
    AudioBook() {}
    ...
}
```

A subclass entity inherits every persistent attribute of every entity it extends.

A root entity may also extend another class and inherit mapped attributes from the other class. But in this case, the class which declares the mapped attributes must be annotated `@MappedSuperclass`.

```
@MappedSuperclass
class Versioned {
    ...
}

@Entity
class Book extends Versioned {
    ...
}
```

A root entity class must declare an attribute annotated `@Id`, or inherit one from a `@MappedSuperclass`. A subclass entity always inherits the identifier attribute of the root entity. It may not declare its own `@Id` attribute.

3.4. Identifier attributes

An identifier attribute is usually a field:

```
@Entity
class Book {
    Book() {}

    @Id
    Long id;

    ...
}
```

But it may be a property:

```
@Entity
class Book {
    Book() {}

    private Long id;

    @Id
    Long getId() { return id; }
    void setId(Long id) { this.id = id; }

    ...
}
```

An identifier attribute must be annotated `@Id` or `@EmbeddedId`.

Identifier values may be:

- assigned by the application, that is, by your Java code, or
- generated and assigned by Hibernate.

We'll discuss the second option first.

3.5. Generated identifiers

An identifier is often system-generated, in which case it should be annotated `@GeneratedValue`:

```
@Id @GeneratedValue
Long id;
```



System-generated identifiers, or *surrogate keys* make it easier to evolve or refactor the relational data model. If you have the freedom to define the relational schema, we recommend the use of surrogate keys. On the other hand, if, as is more common, you're working with a pre-existing database schema, you might not have the option.

JPA defines the following strategies for generating ids, which are enumerated by `GenerationType`:

Table 11. Standard id generation strategies

Strategy	Java type	Implementation
<code>GenerationType.UUID</code>	UUID or String	A Java UUID
<code>GenerationType.IDENTITY</code>	Long or Integer	An identity or autoincrement column
<code>GenerationType.SEQUENCE</code>	Long or Integer	A database sequence
<code>GenerationType.TABLE</code>	Long or Integer	A database table
<code>GenerationType.AUTO</code>	Long or Integer	Selects SEQUENCE, TABLE, or UUID based on the identifier type and capabilities of the database

For example, this UUID is generated in Java code:

```
@Id @GeneratedValue UUID id; // AUTO strategy selects UUID based on the field type
```

This id maps to a SQL identity, auto_increment, or bigserial column:

```
@Id @GeneratedValue(strategy=IDENTITY) Long id;
```

The @SequenceGenerator and @TableGenerator annotations allow further control over SEQUENCE and TABLE generation respectively.

Consider this sequence generator:

```
@SequenceGenerator(name="bookSeq", sequenceName="seq_book", initialValue = 5, allocationSize=10)
```

Values are generated using a database sequence defined as follows:

```
create sequence seq_book start with 5 increment by 10
```

Notice that Hibernate doesn't have to go to the database every time a new identifier is needed. Instead, a given process obtains a block of ids, of size allocationSize, and only needs to hit the database each time the block is exhausted. Of course, the downside is that generated identifiers are not contiguous.



If you let Hibernate export your database schema, the sequence definition will have the right start with and increment values. But if you're working with a database schema managed outside Hibernate, make sure the initialValue and allocationSize members of @SequenceGenerator match the start with and increment specified in the DDL.

Any identifier attribute may now make use of the generator named bookSeq:

```
@Id
@GeneratedValue(strategy=SEQUENCE, generator="bookSeq") // reference to generator defined elsewhere
Long id;
```

Actually, it's extremely common to place the @SequenceGenerator annotation on the @Id attribute that makes use of it:

```
@Id
@GeneratedValue(strategy=SEQUENCE, generator="bookSeq") // reference to generator defined below
@SequenceGenerator(name="bookSeq", sequenceName="seq_book", initialValue = 5, allocationSize=10)
Long id;
```



JPA id generators may be shared between entities. A @SequenceGenerator or @TableGenerator must have a name, and may be shared between multiple id attributes. This fits somewhat uncomfortably with the common practice of annotating the @Id attribute which makes use of the generator!

As you can see, JPA provides quite adequate support for the most common strategies for system-generated ids. However, the annotations themselves are a bit more intrusive than they should be, and there's no well-defined way to extend this framework to support custom strategies for id generation. Nor may @GeneratedValue be used on a property not annotated @Id. Since custom id generation is a rather common requirement, Hibernate provides a very carefully-designed framework for user-defined Generators, which we'll discuss in [User-defined generators](#).

3.6. Natural keys as identifiers

Not every identifier attribute maps to a (system-generated) surrogate key. Primary keys which are meaningful to the user of the system are called *natural keys*.

When the primary key of a table is a natural key, we don't annotate the identifier attribute @GeneratedValue, and it's the responsibility of the application code to assign a value to the identifier attribute.

```
@Entity
class Book {
    @Id
    String isbn;

    ...
}
```

Of particular interest are natural keys which comprise more than one database column, and such natural keys are called *composite keys*.

3.7. Composite identifiers

If your database uses composite keys, you'll need more than one identifier attribute. There are two ways to map composite keys in JPA:

- using an `@IdClass`, or
- using an `@EmbeddedId`.

Perhaps the most immediately-natural way to represent this in an entity class is with multiple fields annotated `@Id`, for example:

```
@Entity
@IdClass(BookId.class)
class Book {
    Book() {}

    @Id
    String isbn;

    @Id
    int printing;

    ...
}
```

But this approach comes with a problem: what object can we use to identify a `Book` and pass to methods like `find()` which accept an identifier?

The solution is to write a separate class with fields that match the identifier attributes of the entity. The `@IdClass` annotation of the `Book` entity identifies the id class to use for that entity:

```
class BookId {

    String isbn;
    int printing;

    BookId() {}

    BookId(String isbn, int printing) {
        this.isbn = isbn;
        this.printing = printing;
    }

    @Override
    public boolean equals(Object other) {
        if (other instanceof BookId) {
            BookId bookId = (BookId) other;
            return bookId.isbn.equals(isbn)
                && bookId.printing == printing;
        }
        else {
            return false;
        }
    }

    @Override
    public int hashCode() {
        return isbn.hashCode();
    }
}
```

Every id class should override `equals()` and `hashCode()`.

This is not our preferred approach. Instead, we recommend that the `BookId` class be declared as an `@Embeddable` type:

```
@Embeddable
class BookId {

    String isbn;
```

```

    int printing;

    BookId() {}

    BookId(String isbn, int printing) {
        this.isbn = isbn;
        this.printing = printing;
    }

    ...
}

```

We'll learn more about [Embeddable objects](#) below.

Now the entity class may reuse this definition using `@EmbeddedId`, and the `@IdClass` annotation is no longer required:

```

@Entity
class Book {
    Book() {}

    @EmbeddedId
    BookId bookId;

    ...
}

```

This second approach eliminates some duplicated code.

Either way, we may now use `BookId` to obtain instances of `Book`:

```

Book book = session.find(Book.class, new BookId(isbn, printing));

```

3.8. Version attributes

An entity may have an attribute which is used by Hibernate for optimistic lock checking. A version attribute is usually of type `Integer`, `Short`, `Long`, `LocalDateTime`, `OffsetDateTime`, `ZonedDateTime`, or `Instant`.

```

@Version
LocalDateTime lastUpdated;

```

Version attributes are automatically assigned by Hibernate when an entity is made persistent, and automatically incremented or updated each time the entity is updated.



If an entity doesn't have a version number, which often happens when mapping legacy data, we can still do optimistic locking. The `@OptimisticLocking` annotation lets us specify that optimistic locks should be checked by validating the values of ALL fields, or only the DIRTY fields of the entity. And the `@OptimisticLock` annotation lets us selectively exclude certain fields from optimistic locking.

The `@Id` and `@Version` attributes we've already seen are just specialized examples of *basic attributes*.

3.9. Natural id attributes

Even when an entity has a surrogate key, it should always be possible to write down a combination of fields which uniquely identifies an instance of the entity, from the point of view of the user of the system. This combination of fields is its natural key. Above, we [considered](#) the case where the natural key coincides with the primary key. Here, the natural key is a second unique key of the entity, distinct from its surrogate primary key.



If you can't identify a natural key, it might be a sign that you need to think more carefully about some aspect of your data model. If an entity doesn't have a meaningful unique key, then it's impossible to say what event or object it represents in the "real world" outside your program.

Since it's *extremely* common to retrieve an entity based on its natural key, Hibernate has a way to mark the attributes of the entity which make up its natural key. Each attribute must be annotated `@NaturalId`.

```

@Entity
class Book {

```

```

Book() {}

@Id @GeneratedValue
Long id; // the system-generated surrogate key

@NaturalId
String isbn; // belongs to the natural key

@NaturalId
int printing; // also belongs to the natural key

...
}

```

Hibernate automatically generates a UNIQUE constraint on the columns mapped by the annotated fields.



Consider using the natural id attributes to implement `equals()` and `hashCode()`.

The payoff for doing this extra work, as we will see [much later](#), is that we can take advantage of optimized natural id lookups that make use of the second-level cache.

Note that even when you've identified a natural key, we still recommend the use of a generated surrogate key in foreign keys, since this makes your data model *much* easier to change.

3.10. Basic attributes

A *basic* attribute of an entity is a field or property which maps to a single column of the associated database table. The JPA specification defines a quite limited set of basic types:

Table 12. JPA-standard basic attribute types

Classification	Package	Types
Primitive types		boolean, int, double, etc
Primitive wrappers	java.lang	Boolean, Integer, Double, etc
Strings	java.lang	String
Arbitrary-precision numeric types	java.math	BigInteger, BigDecimal
Date/time types	java.time	LocalDate, LocalTime, LocalDateTime, OffsetDateTime, Instant
Deprecated date/time types ☠	java.util	Date, Calendar
Deprecated JDBC date/time types ☠	java.sql	Date, Time, Timestamp
Binary and character arrays		byte[], char[]
UUIDs	java.util	UUID
Enumerated types		Any enum
Serializable types		Any type which implements <code>java.io.Serializable</code>



We're begging you to use types from the `java.time` package instead of anything which inherits `java.util.Date`.



Serializing a Java object and storing its binary representation in the database is usually wrong. As we'll soon see in [Embeddable objects](#), Hibernate has much better ways to handle complex Java objects.

Hibernate slightly extends this list with the following types:

Table 13. Additional basic attribute types in Hibernate

Classification	Package	Types
Additional date/time types	java.time	Duration, ZoneId, ZoneOffset, Year, and even ZonedDateTime
JDBC LOB types	java.sql	Blob, Clob, NClob
Java class object	java.lang	Class
Miscellaneous types	java.util	Currency, URL, TimeZone

The `@Basic` annotation explicitly specifies that an attribute is basic, but it's often not needed, since attributes are assumed basic by default. On the other hand, if a non-primitively-typed attribute cannot be null, use of `@Basic(optional=false)` is highly recommended.

```
@Basic(optional=false) String firstName;
@Basic(optional=false) String lastName;
String middleName; // may be null
```

Note that primitively-typed attributes are inferred NOT NULL by default.

How to make a column not null in JPA

There are two standard ways to add a NOT NULL constraint to a mapped column in JPA:

- using `@Basic(optional=false)`, or
- using `@Column(nullable=false)`.

You might wonder what the difference is.

Well, it's perhaps not obvious to a casual user of the JPA annotations, but they actually come in two "layers":

- annotations like `@Entity`, `@Id`, and `@Basic` belong to the *logical* layer, the subject of the current chapter—they specify the semantics of your Java domain model, whereas
- annotations like `@Table` and `@Column` belong to the *mapping* layer, the topic of the [next chapter](#)—they specify how elements of the domain model map to objects in the relational database.

Information may be inferred from the logical layer down to the mapping layer, but is never inferred in the opposite direction.

Now, the `@Column` annotation, to whom we'll be properly [introduced](#) a bit later, belongs to the *mapping* layer, and so its `nullable` member only affects schema generation (resulting in a not null constraint in the generated DDL). On the other hand, the `@Basic` annotation belongs to the logical layer, and so an attribute marked `optional=false` is checked by Hibernate before it even writes an entity to the database. Note that:

- `optional=false` implies `nullable=false`, but
- `nullable=false` does *not* imply `optional=false`.

Therefore, we prefer `@Basic(optional=false)` to `@Column(nullable=false)`.



But wait! An even better solution is to use the `@NotNull` annotation from Bean Validation. Just add Hibernate Validator to your project build, as described in [Optional dependencies](#).

3.11. Enumerated types

We included Java `enums` on the list above. An enumerated type is considered a sort of basic type, but since most databases don't have a native `ENUM` type, JPA provides a special `@Enumerated` annotation to specify how the enumerated values should be represented in the database:

- by default, an enum is stored as an integer, the value of its `ordinal()` member, but
- if the attribute is annotated `@Enumerated(STRING)`, it will be stored as a string, the value of its `name()` member.

```
//here, an ORDINAL encoding makes sense
@Enumerated
@Basic(optional=false)
DayOfWeek dayOfWeek;
```

```
//but usually, a STRING encoding is better
@Enumerated(EnumType.STRING)
@Basic(optional=false)
Status status;
```

In Hibernate 6, an enum annotated `@Enumerated(String)` is mapped to:

- a VARCHAR column type with a CHECK constraint on most databases, or
- an ENUM column type on MySQL.

Any other enum is mapped to a TINYINT column with a CHECK constraint.



JPA picks the wrong default here. In most cases, storing an integer encoding of the enum value makes the relational data harder to interpret.

Even considering `DayOfWeek`, the encoding to integers is ambiguous. If you check `java.time.DayOfWeek`, you'll notice that SUNDAY is encoded as 6. But in the country I was born, SUNDAY is the *first* day of the week!

So we prefer `@Enumerated(String)` for most enum attributes.

An interesting special case is PostgreSQL. Postgres supports *named* ENUM types, which must be declared using a DDL `CREATE TYPE` statement. Sadly, these ENUM types aren't well-integrated with the language nor well-supported by the Postgres JDBC driver, so Hibernate doesn't use them by default. But if you would like to use a named enumerated type on Postgres, just annotate your enum attribute like this:

```
@JdbcTypeCode(SqlTypes.NAMED_ENUM)
@Basic(optional=false)
Status status;
```

The limited set of pre-defined basic attribute types can be stretched a bit further by supplying a *converter*.

3.12. Converters

A JPA `AttributeConverter` is responsible for:

- converting a given Java type to one of the types listed above, and/or
- perform any other sort of pre- and post-processing you might need to perform on a basic attribute value before writing and reading it to or from the database.

Converters substantially widen the set of attribute types that can be handled by JPA.

There are two ways to apply a converter:

- the `@Convert` annotation applies an `AttributeConverter` to a particular entity attribute, or
- the `@Converter` annotation (or, alternatively, the `@ConverterRegistration` annotation) registers an `AttributeConverter` for automatic application to all attributes of a given type.

For example, the following converter will be automatically applied to any attribute of type `BitSet`, and takes care of persisting the `BitSet` to a column of type `varbinary`:

```
@Converter(autoApply = true)
public static class EnumSetConverter implements AttributeConverter<EnumSet<DayOfWeek>, Integer> {
    @Override
    public Integer convertToDatabaseColumn(EnumSet<DayOfWeek> enumSet) {
        int encoded = 0;
        var values = DayOfWeek.values();
        for (int i = 0; i < values.length; i++) {
            if (enumSet.contains(values[i])) {
                encoded |= 1 << i;
            }
        }
        return encoded;
    }

    @Override
    public EnumSet<DayOfWeek> convertToEntityAttribute(Integer encoded) {
        var set = EnumSet.noneOf(DayOfWeek.class);
        var values = DayOfWeek.values();
        for (int i = 0; i < values.length; i++) {
            if ((encoded > 0) & (1 << i) == encoded) {
                set.add(values[i]);
            }
        }
        return set;
    }
}
```

```

        if (((1<<i) & encoded) != 0) {
            set.add(values[i]);
        }
    }
    return set;
}
}

```

On the other hand, if we *don't* set `autoapply=true`, then we must explicitly apply the converter using the `@Convert` annotation:

```

@Convert(converter = BitSetConverter.class)
@Basic(optional = false)
BitSet bitset;

```

All this is nice, but it probably won't surprise you that Hibernate goes beyond what is required by JPA.

3.13. Compositional basic types

Hibernate considers a "basic type" to be formed by the marriage of two objects:

- a `JavaType`, which models the semantics of a certain Java class, and
- a `JdbcType`, representing a SQL type which is understood by JDBC.

When mapping a basic attribute, we may explicitly specify a `JavaType`, a `JdbcType`, or both.

JavaType

An instance of `org.hibernate.type.descriptor.java.JavaType` represents a particular Java class. It's able to:

- compare instances of the class to determine if an attribute of that class type is dirty (modified),
- produce a useful hash code for an instance of the class,
- coerce values to other types, and, in particular,
- convert an instance of the class to one of several other equivalent Java representations at the request of its partner `JdbcType`.

For example, `IntegerJavaType` knows how to convert an `Integer` or `int` value to the types `Long`, `BigInteger`, and `String`, among others.

We may explicitly specify a Java type using the `@JavaType` annotation, but for the built-in `JavaTypes` this is never necessary.

```

@JavaType(LongJavaType.class) // not needed, this is the default JavaType for long
long currentTimeMillis;

```

For a user-written `JavaType`, the annotation is more useful:

```

@JavaType(BitSetJavaType.class)
BitSet bitSet;

```

Alternatively, the `@JavaTypeRegistration` annotation may be used to register `BitSetJavaType` as the default `JavaType` for `BitSet`.

JdbcType

A `org.hibernate.type.descriptor.jdbc.JdbcType` is able to read and write a single Java type from and to JDBC.

For example, `VarcharJdbcType` takes care of:

- writing Java strings to JDBC `PreparedStatement`s by calling `setString()`, and
- reading Java strings from JDBC `ResultSet`s using `getString()`.

By pairing `LongJavaType` with `VarcharJdbcType` in holy matrimony, we produce a basic type which maps `Long`s and primitive `long`s to the SQL type `VARCHAR`.

We may explicitly specify a JDBC type using the `@JdbcType` annotation.

```

@JdbcType(VarcharJdbcType.class)
long currentTimeMillis;

```

Alternatively, we may specify a JDBC type code:


```
@JdbcTypeCode(Types.VARCHAR)
long currentTimeMillis;
```

The `@JdbcTypeRegistration` annotation may be used to register a user-written `JdbcType` as the default for a given SQL type code.

JDBC types and JDBC type codes

The types defined by the JDBC specification are enumerated by the integer type codes in the class `java.sql.Types`. Each JDBC type is an abstraction of a commonly-available type in SQL. For example, `Types.VARCHAR` represents the SQL type VARCHAR (or VARCHAR2 on Oracle).

Since Hibernate understand more SQL types than JDBC, there's an extended list of integer type codes in the class `org.hibernate.type.SqlTypes`. For example, `SqlTypes.GEOMETRY` represents the spatial data type GEOMETRY.

AttributeConverter

If a given `JavaType` doesn't know how to convert its instances to the type required by its partner `JdbcType`, we must help it out by providing a JPA `AttributeConverter` to perform the conversion.

For example, to form a basic type using `LongJavaType` and `TimestampJdbcType`, we would provide an `AttributeConverter<Long, Timestamp>`.

```
@JdbcType(TimestampJdbcType.class)
@Convert(converter = LongToTimestampConverter.class)
long currentTimeMillis;
```

Let's abandon our analogy right here, before we start calling this basic type a "throuple".

3.14. Embeddable objects

An embeddable object is a Java class whose state maps to multiple columns of a table, but which doesn't have its own persistent identity. That is, it's a class with mapped attributes, but no `@Id` attribute.

An embeddable object can only be made persistent by assigning it to the attribute of an entity. Since the embeddable object does not have its own persistent identity, its lifecycle with respect to persistence is completely determined by the lifecycle of the entity to which it belongs.

An embeddable class must be annotated `@Embeddable` instead of `@Entity`.

```
@Embeddable
class Name {

    @Basic(optional=false)
    String firstName;

    @Basic(optional=false)
    String lastName;

    String middleName;

    Name() {}

    Name(String firstName, String middleName, String lastName) {
        this.firstName = firstName;
        this.middleName = middleName;
        this.lastName = lastName;
    }

    ...
}
```

An embeddable class must satisfy the same requirements that entity classes satisfy, with the exception that an embeddable class has no `@Id` attribute. In particular, it must have a constructor with no parameters.

Alternatively, an embeddable type may be defined as a Java record type:

```
@Embeddable
```

```
record Name(String firstName, String middleName, String lastName) {}
```

In this case, the requirement for a constructor with no parameters is relaxed.



Unfortunately, as of May 2023, Java record types still cannot be used as `@EmbeddedIds`.

We may now use our `Name` class (or record) as the type of an entity attribute:

```
@Entity
class Author {
    @Id @GeneratedValue
    Long id;

    Name name;

    ...
}
```

Embeddable types can be nested. That is, an `@Embeddable` class may have an attribute whose type is itself a different `@Embeddable` class.



JPA provides an `@Embedded` annotation to identify an attribute of an entity that refers to an embeddable type. This annotation is completely optional, and so we don't usually use it.

On the other hand a reference to an embeddable type is *never* polymorphic. One `@Embeddable` class `F` may inherit a second `@Embeddable` class `E`, but an attribute of type `E` will always refer to an instance of that concrete class `E`, never to an instance of `F`.

Usually, embeddable types are stored in a "flattened" format. Their attributes map columns of the table of their parent entity. Later, in [Mapping embeddable types to UDTs or to JSON](#), we'll see a couple of different options.

An attribute of embeddable type represents a relationship between a Java object with a persistent identity, and a Java object with no persistent identity. We can think of it as a whole/part relationship. The embeddable object belongs to the entity, and can't be shared with other entity instances. And it exists for only as long as its parent entity exists.

Next we'll discuss a different kind of relationship: a relationship between Java objects that each have their persistent identity and persistence lifecycle.

3.15. Associations


An *association* is a relationship between entities. We usually classify associations based on their *multiplicity*. If `E` and `F` are both entity classes, then:

- a *one-to-one* association relates at most one unique instance `E` with at most one unique instance of `F`,
- a *many-to-one* association relates zero or more instances of `E` with a unique instance of `F`, and
- a *many-to-many* association relates zero or more instances of `E` with zero or more instance of `F`.

An association between entity classes may be either:

- *unidirectional*, navigable from `E` to `F` but not from `F` to `E`, or
- *bidirectional*, and navigable in either direction.

In this example data model, we can see the sorts of associations which are possible:





An astute observer of the diagram above might notice that the relationship we've presented as a unidirectional one-to-one association could reasonably be represented in Java using subtyping. This is quite normal. A one-to-one association is the usual way we implement subtyping in a fully-normalized relational model. It's related to the JOINED inheritance mapping strategy.

There are three annotations for mapping associations: @ManyToOne, @OneToMany, and @ManyToMany. They share some common annotation members:

Table 14. Association-defining annotation members

Member	Interpretation	Default value
cascade	Persistence operations which should cascade to the associated entity; a list of CascadeTypes	{}
fetch	Whether the association is eagerly fetched or may be proxied	• LAZY for @OneToMany and @ManyToMany • EAGER for @ManyToOne ☹️☹️☹️
targetEntity	The associated entity class	Determined from the attribute type declaration
optional	For a @ManyToOne or @OneToOne association, whether the association can be null	true
mappedBy	For a bidirectional association, an attribute of the associated entity which maps the association	By default, the association is assumed unidirectional

We'll explain the effect of these members as we consider the various types of association mapping.

Let's begin with the most common association multiplicity.

3.16. Many-to-one

A many-to-one association is the most basic sort of association we can imagine. It maps completely naturally to a foreign key in the database. Almost all the associations in your domain model are going to be of this form.



Later, we'll see how to map a many-to-one association to an **association table**.

The @ManyToOne annotation marks the "to one" side of the association, so a unidirectional many-to-one association looks like this:

```
class Book {
    @Id @GeneratedValue
    Long id;

    @ManyToOne(fetch=LAZY)
    Publisher publisher;

    ...
}
```

Here, the Book table has a foreign key column holding the identifier of the associated Publisher.



A very unfortunate misfeature of JPA is that @ManyToOne associations are fetched eagerly by default. This is almost never what we want. Almost all associations should be lazy. The only scenario in which fetch=EAGER makes sense is if we think there's always a *very* high probability that the **associated object will be found in the second-level cache**. Whenever this isn't the case, remember to explicitly specify fetch=LAZY.

Most of the time, we would like to be able to easily navigate our associations in both directions. We do need a way to get the Publisher of a given Book, but we would also like to be able to obtain all the Books belonging to a given publisher.

To make this association bidirectional, we need to add a collection-valued attribute to the Publisher class, and annotate it @OneToMany.



Hibernate needs to **proxy** unfetched associations at runtime. Therefore, the many-valued side must be declared using an interface type like Set or List, and never using a concrete type like HashSet or ArrayList.

To indicate clearly that this is a bidirectional association, and to reuse any mapping information already specified in the `Book` entity, we must use the `mappedBy` annotation member to refer back to `Book.publisher`.

```
@Entity
class Publisher {
    @Id @GeneratedValue
    Long id;

    @OneToMany(mappedBy="publisher")
    Set<Book> books;

    ...
}
```

The `Publisher.books` field is called the *unowned* side of the association.

Now, we passionately *hate* the stringly-typed `mappedBy` reference to the owning side of the association. Thankfully, the [Metamodel Generator](#) gives us a way to make it a bit more typesafe:

```
@OneToMany(mappedBy=Book_.PUBLISHER) // get used to doing it this way!
Set<Book> books;
```

We're going to use this approach for the rest of the Introduction.

To modify a bidirectional association, we must change the *owning side*.



Changes made to the unowned side of an association are never synchronized to the database. If we desire to change an association in the database, we must change it from the owning side. Here, we must set `Book.publisher`.

In fact, it's often necessary to change *both sides* of a bidirectional association. For example, if the collection `Publisher.books` was stored in the second-level cache, we must also modify the collection, to ensure that the second-level cache remains synchronized with the database.

That said, it's *not* a hard requirement to update the unowned side, at least if you're sure you know what you're doing.



In principle Hibernate *does* allow you to have a unidirectional one-to-many, that is, a `@OneToMany` with no matching `@ManyToOne` on the other side. In practice, this mapping is unnatural, and just doesn't work very well. Avoid it.

Here we've used `Set` as the type of the collection, but Hibernate also allows the use of `List` or `Collection` here, with almost no difference in semantics. In particular, the `List` may not contain duplicate elements, and its order will not be persistent.

```
@OneToMany(mappedBy=Book_.PUBLISHER)
Collection<Book> books;
```

We'll see how to map a collection with a persistent order [much later](#).

Set, List, or Collection?

A one-to-many association mapped to a foreign key can never contain duplicate elements, so `Set` seems like the most semantically correct Java collection type to use here, and so that's the conventional practice in the Hibernate community.

The catch associated with using a set is that we must carefully ensure that `Book` has a high-quality implementation of `equals()` and `hashCode()`. Now, that's not necessarily a bad thing, since a quality `equals()` is independently useful.

But what if we used `Collection` or `List` instead? Then our code would be much less sensitive to how `equals()` and `hashCode()` were implemented.

In the past, we were perhaps too dogmatic in recommending the use of `Set`. Now? I guess we're happy to let you guys decide. In hindsight, we could have done more to make clear that this was always a viable option.

3.17. One-to-one (first way)

The simplest sort of one-to-one association is almost exactly like a `@ManyToOne` association, except that it maps to a foreign key column with a `UNIQUE` constraint.



Later, we'll see how to map a one-to-one association to an [association table](#).

A one-to-one association must be annotated `@OneToOne`:

```
@Entity
class Author {
    @Id @GeneratedValue
    Long id;

    @OneToOne(optional=false, fetch=LAZY)
    Person author;

    ...
}
```

Here, the `Author` table has a foreign key column holding the identifier of the associated `Publisher`.



A one-to-one association often models a "type of" relationship. In our example, an `Author` is a type of `Person`. An alternative—and often more natural—way to represent "type of" relationships in Java is via [entity class inheritance](#).

We can make this association bidirectional by adding a reference back to the `Author` in the `Person` entity:

```
@Entity
class Person {
    @Id @GeneratedValue
    Long id;

    @OneToOne(mappedBy = Author_.PERSON)
    Author author;

    ...
}
```

`Person.author` is the unowned side, because it's the side marked `mappedBy`.

Lazy fetching for one-to-one associations

Notice that we did not declare the unowned end of the association `fetch=LAZY`. That's because:

1. not every `Person` has an associated `Author`, and
2. the foreign key is held in the table mapped by `Author`, not in the table mapped by `Person`.

Therefore, Hibernate can't tell if the reference from `Person` to `Author` is `null` without fetching the associated `Author`.

On the other hand, if every `Person` was an `Author`, that is, if the association were non-optional, we would not have to consider the possibility of `null` references, and we would map it like this:

```
@OneToOne(optional=false, mappedBy = Author_.PERSON, fetch=LAZY)
Author author;
```

This is not the only sort of one-to-one association.

3.18. One-to-one (second way)

An arguably more elegant way to represent such a relationship is to share a primary key between the two tables.

To use this approach, the `Author` class must be annotated like this:

```
@Entity
class Author {
    @Id
    Long id;

    @OneToOne(optional=false, fetch=LAZY)
    @MapsId
    Person person;
```

```

    Person author;

    ...
}

```

Notice that, compared with the previous mapping:

- the `@Id` attribute is no longer a `@GeneratedValue` and,
- instead, the `author` association is annotated `@MapsId`.

This lets Hibernate know that the association to `Person` is the source of primary key values for `Author`.

Here, there's no extra foreign key column in the `Author` table, since the `id` column holds the identifier of `Person`. That is, the primary key of the `Author` table does double duty as the foreign key referring to the `Person` table.

The `Person` class doesn't change. If the association is bidirectional, we annotate the unowned side `@OneToOne(mappedBy = Author_.PERSON)` just as before.

3.19. Many-to-many

A unidirectional many-to-many association is represented as a collection-valued attribute. It always maps to a separate *association table* in the database.

It tends to happen that a many-to-many association eventually turns out to be an entity in disguise.



Suppose we start with a nice clean many-to-many association between `Author` and `Book`. Later on, it's quite likely that we'll discover some additional information which comes attached to the association, so that the association table needs some extra columns.

For example, imagine that we needed to report the percentage contribution of each author to a book. That information naturally belongs to the association table. We can't easily store it as an attribute of `Book`, nor as an attribute of `Author`.

When this happens, we need to change our Java model, usually introducing a new entity class which maps the association table directly. In our example, we might call this entity something like `BookAuthorship`, and it would have `@OneToMany` associations to both `Author` and `Book`, along with the `contribution` attribute.

We can evade the disruption occasioned by such "discoveries" by simply avoiding the use of `@ManyToMany` right from the start. There's little downside to representing every—or at least *almost* every—logical many-to-many association using an intermediate entity.

A many-to-many association must be annotated `@ManyToMany`:

```

@Entity
class Book {
    @Id @GeneratedValue
    Long id;

    @ManyToMany
    Set<Author> authors;

    ...
}

```

If the association is bidirectional, we add a very similar-looking attribute to `Book`, but this time we must specify `mappedBy` to indicate that this is the unowned side of the association:

```

@Entity
class Book {
    @Id @GeneratedValue
    Long id;

    @ManyToMany(mappedBy=Author_.BOOKS)
    Set<Author> authors;

    ...
}

```

Remember, if we wish to the modify the collection we must [change the owning side](#).

We've again used `Sets` to represent the association. As before, we have the option to use `Collection` or `List`. But in this case it *does* make a difference to the semantics of the association.



A many-to-many association represented as a `Collection` or `List` may contain duplicate elements. However, as before, the order of the elements is not persistent. That is, the collection is a *bag*, not a set.

3.20. Collections of basic values and embeddable objects

We've now seen the following kinds of entity attribute:

Kind of entity attribute	Kind of reference	Multiplicity	Examples
Single-valued attribute of basic type	Non-entity	At most one	@Basic String name
Single-valued attribute of embeddable type	Non-entity	At most one	@Embedded Name name
Single-valued association	Entity	At most one	@ManyToOne Publisher publisher @OneToOne Person person
Many-valued association	Entity	Zero or more	@OneToMany Set<Book> books @ManyToMany Set<Author> authors

Scanning this taxonomy, you might ask: does Hibernate have multivalued attributes of basic or embeddable type?

Well, actually, we've already seen that it does, at least in two special cases. So first, let's [recall](#) that JPA treats `byte[]` and `char[]` arrays as basic types. Hibernate persists a `byte[]` or `char[]` array to a `VARBINARY` or `VARCHAR` column, respectively.

But in this section we're really concerned with cases *other* than these two special cases. So then, *apart from* `byte[]` and `char[]`, does Hibernate have multivalued attributes of basic or embeddable type?

And the answer again is that *it does*. Indeed, there are two different ways to handle such a collection, by mapping it:

- to a column of SQL `ARRAY` type (assuming the database has an `ARRAY` type), or
- to a separate table.

So we may expand our taxonomy with:

Kind of entity attribute	Kind of reference	Multiplicity	Examples
<code>byte[]</code> and <code>char[]</code> arrays	Non-entity	Zero or more	<code>byte[]</code> image <code>char[]</code> text
Collection of basic-typed elements	Non-entity	Zero or more	@Array String[] names @ElementCollection Set<String> names
Collection of embeddable elements	Non-entity	Zero or more	@ElementCollection Set<Name> names

There's actually two new kinds of mapping here: `@Array` mappings, and `@ElementCollection` mappings.



These sorts of mappings are overused.

There *are* situations where we think it's appropriate to use a collection of basic-typed values in our entity class. But such situations are rare. Almost every many-valued relationship should map to a foreign key association between separate tables. And almost every table should be mapped by an entity class.

The features we're about to meet in the next two subsections are used much more often by beginners than they're used by experts. So if you're a beginner, you'll save yourself some hassle by staying away from these features for now.

We'll talk about `@Array` mappings first.

3.21. Collections mapped to SQL arrays

Let's consider a calendar event which repeats on certain days of the week. We might represent this in our `Event` entity as an attribute of type `DayOfWeek[]` or `List<DayOfWeek>`. Since the number of elements of this array or list is upper bounded by 7, this is a reasonable case for the use of an `ARRAY`-typed column. It's hard to see much value in storing this collection in a separate table.

Learning to not hate SQL arrays

For a long time, we thought arrays were a kind of weird and warty thing to add to the relational model, but recently we've come to realize that this view was overly closed-minded. Indeed, we might choose to view SQL `ARRAY` types as a generalization of `VARCHAR` and `VARBINARY` to generic "element" types. And from this point of view, SQL arrays look quite attractive, at least for certain problems. If we're comfortable mapping `byte[]` to `VARBINARY(255)`, why would we shy away from mapping `DayOfWeek[]` to `TINYINT ARRAY[7]`?

Unfortunately, JPA doesn't define a standard way to map SQL arrays, but here's how we can do it in Hibernate:

```
@Entity
class Event {
    @Id @GeneratedValue
    Long id;
    ...
    @Array(length=7)
    DayOfWeek[] daysOfWeek; // stored as a SQL ARRAY type
    ...
}
```

The `@Array` annotation is optional, but it's important to limit the amount of storage space the database allocates to the `ARRAY` column.



Now for the gotcha: not every database has a SQL `ARRAY` type, and some that *do* have an `ARRAY` type don't allow it to be used as a column type.

In particular, neither DB2 nor SQL Server have array-typed columns. On these databases, Hibernate falls back to something much worse: it uses Java serialization to encode the array to a binary representation, and stores the binary stream in a `VARBINARY` column. Quite clearly, this is terrible. You can ask Hibernate to do something *slightly* less terrible by annotating the attribute `@JdbcTypeCode(SqlTypes.JSON)`, so that the array is serialized to JSON instead of binary format. But at this point it's better to just admit defeat and use an `@ElementCollection` instead.

Alternatively, we could store this array or list in a separate table.

3.22. Collections mapped to a separate table

JPA *does* define a standard way to map a collection to an auxiliary table: the `@ElementCollection` annotation.

```
@Entity
class Event {
    @Id @GeneratedValue
    Long id;
    ...
    @ElementCollection
    DayOfWeek[] daysOfWeek; // stored in a dedicated table
    ...
}
```

Actually, we shouldn't use an array here, since array types can't be proxied, and so the JPA specification doesn't even say they're supported. Instead, we should use `Set`, `List`, or `Map`.

```
@Entity
class Event {
    @Id @GeneratedValue
    Long id;
    ...
    @ElementCollection
    List<DayOfWeek> daysOfWeek; // stored in a dedicated table
    ...
}
```


Here, each collection elements are stored as separate rows of the auxiliary table. By default, this table has the following definition:

```
create table Event_daysOfWeek (  
    Event_id bigint not null,  
    daysOfWeek tinyint check (daysOfWeek between 0 and 6),  
    daysOfWeek_ORDER integer not null,  
    primary key (Event1_id, daysOfWeek_ORDER)  
)
```

Which is fine, but it's still a mapping we prefer to avoid.



@ElementCollection is one of our least-favorite features of JPA. Even the name of the annotation is bad.

The code above results in a table with three columns:

- a foreign key of the Event table,
- a TINYINT encoding the enum, and
- an INTEGER encoding the ordering of elements in the array.

Instead of a surrogate primary key, it has a composite key comprising the foreign key of Event and the order column.

When—inevitably—we find that we need to add a fourth column to that table, our Java code must change completely. Most likely, we'll realize that we need to add a separate entity after all. So this mapping isn't very robust in the face of minor changes to our data model.

There's much more we could say about "element collections", but we won't say it, because we don't want to hand you the gun you'll shoot your foot with.

3.23. Summary of annotations

Let's pause to remember the annotations we've met so far.

Table 15. Declaring entities and embeddable types

Annotation	Purpose	JPA-standard
@Entity	Declare an entity class	✓
@MappedSuperclass	Declare a non-entity class with mapped attributes inherited by an entity	✓
@Embeddable	Declare an embeddable type	✓
@IdClass	declare the identifier class for an entity with multiple @Id attributes	✓

Table 16. Declaring basic attributes

Annotation	Purpose		JPA-standard
@Id	Declare a basic-typed identifier attribute		✓
@Version	Declare a version attribute		✓
@Basic	Declare a basic attribute	Default	✓
@EmbeddedId	Declare an embeddable-typed identifier attribute		✓
@Embedded	Declare an embeddable-typed attribute	Inferred	✓
@Enumerated	Declare an enum-typed attribute and specify how it is encoded	Inferred	✓
@Array	Declare that an attribute maps to a SQL ARRAY, and specify the length	Inferred	✗
@ElementCollection	Declare that a collection is mapped to a dedicated table		✓

Table 17. Converters and compositional basic types

Annotation	Purpose	JPA-standard
@Converter	Register an <code>AttributeConverter</code>	✓
@Convert	Apply a converter to an attribute	✓
@JavaType	Explicitly specify an implementation of <code>JavaType</code> for a basic attribute	✗
@JdbcType	Explicitly specify an implementation of <code>JdbcType</code> for a basic attribute	✗
@JdbcTypeCode	Explicitly specify a JDBC type code used to determine the <code>JdbcType</code> for a basic attribute	✗
@JavaTypeRegistration	Register a <code>JavaType</code> for a given Java type	✗
@JdbcTypeRegistration	Register a <code>JdbcType</code> for a given JDBC type code	✗

Table 18. System-generated identifiers

Annotation	Purpose	JPA-standard
@GeneratedValue	Specify that an identifier is system-generated	✓
@SequenceGenerator	Define an id generated backed by on a database sequence	✓
@TableGenerator	Define an id generated backed by a database table	✓
@IdGeneratorType	Declare an annotation that associates a custom <code>Generator</code> with each <code>@Id</code> attribute it annotates	✗
@ValueGenerationType	Declare an annotation that associates a custom <code>Generator</code> with each <code>@Basic</code> attribute it annotates	✗

Table 19. Declaring entity associations

Annotation	Purpose	JPA-standard
@ManyToOne	Declare the single-valued side of a many-to-one association (the owning side)	✓
@OneToMany	Declare the many-valued side of a many-to-one association (the unowned side)	✓
@ManyToMany	Declare either side of a one-to-one association	✓
@OneToOne	Declare either side of a one-to-one association	✓
@MapsId	Declare that the owning side of a <code>@OneToOne</code> association maps the primary key column	✓

Phew! That's already a lot of annotations, and we have not even started with the annotations for O/R mapping!

3.24. equals() and hashCode()

Entity classes should override `equals()` and `hashCode()`, especially when associations are represented as sets.

People new to Hibernate or JPA are often confused by exactly which fields should be included in the `hashCode()`. Even people with plenty of experience often argue quite religiously that one or another approach is the only right way. The truth is, there's no unique right way to do it, but there are some constraints. So please keep the following principles in mind:

- You should not include a mutable field in the hashcode, since that would require rehashing every collection containing the entity whenever the field is mutated.
- It's not completely wrong to include a generated identifier (surrogate key) in the hashcode, but since the identifier is not generated until the entity instance is made persistent, you must take great care to not add it to any hashed collection before the identifier is generated. We therefore advise against including any database-generated field in the hashcode.

It's OK to include any immutable, non-generated field in the hashcode.



We therefore recommend identifying a **natural key** for each entity, that is, a combination of fields that uniquely identifies an instance of the entity, from the perspective of the data model of the program. The natural key should correspond to a unique constraint on the database, and to the fields which are included in `equals()` and `hashCode()`.

```
@Entity
class Book {

    @Id @GeneratedValue
    Long id;

    @NaturalId
    @Basic(optional=false)
    String isbn;

    ...

    @Override
    public boolean equals(Object other) {
        return other instanceof Book
            && ((Book) other).isbn.equals(isbn);
    }
    @Override
    public int hashCode() {
        return isbn.hashCode();
    }
}
```

That said, an implementation of `equals()` and `hashCode()` based on the generated identifier of the entity can work *if you're careful*.

Chapter 4. Object/relational mapping

Given a domain model—that is, a collection of entity classes decorated with all the fancy annotations we [just met](#) in the previous chapter—Hibernate will happily go away and infer a complete relational schema, and even [export it to your database](#) if you ask politely.

The resulting schema will be entirely sane and reasonable, though if you look closely, you'll find some flaws. For example, every VARCHAR column will have the same length, VARCHAR(255).

But the process I just described—which we call *top down* mapping—simply doesn't fit the most common scenario for the use of O/R mapping. It's only rarely that the Java classes precede the relational schema. Usually, *we already have a relational schema*, and we're constructing our domain model around the schema. This is called *bottom up* mapping.



Developers often refer to a pre-existing relational database as "legacy" data. This tends to conjure images of bad old "legacy apps" written in COBOL or something. But legacy data is valuable, and learning to work with it is important.

Especially when mapping bottom up, we often need to customize the inferred object/relational mappings. This is a somewhat tedious topic, and so we don't want to spend too many words on it. Instead, we'll quickly skim the most important mapping annotations.

Hibernate SQL case convention

Computers have had lowercase letters for rather a long time now. Most developers learned long ago that text written in MixedCase, camelCase, or even snake_case is easier to read than text written in SHOUTYCASE. This is just as true of SQL as it is of any other language.

Therefore, for over twenty years, the convention on the Hibernate project has been that:

- query language identifiers are written in lowercase,
- table names are written in MixedCase, and
- column names are written in camelCase.

That is to say, we simply adopted Java's excellent conventions and applied them to SQL.

Now, there's no way we can force you to follow this convention, even if we wished to. Hell, you can easily write a `PhysicalNamingStrategy` which makes table and column names ALL UGLY AND SHOUTY LIKE THIS IF YOU PREFER. But, *by default*, it's the convention Hibernate follows, and it's frankly a pretty reasonable one.

4.1. Mapping entity inheritance hierarchies

In [Entity class inheritance](#) we saw that entity classes may exist within an inheritance hierarchy. There's three basic strategies for mapping an entity hierarchy to relational tables. Let's put them in a table, so we can more easily compare the points of difference between them.

Table 20. Entity inheritance mapping strategies

Strategy	Mapping	Polymorphic queries	Constraints	Normalization	When to use it
SINGLE_TABLE	Map every class in the hierarchy to the same table, and uses the value of a <i>discriminator column</i> to determine which concrete class each row represents.	To retrieve instances of a given class, we only need to query the one table.	Attributes declared by subclasses map to columns without NOT NULL constraints. 🤖 Any association may have a FOREIGN KEY constraint. 🤖	Subclass data is denormalized. 🤖	Works well when subclasses declare few or no additional attributes.

Strategy	Mapping	Polymorphic queries	Constraints	Normalization	When to use it
JOINED	Map every class in the hierarchy to a separate table, but each table only maps the attributes declared by the class itself. Optionally, a discriminator column may be used.	To retrieve instances of a given class, we must JOIN the table mapped by the class with: - all tables mapped by its superclasses and - all tables mapped by its subclasses.	Any attribute may map to a column with a NOT NULL constraint. 🤖 Any association may have a FOREIGN KEY constraint. 🤖	The tables are normalized. 🤖	The best option when we care a lot about constraints and normalization.
TABLE_PER_CLASS	Map every concrete class in the hierarchy to a separate table, but denormalize all inherited attributes into the table.	To retrieve instances of a given class, we must take a UNION over the table mapped by the class and the tables mapped by its subclasses.	Associations targeting a superclass cannot have a corresponding FOREIGN KEY constraint in the database. 🤖🤖 Any attribute may map to a column with a NOT NULL constraint. 🤖	Superclass data is denormalized. 🤖	Not very popular. From a certain point of view, competes with @MappedSuperclass.

The three mapping strategies are enumerated by `InheritanceType`. We specify an inheritance mapping strategy using the `@Inheritance` annotation.

For mappings with a *discriminator column*, we should:

- specify the discriminator column name and type by annotating the root entity `@DiscriminatorColumn`, and
- specify the values of this discriminator by annotating each entity in the hierarchy `@DiscriminatorValue`.

For single table inheritance we always need a discriminator:

```
@Entity
@DiscriminatorColumn(discriminatorType=CHAR, name="kind")
@DiscriminatorValue('P')
class Person { ... }

@Entity
@DiscriminatorValue('A')
class Author { ... }
```

We don't need to explicitly specify `@Inheritance(strategy=SINGLE_TABLE)`, since that's the default.

For JOINED inheritance we don't need a discriminator:

```
@Entity
@Inheritance(strategy=JOINED)
class Person { ... }

@Entity
class Author { ... }
```



However, we can add a discriminator column if we like, and in that case the generated SQL for polymorphic queries will be slightly simpler.

Similarly, for `TABLE_PER_CLASS` inheritance we have:

```
@Entity
@Inheritance(strategy=TABLE_PER_CLASS)
```

```
class Person { ... }
```

```
@Entity  
class Author { ... }
```



Hibernate doesn't allow discriminator columns for TABLE_PER_CLASS inheritance mappings, since they would make no sense, and offer no advantage.

Notice that in this last case, a polymorphic association like:

```
@ManyToOne Person person;
```

is a bad idea, since it's impossible to create a foreign key constraint that targets both mapped tables.

4.2. Mapping to tables

The following annotations specify exactly how elements of the domain model map to tables of the relational model:

Table 21. Annotations for mapping tables

Annotation	Purpose
@Table	Map an entity class to its primary table
@SecondaryTable	Define a secondary table for an entity class
@JoinTable	Map a many-to-many or many-to-one association to its association table
@CollectionTable	Map an @ElementCollection to its table

The first two annotations are used to map an entity to its *primary table* and, optionally, one or more *secondary tables*.

4.3. Mapping entities to tables

By default, an entity maps to a single table, which may be specified using @Table:

```
@Entity  
@Table(name="People")  
class Person { ... }
```

However, the @SecondaryTable annotation allows us to spread its attributes across multiple *secondary tables*.

```
@Entity  
@Table(name="Books")  
@SecondaryTable(name="Editions")  
class Book { ... }
```

The @Table annotation can do more than just specify a name:

Table 22. @Table annotation members

Annotation member	Purpose
name	The name of the mapped table
schema 🐼	The schema to which the table belongs
catalog 🐼	The catalog to which the table belongs
uniqueConstraints	One or more @UniqueConstraint annotations declaring multi-column unique constraints
indexes	One or more @Index annotations each declaring an index



It only makes sense to explicitly specify the schema in annotations if the domain model is spread across multiple schemas.

Otherwise, it's a bad idea to hardcode the schema (or catalog) in a `@Table` annotation. Instead:

- set the configuration property `hibernate.default_schema` (or `hibernate.default_catalog`), or
- simply specify the schema in the JDBC connection URL.

The `@SecondaryTable` annotation is even more interesting:

Table 23. `@SecondaryTable` annotation members

Annotation member	Purpose
<code>name</code>	The name of the mapped table
<code>schema</code> 🧠	The schema to which the table belongs
<code>catalog</code> 🧠	The catalog to which the table belongs
<code>uniqueConstraints</code>	One or more <code>@UniqueConstraint</code> annotations declaring multi-column unique constraints
<code>indexes</code>	One or more <code>@Index</code> annotations each declaring an index
<code>pkJoinColumns</code>	One or more <code>@PrimaryKeyJoinColumn</code> annotations, specifying primary key column mappings
<code>foreignKey</code>	An <code>@ForeignKey</code> annotation specifying the name of the FOREIGN KEY constraint on the <code>@PrimaryKeyJoinColumn</code> s



Using `@SecondaryTable` on a subclass in a `SINGLE_TABLE` entity inheritance hierarchy gives us a sort of mix of `SINGLE_TABLE` with `JOINED` inheritance.

4.4. Mapping associations to tables

The `@JoinTable` annotation specifies an *association table*, that is, a table holding foreign keys of both associated entities. This annotation is usually used with `@ManyToMany` associations:

```
@Entity
class Book {
    ...

    @ManyToMany
    @JoinTable(name="BooksAuthors")
    Set<Author> authors;

    ...
}
```

But it's even possible to use it to map a `@ManyToOne` or `@OneToOne` association to an association table.

```
@Entity
class Book {
    ...

    @ManyToOne(fetch=LAZY)
    @JoinTable(name="BookPublisher")
    Publisher publisher;

    ...
}
```

Here, there should be a `UNIQUE` constraint on one of the columns of the association table.

```
@Entity
class Author {
    ...
}
```

```

    @OneToOne(optional=false, fetch=LAZY)
    @JoinTable(name="AuthorPerson")
    Person author;
    ...
}

```

Here, there should be a `UNIQUE` constraint on *both* columns of the association table.

Table 24. `@JoinTable` annotation members

Annotation member	Purpose
<code>name</code>	The name of the mapped association table
<code>schema</code> 🐼	The schema to which the table belongs
<code>catalog</code> 🐼	The catalog to which the table belongs
<code>uniqueConstraints</code>	One or more <code>@UniqueConstraint</code> annotations declaring multi-column unique constraints
<code>indexes</code>	One or more <code>@Index</code> annotations each declaring an index
<code>joinColumns</code>	One or more <code>@JoinColumn</code> annotations, specifying foreign key column mappings to the table of the owning side
<code>inverseJoinColumns</code>	One or more <code>@JoinColumn</code> annotations, specifying foreign key column mappings to the table of the unowned side
<code>foreignKey</code>	An <code>@ForeignKey</code> annotation specifying the name of the <code>FOREIGN KEY</code> constraint on the <code>joinColumns</code>
<code>inverseForeignKey</code>	An <code>@ForeignKey</code> annotation specifying the name of the <code>FOREIGN KEY</code> constraint on the <code>inverseJoinColumns</code>

To better understand these annotations, we must first discuss column mappings in general.

4.5. Mapping to columns

These annotations specify how elements of the domain model map to columns of tables in the relational model:

Table 25. Annotations for mapping columns

Annotation	Purpose
<code>@Column</code>	Map an attribute to a column
<code>@JoinColumn</code>	Map an association to a foreign key column
<code>@PrimaryKeyJoinColumn</code>	Map the primary key used to join a secondary table with its primary, or a subclass table in <code>JOINED</code> inheritance with its root class table
<code>@OrderColumn</code>	Specifies a column that should be used to maintain the order of a <code>List</code> .
<code>@MapKeyColumn</code>	Specified a column that should be used to persist the keys of a <code>Map</code> .

We use the `@Column` annotation to map basic attributes.

4.6. Mapping basic attributes to columns

The `@Column` annotation is not only useful for specifying the column name.

Table 26. `@Column` annotation members

Annotation member	Purpose
<code>name</code>	The name of the mapped column
<code>table</code>	The name of the table to which this column belongs

Annotation member	Purpose
length	The length of a VARCHAR, CHAR, or VARBINARY column type
precision	The decimal digits of precision of a FLOAT, DECIMAL, NUMERIC, or TIME, or TIMESTAMP column type
scale	The scale of a DECIMAL or NUMERIC column type, the digits of precision that occur to the right of the decimal point
unique	Whether the column has a UNIQUE constraint
nullable	Whether the column has a NOT NULL constraint
insertable	Whether the column should appear in generated SQL INSERT statements
updatable	Whether the column should appear in generated SQL UPDATE statements
columnDefinition 	A DDL fragment that should be used to declare the column



We no longer recommend the use of `columnDefinition` since it results in unportable DDL. Hibernate has much better ways to customize the generated DDL using techniques that result in portable behavior across different databases.

Here we see four different ways to use the `@Column` annotation:

```

@Entity
@Table(name="Books")
@SecondaryTable(name="Editions")
class Book {
    @Id @GeneratedValue
    @Column(name="bookId") // customize column name
    Long id;

    @Column(length=100, nullable=false) // declare column as VARCHAR(100) NOT NULL
    String title;

    @Column(length=17, unique=true, nullable=false) // declare column as VARCHAR(17) NOT NULL UNIQUE
    String isbn;

    @Column(table="Editions", updatable=false) // column belongs to the secondary table, and is never updated
    int edition;
}

```

We don't use `@Column` to map associations.

4.7. Mapping associations to foreign key columns

The `@JoinColumn` annotation is used to customize a foreign key column.

Table 27. `@JoinColumn` annotation members

Annotation member	Purpose
name	The name of the mapped foreign key column
table	The name of the table to which this column belongs
referencedColumnName	The name of the column to which the mapped foreign key column refers
unique	Whether the column has a UNIQUE constraint
nullable	Whether the column has a NOT NULL constraint
insertable	Whether the column should appear in generated SQL INSERT statements
updatable	Whether the column should appear in generated SQL UPDATE statements

Annotation member	Purpose
columnDefinition 	A DDL fragment that should be used to declare the column
foreignKey	A @ForeignKey annotation specifying the name of the FOREIGN KEY constraint

A foreign key column doesn't necessarily have to refer to the primary key of the referenced table. It's quite acceptable for the foreign key to refer to any other unique key of the referenced entity, even to a unique key of a secondary table.

Here we see how to use @JoinColumn to define a @ManyToOne association mapping a foreign key column which refers to the @NaturalId of Book:

```
@Entity
@Table(name="Items")
class Item {
    ...

    @ManyToOne(optional=false) // implies nullable=false
    @JoinColumn(name = "bookIsbn", referencedColumnName = "isbn", // a reference to a non-PK column
                foreignKey = @ForeignKey(name="ItemsToBooksBySsn")) // supply a name for the FK constraint
    Book book;

    ...
}
```

In case this is confusing:

- bookIsbn is the name of the foreign key column in the Items table,
- it refers to a unique key isbn in the Books table, and
- it has a foreign key constraint named ItemsToBooksBySsn.

Note that the foreignkey member is completely optional and only affects DDL generation.



If you don't supply an explicit name using @ForeignKey, Hibernate will generate a quite ugly name. The reason for this is that the maximum length of foreign key names on some databases is extremely constrained, and we need to avoid collisions. To be fair, this is perfectly fine if you're only using the generated DDL for testing.

For composite foreign keys we might have multiple @JoinColumn annotations:

```
@Entity
@Table(name="Items")
class Item {
    ...

    @ManyToOne(optional=false)
    @JoinColumn(name = "bookIsbn", referencedColumnName = "isbn")
    @JoinColumn(name = "bookPrinting", referencedColumnName = "printing")
    Book book;

    ...
}
```

If we need to specify the @ForeignKey, this starts to get a bit messy:

```
@Entity
@Table(name="Items")
class Item {
    ...

    @ManyToOne(optional=false)
    @JoinColumns(value = {@JoinColumn(name = "bookIsbn", referencedColumnName = "isbn"),
                        @JoinColumn(name = "bookPrinting", referencedColumnName = "printing")},
                foreignKey = @ForeignKey(name="ItemsToBooksBySsn"))
    Book book;

    ...
}
```

For associations mapped to a `@JoinTable`, fetching the association requires two joins, and so we must declare the `@JoinColumns` inside the `@JoinTable` annotation:

```
@Entity
class Book {
    @Id @GeneratedValue
    Long id;

    @ManyToMany
    @JoinTable(joinColumns=@JoinColumn(name="bookId"),
              inverseJoinColumns=@JoinColumn(name="authorId"),
              foreignKey=@ForeignKey(name="BooksToAuthors"))
    Set<Author> authors;

    ...
}
```

Again, the `foreignKey` member is optional.

4.8. Mapping primary key joins between tables

The `@PrimaryKeyJoinColumn` is a special-purpose annotation for mapping:

- the primary key column of a `@SecondaryTable`—which is also a foreign key referencing the primary table, or
- the primary key column of the primary table mapped by a subclass in a `JOINED` inheritance hierarchy—which is also a foreign key referencing the primary table mapped by the root entity.

Table 28. `@PrimaryKeyJoinColumn` annotation members

Annotation member	Purpose
<code>name</code>	The name of the mapped foreign key column
<code>referencedColumnName</code>	The name of the column to which the mapped foreign key column refers
<code>columnDefinition</code> 🐼	A DDL fragment that should be used to declare the column
<code>foreignKey</code>	A <code>@ForeignKey</code> annotation specifying the name of the FOREIGN KEY constraint

When mapping a subclass table primary key, we place the `@PrimaryKeyJoinColumn` annotation on the entity class:

```
@Entity
@Table(name="People")
@Inheritance(strategy=JOINED)
class Person { ... }

@Entity
@Table(name="Authors")
@PrimaryKeyJoinColumn(name="personId") // the primary key of the Authors table
class Author { ... }
```

But to map a secondary table primary key, the `@PrimaryKeyJoinColumn` annotation must occur inside the `@SecondaryTable` annotation:

```
@Entity
@Table(name="Books")
@SecondaryTable(name="Editions",
               pkJoinColumns = @PrimaryKeyJoinColumn(name="bookId")) // the primary key of the Editions table
class Book {
    @Id @GeneratedValue
    @Column(name="bookId") // the name of the primary key of the Books table
    Long id;

    ...
}
```

4.9. Column lengths and adaptive column types

Hibernate automatically adjusts the column type used in generated DDL based on the column length specified by the `@Column` annotation. So we don't usually need to explicitly specify that a column should be of type `TEXT` or `CLOB`—or worry about the parade of `TINYTEXT`, `MEDIUMTEXT`, `TEXT`, `LONGTEXT` types on MySQL—because Hibernate will automatically select one of those types if required to accommodate a string of the length we specify.

The constant values defined in the class `org.hibernate.Length` are very helpful here:

Table 29. Predefined column lengths

Constant	Value	Description
DEFAULT	255	The default length of a <code>VARCHAR</code> or <code>VARBINARY</code> column when none is explicitly specified
LONG	32600	The largest column length for a <code>VARCHAR</code> or <code>VARBINARY</code> that is allowed on every database Hibernate supports
LONG16	32767	The maximum length that can be represented using 16 bits (but this length is too large for a <code>VARCHAR</code> or <code>VARBINARY</code> column on for some database)
LONG32	2147483647	The maximum length for a Java string

We can use these constants in the `@Column` annotation:

```
@Column(length=LONG)
String text;

@Column(length=LONG32)
byte[] binaryData;
```

This is usually all you need to do to make use of large object types in Hibernate.

4.10. LOBs

JPA provides a `@Lob` annotation which specifies that a field should be persisted as a `BLOB` or `CLOB`.

Semantics of the `@Lob` annotation

What the spec actually says is that the field should be persisted

...as a large object to a database-supported large object type.

It's quite unclear what this means, and the spec goes on to say that

...the treatment of the `Lob` annotation is provider-dependent...

which doesn't help much.

Hibernate interprets this annotation in what we think is the most reasonable way. In Hibernate, an attribute annotated `@Lob` will be written to JDBC using the `setClob()` or `setBlob()` method of `PreparedStatement`, and will be read from JDBC using the `getClob()` or `getBlob()` method of `ResultSet`.

Now, the use of these JDBC methods is usually unnecessary! JDBC drivers are perfectly capable of converting between `String` and `CLOB` or between `byte[]` and `BLOB`. So unless you specifically need to use these JDBC LOB APIs, you *don't* need the `@Lob` annotation.

Instead, as we just saw in [Column lengths and adaptive column types](#), all you need is to specify a large enough column length to accommodate the data you plan to write to that column.



Unfortunately, the driver for PostgreSQL doesn't allow `BYTEA` or `TEXT` columns to be read via the JDBC LOB APIs.

This limitation of the Postgres driver has resulted in a whole cottage industry of bloggers and stackoverflow question-answerers recommending convoluted ways to hack the Hibernate `Dialect` for Postgres to allow an attribute annotated `@Lob` to be written using `setString()` and read using `getString()`.

But simply removing the `@Lob` annotation has exactly the same effect.

Conclusion:

- on PostgreSQL, `@Lob` always means the `OID` type,
- `@Lob` should never be used to map columns of type `BYTEA` or `TEXT`, and
- please don't believe everything you read on stackoverflow.

Finally, as an alternative, Hibernate lets you declare an attribute of type `java.sql.Blob` or `java.sql.Clob`.

```
@Entity
class Book {
    ...
    Clob text;
    Blob coverArt;
    ...
}
```

The advantage is that a `java.sql.Clob` or `java.sql.Blob` can in principle index up to 2^{63} characters or bytes, much more data than you can fit in a Java `String` or `byte[]` array (or in your computer).

To assign a value to these fields, we'll need to use a `LobHelper`. We can get one from the `Session`:

```
LobHelper helper = session.getLobHelper();
book.text = helper.createClob(text);
book.coverArt = helper.createBlob(image);
```

In principle, the `Blob` and `Clob` objects provide efficient ways to read or stream LOB data from the server.

```
Book book = session.find(Book.class, bookId);
String text = book.text.getSubString(1, textLength);
InputStream bytes = book.images.getBinaryStream();
```

Of course, the behavior here depends very much on the JDBC driver, and so we really can't promise that this is a sensible thing to do on your database.

4.11. Mapping embeddable types to UDTs or to JSON

There's a couple of alternative ways to represent an embeddable type on the database side.

Embeddables as UDTs

First, a really nice option, at least in the case of Java record types, and for databases which support *user-defined types* (UDTs), is to define a UDT which represents the record type. Hibernate 6 makes this really easy. Just annotate the record type, or the attribute which holds a reference to it, with the new `@Struct` annotation:

```
@Embeddable
@Struct(name="PersonName")
record Name(String firstName, String middleName, String lastName) {}

@Entity
class Person {
    ...
    Name name;
    ...
}
```

This results in the following UDT:

```
create type PersonName as (firstName varchar(255), middleName varchar(255), lastName varchar(255))
```

And the name column of the Author table will have the type PersonName.

Embeddables to JSON

A second option that's available is to map the embeddable type to a JSON (or JSONB) column. Now, this isn't something we would exactly *recommend* if you're defining a data model from scratch, but it's at least useful for mapping pre-existing tables with JSON-typed columns. Since embeddable types are nestable, we can map some JSON formats this way, and even query JSON properties using HQL.



At this time, JSON arrays are not supported!

To map an attribute of embeddable type to JSON, we must annotate the attribute `@JdbcTypeCode(SqlTypes.JSON)`, instead of annotating the embeddable type. But the embeddable type `Name` should still be annotated `@Embeddable` if we want to query its attributes using HQL.

```
@Embeddable
record Name(String firstName, String middleName, String lastName) {}

@Entity
class Person {
    ...
    @JdbcTypeCode(SqlTypes.JSON)
    Name name;
    ...
}
```

We also need to add Jackson or an implementation of JSONB—for example, Yasson—to our runtime classpath. To use Jackson we could add this line to our Gradle build:

```
runtimeOnly 'com.fasterxml.jackson.core:jackson-databind:{jacksonVersion}'
```

Now the name column of the Author table will have the type jsonb, and Hibernate will automatically use Jackson to serialize a Name to and from JSON format.

4.12. Summary of SQL column type mappings

So, as we've seen, there are quite a few annotations that affect the mapping of Java types to SQL column types in DDL. Here we summarize the ones we've just seen in the second half of this chapter, along with some we already mentioned in earlier chapters.

Table 30. Annotations for mapping SQL column types

Annotation	Interpretation
@Enumerated	Specify how an enum type should be persisted
@Nationalized	Use a nationalized character type: NCHAR, NVARCHAR, or NCLOB
@Lob 🐼	Use JDBC LOB APIs to read and write the annotated attribute
@Array	Map a collection to a SQL ARRAY type of the specified length
@Struct	Map an embeddable to a SQL UDT with the given name
@TimeZoneStorage	Specify how the time zone information should be persisted
@JdbcType or @JdbcTypeCode	Use an implementation of JdbcType to map an arbitrary SQL type

In addition, there are some configuration properties which have a *global* affect on how basic types map to SQL column types:

Table 31. Type mapping settings

Configuration property name	Purpose
hibernate.use_nationalized_character_data	Enable use of nationalized character types by default
hibernate.type.preferred_boolean_jdbc_type	Specify the default SQL column type for mapping boolean

Configuration property name	Purpose
hibernate.type.preferred_uuid_jdbc_type	Specify the default SQL column type for mapping UUID
hibernate.type.preferred_duration_jdbc_type	Specify the default SQL column type for mapping Duration
hibernate.type.preferred_instant_jdbc_type	Specify the default SQL column type for mapping Instant
hibernate.timezone.default_storage	Specify the default strategy for storing time zone information



These are *global* settings and thus quite clumsy. We recommend against messing with any of these settings unless you have a really good reason for it.

There's one more topic we would like to cover in this chapter.

4.13. Mapping to formulas

Hibernate lets us map an attribute of an entity to a SQL formula involving columns of the mapped table. Thus, the attribute is a sort of "derived" value.

Table 32. Annotations for mapping formulas

Annotation	Purpose
@Formula	Map an attribute to a SQL formula
@JoinFormula	Map an association to a SQL formula
@DiscriminatorFormula	Use a SQL formula as the discriminator in single table inheritance .

For example:

```
@Entity
class Order {
    ...
    @Column(name = "sub_total", scale=2, precision=8)
    BigDecimal subTotal;

    @Column(name = "tax", scale=4, precision=4)
    BigDecimal taxRate;

    @Formula("sub_total * (1.0 + tax)")
    BigDecimal totalWithTax;
    ...
}
```

Chapter 5. Interacting with the database

To interact with the database, that is, to execute queries, or to insert, update, or delete data, we need an instance of one of the following objects:

- a JPA `EntityManager`,
- a Hibernate `Session`, or
- a Hibernate `StatelessSession`.

The `Session` interface extends `EntityManager`, and so the only difference between the two interfaces is that `Session` offers a few more operations.



Actually, in Hibernate, every `EntityManager` is a `Session`, and you can narrow it like this:

```
Session session = entityManager.unwrap(Session.class);
```

An instance of `Session` (or of `EntityManager`) is a *stateful session*. It mediates the interaction between your program and the database via a operations on a *persistence context*.

In this chapter, we're not going to talk much about `StatelessSession`. We'll come back to [this very useful API](#) when we talk about performance. What you need to know for now is that a stateless session doesn't have a persistence context.



Still, we should let you know that some people prefer to use `StatelessSession` everywhere. It's a simpler programming model, and lets the developer interact with the database more *directly*.


Stateful sessions certainly have their advantages, but they're more difficult to reason about, and when something goes wrong, the error messages can be more difficult to understand.

5.1. Persistence Contexts

A persistence context is a sort of cache; we sometimes call it the "first-level cache", to distinguish it from the [second-level cache](#). For every entity instance read from the database within the scope of a persistence context, and for every new entity made persistent within the scope of the persistence context, the context holds a unique mapping from the identifier of the entity instance to the instance itself.

Thus, an entity instance may be in one of three states with respect to a given persistence context:

1. *transient* — never persistent, and not associated with the persistence context,
2. *persistent* — currently associated with the persistence context, or
3. *detached* — previously persistent in another session, but not currently associated with *this* persistence context.



At any given moment, an instance may be associated with at most one persistence context.

The lifetime of a persistence context usually corresponds to the lifetime of a transaction, though it's possible to have a persistence context that spans several database-level transactions that form a single logical unit of work.



A persistence context—that is, a `Session` or `EntityManager`—absolutely positively **must not be shared between multiple threads or between concurrent transactions**.

If you accidentally leak a session across threads, you will suffer.

Container-managed persistence contexts

In a container environment, the lifecycle of a persistence context scoped to the transaction will usually be managed for you.

There are several reasons we like persistence contexts.

1. They help avoid *data aliasing*: if we modify an entity in one section of code, then other code executing within the same persistence context will see our modification.
2. They enable *automatic dirty checking*: after modifying an entity, we don't need to perform any explicit operation to ask Hibernate to propagate that change back to the database. Instead, the change will be automatically synchronized with the database when the session is **flushed**.
3. They can improve performance by avoiding a trip to the database when a given entity instance is requested repeatedly in a given unit of work.
4. They make it possible to *transparently batch* together multiple database operations.

A persistence context also allows us to detect circularities when performing operations on graphs of entities. (Even in a stateless session, we need some sort of temporary cache of the entity instances we've visited while executing a query.)

On the other hand, stateful sessions come with some very important restrictions, since:

- persistence contexts aren't threadsafe, and can't be shared across threads, and
- a persistence context can't be reused across unrelated transactions, since that would break the isolation and atomicity of the transactions.

Furthermore, a persistence context holds a hard references to all its entities, preventing them from being garbage collected. Thus, the session must be discarded once a unit of work is complete.



If you don't completely understand the previous passage, go back and re-read it until you do. A great deal of human suffering has resulted from users mismanaging the lifecycle of the `Hibernate Session` or `JPA EntityManager`.

We'll conclude by noting that whether a persistence context helps or harms the performance of a given unit of work depends greatly on the nature of the unit of work. For this reason Hibernate provides both stateful and stateless sessions.

5.2. Creating a session

Sticking with standard JPA-defined APIs, we saw how to obtain an `EntityManagerFactory` in [Configuration using JPA XML](#). It's quite unsurprising that we may use this object to create an `EntityManager`:

```
EntityManager entityManager = entityManagerFactory.createEntityManager();
```

When we're finished with the `EntityManager`, we should explicitly clean it up:

```
entityManager.close();
```

On the other hand, if we're starting from a `SessionFactory`, as described in [Configuration using Hibernate API](#), we may use:

```
Session session = sessionFactory.openSession();
```

But we still need to clean up:

```
session.close();
```

Injecting the `EntityManager`

If you're writing code for some sort of container environment, you'll probably obtain the `EntityManager` by some sort of dependency injection. For example, in Java (or Jakarta) EE you would write:

```
@PersistenceContext EntityManager entityManager;
```

In Quarkus, injection is handled by CDI:

```
@Inject EntityManager entityManager;
```

Outside a container environment, we'll also have to write code to manage database transactions.

5.3. Managing transactions

Using JPA-standard APIs, the `EntityManager` interface allows us to control database transactions. The idiom we recommend is the following:

```
EntityManager entityManager = entityManagerFactory.createEntityManager();
EntityTransaction tx = entityManager.getTransaction();
try {
    tx.begin();
    //do some work
    ...
    tx.commit();
}
catch (Exception e) {
    if (tx.isActive()) tx.rollback();
    throw e;
}
finally {
    entityManager.close();
}
```

Using Hibernate's native APIs we might write something really similar, but since this sort of code is extremely tedious, we have a much nicer option:

```
sessionFactory.inTransaction(session -> {
    //do the work
    ...
});
```

Container-managed transactions

In a container environment, the container itself is usually responsible for managing transactions. In Java EE or Quarkus, you'll probably indicate the boundaries of the transaction using the `@Transactional` annotation.

5.4. Operations on the persistence context

Of course, the main reason we need an `EntityManager` is to do stuff to the database. The following important operations let us interact with the persistence context and schedule modifications to the data:

Table 33. Methods for modifying data and managing the persistence context

Method name and parameters	Effect
<code>persist(Object)</code>	Make a transient object persistent and schedule a SQL <code>insert</code> statement for later execution
<code>remove(Object)</code>	Make a persistent object transient and schedule a SQL <code>delete</code> statement for later execution
<code>merge(Object)</code>	Copy the state of a given detached object to a corresponding managed persistent instance and return the persistent object
<code>detach(Object)</code>	Disassociate a persistent object from a session without affecting the database
<code>clear()</code>	Empty the persistence context and detach all its entities
<code>flush()</code>	Detect changes made to persistent objects association with the session and synchronize the database state with the state of the session by executing SQL <code>insert</code> , <code>update</code> , and <code>delete</code> statements

Notice that `persist()` and `remove()` have no immediate effect on the database, and instead simply schedule a command for later execution. Also notice that there's no `update()` operation for a stateful session. Modifications are automatically detected when the session is [flushed](#).

On the other hand, the following operations all result in immediate access to the database:

Table 34. Methods for reading and locking data

Method name and parameters	Effect
<code>find(Class, Object)</code>	Obtain a persistent object given its type and its id
<code>find(Class, Object, LockModeType)</code>	Obtain a persistent object given its type and its id, requesting the given optimistic or pessimistic lock mode
<code>getReference(Class, id)</code>	Obtain a reference to a persistent object given its type and its id, without actually loading its state from the database
<code>getReference(Object)</code>	Obtain a reference to a persistent object with the same identity as the given detached instance, without actually loading its state from the database
<code>refresh(Object)</code>	Refresh the persistent state of an object using a new SQL <code>select</code> to retrieve its current state from the database
<code>refresh(Object, LockModeType)</code>	Refresh the persistent state of an object using a new SQL <code>select</code> to retrieve its current state from the database, requesting the given optimistic or pessimistic lock mode
<code>lock(Object, LockModeType)</code>	Obtain an optimistic or pessimistic lock on a persistent object

Any of these operations might throw an exception. Now, if an exception occurs while interacting with the database, there's no good way to resynchronize the state of the current persistence context with the state held in database tables.

Therefore, a session is considered to be unusable after any of its methods throws an exception.



The persistence context is fragile. If you receive an exception from Hibernate, you should immediately close and discard the current session. Open a new session if you need to, but throw the bad one away first.

Each of the operations we've seen so far affects a single entity instance passed as an argument. But there's a way to set things up so that an operation will propagate to associated entities.

5.5. Cascading persistence operations

It's quite often the case that the lifecycle of a *child* entity is completely dependent on the lifecycle of some *parent*. This is especially common for many-to-one and one-to-one associations, though it's very rare for many-to-many associations.

For example, it's quite common to make an *Order* and all its *Items* persistent in the same transaction, or to delete a *Project* and its *Files* at once. This sort of relationship is sometimes called a *whole/part*-type relationship.

Cascading is a convenience which allows us to propagate one of the operations listed in [Operations on the persistence context](#) from a parent to its children. To set up cascading, we specify the `cascade` member of one of the association mapping annotations, usually `@OneToMany` or `@OneToOne`.

```
@Entity
class Order {
    ...
    @OneToMany(mappedBy=Item_.ORDER,
        // cascade persist(), remove(), and refresh() from Order to Item
        cascade={PERSIST, REMOVE, REFRESH},
        // also remove() orphaned Items
        orphanRemoval=true)
    private Set<Item> items;
    ...
}
```

Orphan removal indicates that an *Item* should be automatically deleted if it is removed from the set of items belonging to its parent *Order*.

5.6. Proxies and lazy fetching

Our data model is a set of interconnected entities, and in Java our whole dataset would be represented as an enormous interconnected graph of objects. It's possible that this graph is disconnected, but more likely it's connected, or composed of a relatively small number of connected subgraphs.

Therefore, when we retrieve an object belonging to this graph from the database and instantiate it in memory, we simply can't

recursively retrieve and instantiate all its associated entities. Quite aside from the waste of memory on the VM side, this process would involve a huge number of round trips to the database server, or a massive multidimensional cartesian product of tables, or both. Instead, we're forced to cut the graph somewhere.

Hibernate solves this problem using *proxies* and *lazy fetching*. A proxy is an object that masquerades as a real entity or collection, but doesn't actually hold any state, because that state has not yet been fetched from the database. When you call a method of the proxy, Hibernate will detect the call and fetch the state from the database before allowing the invocation to proceed to the real entity object or collection.

Now for the gotchas:

1. Hibernate will only do this for an entity which is currently associated with a persistence context. Once the session ends, and the persistence context is cleaned up, the proxy is no longer fetchable, and instead its methods throw the hated `LazyInitializationException`.
2. A round trip to the database to fetch the state of a single entity instance is just about *the least efficient* way to access data. It almost inevitably leads to the infamous *N+1 selects* problem we'll discuss later when we talk about how to [optimize association fetching](#).



We're getting a bit ahead of ourselves here, but let's quickly mention the general strategy we recommend to navigate past these gotchas:

- All associations should be set `fetch=LAZY` to avoid fetching extra data when it's not needed. As we mentioned [earlier](#), this setting is not the default for `@ManyToOne` associations, and must be specified explicitly.
- But strive to avoid writing code which triggers lazy fetching. Instead, fetch all the data you'll need upfront at the beginning of a unit of work, using one of the techniques described in [Association fetching](#), usually, using `join fetch` in HQL or an `EntityGraph`.

It's important to know that some operations which may be performed with an unfetched proxy *don't* require fetching its state from the database. First, we're always allowed to obtain its identifier:

```
var pubId = entityManager.find(Book.class, bookId).getPublisher().getId(); // does not fetch publisher
```

Second, we may create an association to a proxy:

```
book.setPublisher(entityManager.getReference(Publisher.class, pubId)); // does not fetch publisher
```

Sometimes it's useful to test whether a proxy or collection has been fetched from the database. JPA lets us do this using the `PersistenceUnitUtil`:

```
boolean authorsFetched = entityManagerFactory.getPersistenceUnitUtil().isLoaded(book.getAuthors());
```

Hibernate has a slightly easier way to do it:

```
boolean authorsFetched = Hibernate.isInitialized(book.getAuthors());
```

But the static methods of the `Hibernate` class let us do a lot more, and it's worth getting a bit familiar them.

Of particular interest are the operations which let us work with unfetched collections without fetching their state from the database. For example, consider this code:

```
Book book = session.find(Book.class, bookId); // fetch just the Book, leaving authors unfetched
Author authorRef = session.getReference(Author.class, authorId); // obtain an unfetched proxy
boolean isByAuthor = Hibernate.contains(book.getAuthors(), authorRef); // no fetching
```

This code fragment leaves both the set `book.authors` and the proxy `authorRef` unfetched.

Finally, `Hibernate.initialize()` is a convenience method that force-fetches a proxy or collection:

```
Book book = session.find(Book.class, bookId); // fetch just the Book, leaving authors unfetched
Hibernate.initialize(book.getAuthors()); // fetch the Authors
```

But of course, this code is very inefficient, requiring two trips to the database to obtain data that could in principle be retrieved with just one query.

It's clear from the discussion above that we need a way to request that an association be *eagerly* fetched using a database join, thus protecting ourselves from the infamous *N+1 selects*. One way to do that is by passing an `EntityGraph` to `find()`.

5.7. Entity graphs and eager fetching

When an association is mapped `fetch=LAZY`, it won't, by default, be fetched when we call the `find()` method. We may request that an association be fetched eagerly (immediately) by passing an `EntityGraph` to `find()`.

The JPA-standard API for this is a bit unwieldy:

```
var graph = entityManager.createEntityGraph(Book.class);
graph.addSubgraph(Book_.publisher);
Book book = entityManager.find(Book.class, bookId, Map.of(SpecHints.HINT_SPEC_FETCH_GRAPH, graph));
```

This is untypesafe and unnecessarily verbose. Hibernate has a better way:

```
var graph = session.createEntityGraph(Book.class);
graph.addSubgraph(Book_.publisher);
Book book = session.byId(Book.class).withFetchGraph(graph).load(bookId);
```

This code adds a `left outer join` to our SQL query, fetching the associated `Publisher` along with the `Book`.

We may even attach additional nodes to our `EntityGraph`:

```
var graph = session.createEntityGraph(Book.class);
graph.addSubgraph(Book_.publisher);
graph.addPluralSubgraph(Book_.authors).addSubgraph(Author_.person);
Book book = session.byId(Book.class).withFetchGraph(graph).load(bookId);
```

This results in a SQL query with *four* `left outer joins`.



In the code examples above, The classes `Book_` and `Author_` are generated by the [JPA Metamodel Generator](#) we saw earlier. They let us refer to attributes of our model in a completely type-safe way. We'll use them again, below, when we talk about [Criteria queries](#).

JPA specifies that any given `EntityGraph` may be interpreted in two different ways.

- A *fetch graph* specifies exactly the associations that should be eagerly loaded. Any association not belonging to the entity graph is proxied and loaded lazily only if required.
- A *load graph* specifies that the associations in the entity graph are to be fetched in addition to the associations mapped `fetch=EAGER`.

You're right, the names make no sense. But don't worry, if you take our advice, and map your associations `fetch=LAZY`, there's no difference between a "fetch" graph and a "load" graph, so the names don't matter.



JPA even specifies a way to define named entity graphs using annotations. But the annotation-based API is so verbose that it's just not worth using.

5.8. Flushing the session

From time to time, a *flush* operation is triggered, and the session synchronizes dirty state held in memory—that is, modifications to the state of entities associated with the persistence context—with persistent state held in the database. Of course, it does this by executing SQL `INSERT`, `UPDATE`, and `DELETE` statements.

By default, a flush is triggered when:

- the current transaction commits, for example, when `Transaction.commit()` is called,
- before execution of a query whose result would be affected by the synchronization of dirty state held in memory, or
- when the program directly calls `flush()`.



Notice that SQL statements are not usually executed synchronously by methods of the `Session` interface like `persist()` and `remove()`. If synchronous execution of SQL is desired, the `StatelessSession` allows this.

This behavior can be controlled by explicitly setting the flush mode. For example, to disable flushes that occur before query execution, call:

```
entityManager.setFlushMode(FlushModeType.COMMIT);
```

Hibernate allows greater control over the flush mode than JPA:

```
session.setHibernateFlushMode(FlushMode.MANUAL);
```

Since flushing is a somewhat expensive operation (the session must dirty-check every entity in the persistence context), setting the flush mode to `COMMIT` can occasionally be a useful optimization.

Table 35. Flush modes

Hibernate FlushMode	JPA FlushModeType	Interpretation
MANUAL		Never flush automatically
COMMIT	COMMIT	Flush before transaction commit
AUTO	AUTO	Flush before transaction commit, and before execution of a query whose results might be affected by modifications held in memory
ALWAYS		Flush before transaction commit, and before execution of every query

A second way to reduce the cost of flushing is to load entities in *read-only* mode:

- `Session.setDefaultReadOnly(false)` specifies that all entities loaded by a given session should be loaded in read-only mode by default,
- `SelectionQuery.setReadOnly(false)` specifies that every entity returned by a given query should be loaded in read-only mode, and
- `Session.setReadOnly(Object, false)` specifies that a given entity already loaded by the session should be switched to read-only mode.

It's not necessary to dirty-check on entity instance in read-only mode.

5.9. Queries

Hibernate features three complementary ways to write queries:

- the *Hibernate Query Language*, an extremely powerful superset of JPQL, which abstracts most of the features of modern dialects of SQL,
- the JPA *criteria query* API, along with extensions, allowing almost any HQL query to be constructed programmatically via a typesafe API, and, of course
- for when all else fails, *native SQL* queries.

5.10. HQL queries

A full discussion of the query language would require just as much text as the rest of this Introduction. Fortunately, HQL is already described in exhaustive (and exhausting) detail in *A guide to Hibernate Query Language 6*. It doesn't make sense to repeat that information here.

Here we want to see how to execute a query via the `Session` or `EntityManager` API. The method we call depends on what kind of query it is:

- *selection queries* return a result list, but do not modify the data, but
- *mutation queries* modify data, and return the number of modified rows.

Selection queries usually start with the keyword `select` or `from`, whereas mutation queries begin with the keyword `insert`, `update`, or `delete`.

Table 36. Executing HQL

Kind	Session method	EntityManager method	Query execution method
Selection	<code>createSelectionQuery(String, Class)</code>	<code>createQuery(String, Class)</code>	<code>getResultList()</code> , <code>getSingleResult()</code> , or <code>getSingleResultOrNull()</code>
Mutation	<code>createMutationQuery(String)</code>	<code>createQuery(String)</code>	<code>executeUpdate()</code>

So for the `Session` API we would write:

```
List<Book> matchingBooks =
    session.createSelectionQuery("from Book where title like :titleSearchPattern", Book.class)
        .setParameter("titleSearchPattern", titleSearchPattern)
        .getResultList();
```

Or, if we're sticking to the JPA-standard APIs:

```
List<Book> matchingBooks =
    entityManager.createQuery("select b from Book b where b.title like :titleSearchPattern", Book.class)
        .setParameter("titleSearchPattern", titleSearchPattern)
        .getResultList();
```

The only difference between `createSelectionQuery()` and `createQuery()` is that `createSelectionQuery()` throw an exception if passed an insert, delete, or update.

In the query above, `:titleSearchPattern` is called a *named parameter*. We may also identify parameters by a number. These are called *ordinal parameters*.

```
List<Book> matchingBooks =
    session.createSelectionQuery("from Book where title like ?1", Book.class)
        .setParameter(1, titleSearchPattern)
        .getResultList();
```

When a query has multiple parameters, named parameters tend to be easier to read, even if slightly more verbose.



Never concatenate user input with HQL and pass the concatenated string to `createSelectionQuery()`. This would open up the possibility for an attacker to execute arbitrary code on your database server.

If we're expecting a query to return a single result, we can use `getSingleResult()`.

```
Book book =
    session.createSelectionQuery("from Book where isbn = ?1", Book.class)
        .setParameter(1, isbn)
        .getSingleResult();
```

Or, if we're expecting it to return at most one result, we can use `getSingleResultOrNull()`.

```
Book bookOrNull =
    session.createSelectionQuery("from Book where isbn = ?1", Book.class)
        .setParameter(1, isbn)
        .getSingleResultOrNull();
```

The difference, of course, is that `getSingleResult()` throws an exception if there's no matching row in the database, whereas `getSingleResultOrNull()` just returns `null`.

By default, Hibernate dirty checks entities in the persistence context before executing a query, in order to determine if the session should be flushed. If there are many entities association with the persistence context, then this can be an expensive operation.

To disable this behavior, set the flush mode to `COMMIT` or `MANUAL`:

```
Book bookOrNull =
    session.createSelectionQuery("from Book where isbn = ?1", Book.class)
        .setParameter(1, isbn)
        .setHibernateFlushMode(MANUAL)
        .getSingleResult();
```



Setting the flush mode to `COMMIT` or `MANUAL` might cause the query to return stale results.

Occasionally we need to build a query at runtime, from a set of optional conditions. For this, JPA offers an API which allows programmatic construction of a query.

5.11. Criteria queries

Imagine we're implementing some sort of search screen, where the user of our system is offered several different ways to constrain the query result set. For example, we might let them search for books by title and/or the author name. Of course, we could construct a HQL query by string concatenation, but this is a bit fragile, so it's quite nice to have an alternative.

HQL is implemented in terms of criteria objects

Actually, in Hibernate 6, every HQL query is compiled to a criteria query before being translated to SQL. This ensures that the semantics of HQL and criteria queries are identical.

First we need an object for building criteria queries. Using the JPA-standard APIs, this would be a `CriteriaBuilder`, and we get it from the `EntityManagerFactory`:

```
CriteriaBuilder builder = entityManagerFactory.getCriteriaBuilder();
```

But if we have a `SessionFactory`, we get something much better, a `HibernateCriteriaBuilder`:

```
HibernateCriteriaBuilder builder = sessionFactory.getCriteriaBuilder();
```

The `HibernateCriteriaBuilder` extends `CriteriaBuilder` and adds many operations that JPQL doesn't have.



If you're using `EntityManagerFactory`, don't despair, you have two perfectly good ways to obtain the `HibernateCriteriaBuilder` associated with that factory. Either:

```
HibernateCriteriaBuilder builder =  
    entityManagerFactory.unwrap(SessionFactory.class).getCriteriaBuilder();
```

Or simply:

```
HibernateCriteriaBuilder builder =  
    (HibernateCriteriaBuilder) entityManagerFactory.getCriteriaBuilder();
```

We're ready to create a criteria query.

```
CriteriaQuery<Book> query = builder.createQuery(Book.class);  
Root<Book> book = query.from(Book.class);  
Predicate where = builder.conjunction();  
if (titlePattern != null) {  
    where = builder.and(where, builder.like(book.get(Book_.title), titlePattern));  
}  
if (namePattern != null) {  
    Join<Book, Author> author = book.join(Book_.author);  
    where = builder.and(where, builder.like(author.get(Author_.name), namePattern));  
}  
query.select(book).where(where)  
    .orderBy(builder.asc(book.get(Book_.title)));
```

Here, as before, the classes `Book_` and `Author_` are generated by Hibernate's [JPA Metamodel Generator](#).



Notice that we did not bother treating `titlePattern` and `namePattern` as parameters. That's safe because, by default, Hibernate automatically and transparently treats strings passed to the `CriteriaBuilder` as JDBC parameters.

Execution of a criteria query works almost exactly like execution of HQL.

Table 37. Executing criteria queries

Kind	Session method	EntityManager method	Query execution method
Selection	<code>createSelectionQuery(CriteriaQuery)</code>	<code>createQuery(CriteriaQuery)</code>	<code>getResultList()</code> , <code>getSingleResult()</code> , or <code>getSingleResultOrNull()</code>
Mutation	<code>createMutationQuery(CriteriaUpdate)</code> or <code>createQuery(CriteriaDelete)</code>	<code>createQuery(CriteriaUpdate)</code> or <code>createQuery(CriteriaDelete)</code>	<code>executeUpdate()</code>

For example:

```
List<Book> matchingBooks =
```



```
session.createQuery(query)
    .getResultList();
```

Update, insert, and delete queries work similarly:

```
CriteriaDelete<Book> delete = builder.createCriteriaDelete(Book.class);
Root<Book> book = delete.from(Book.class);
delete.where(builder.lt(builder.year(book.get(Book_.publicationDate)), 2000));
session.createMutationQuery(delete).executeUpdate();
```

When all else fails, and sometimes even before that, we're left with the option of writing a query in SQL.

5.12. Native SQL queries

HQL is a powerful language which helps reduce the verbosity of SQL, and significantly increases portability of queries between databases. But ultimately, the true value of ORM is not in avoiding SQL, but in alleviating the pain involved in dealing with SQL result sets once we get them back to our Java program. As we said [right up front](#), Hibernate's generated SQL is meant to be used in conjunction with handwritten SQL, and native SQL queries are one of the facilities we provide to make that easy.

Table 38. Executing SQL

Kind	Session method	EntityManager method	Query execution method
Selection	<code>createNativeQuery(String,Class)</code>	<code>createNativeQuery(String,Class)</code>	<code>getResultList()</code> , <code>getSingleResult()</code> , or <code>getSingleResultOrNull()</code>
Mutation	<code>createNativeMutationQuery(String)</code>	<code>createNativeQuery(String)</code>	<code>executeUpdate()</code>
Stored procedure	<code>createStoredProcedureCall(String)</code>	<code>createStoredProcedureQuery(String)</code>	<code>execute()</code>

For the most simple cases, Hibernate can infer the shape of the result set:

```
Book book =
    session.createNativeQuery("select * from Books where isbn = ?1", Book.class)
        .getSingleResult();

String title =
    session.createNativeQuery("select title from Books where isbn = ?1", String.class)
        .getSingleResult();
```

However, in general, there isn't enough information in the JDBC `ResultSetMetaData` to infer the mapping of columns to entity objects. So for more complicated cases, you'll need to use the `@SqlResultSetMapping` annotation to define a named mapping, and pass the name to `createNativeQuery()`. This gets fairly messy, so we don't want to hurt your eyes by showing you an example of it.

By default, Hibernate doesn't flush the session before execution of a native query. That's because the session is unaware of which modifications held in memory would affect the results of the query.

So if there are any unflushed changes to Books, this query might return stale data:

```
List<Book> books =
    session.createNativeQuery("select * from Books")
        .getResultList()
```

There's two ways to ensure the persistence context is flushed before this query is executed.

Either, we could simply force a flush by set the flush mode to ALWAYS:

```
List<Book> books =
    session.createNativeQuery("select * from Books")
        .setHibernateFlushMode(ALWAYS)
        .getResultList()
```

Or, alternatively, we could tell Hibernate which modified state affects the results of the query:

```
List<Book> books =
    session.createNativeQuery("select * from Books")
```

```
.addSynchronizedEntityClass(Book.class)
.getResultList()
```



You can call stored procedures using `createStoredProcedureQuery()` or `createStoredProcedureCall()`.

5.13. Limits, pagination, and ordering

If a query might return more results than we can handle at one time, we may specify:

- a *limit* on the maximum number of rows returned, and,
- optionally, an *offset*, the first row of an ordered result set to return.



The offset is used to paginate query results.

There's two ways to add a limit or offset to a HQL or native SQL query:

- using the syntax of the query language itself, for example, `offset 10 rows fetch next 20 rows only`, or
- using the methods `setFirstResult()` and `setMaxResults()` of the `SelectionQuery` interface.

If the limit or offset is parameterized, the second option is simpler. For example, this:

```
List<Book> books =
    session.createQuery("from Book where title like ?1 order by title")
        .setParameter(1, titlePattern)
        .setMaxResults(MAX_RESULTS)
        .getResultList();
```

is simpler than:

```
List<Book> books =
    session.createQuery("from Book where title like ?1 order by title fetch first ?2 rows only")
        .setParameter(1, titlePattern)
        .setParameter(2, MAX_RESULTS)
        .getResultList();
```

Hibernate's `SelectionQuery` has a slightly different way to paginate the query results:

```
List<Book> books =
    session.createQuery("from Book where title like ?1 order by title")
        .setParameter(1, titlePattern)
        .setMaxResults(MAX_RESULTS)
        .getResultList();
```

A closely-related issue is ordering, for which Hibernate ORM 6.3 offers an incubating API.

Unfortunately, there's no way to do this using JPA's `TypedQuery` interface.

Table 39. Methods for query limits, pagination, and ordering

Method name	Purpose	JPA-standard
<code>setMaxResults()</code>	Set a limit on the number of results returned by a query	✓
<code>setFirstResult()</code>	Set an offset on the results returned by a query	✓

5.14. Representing projection lists

A *projection list* is the list of things that a query returns, that is, the list of expressions in the `select` clause. Since Java has no tuple types, representing query projection lists in Java has always been a problem for JPA and Hibernate. Traditionally, we've just used `Object[]` most of the time:

```
var results =
    session.createQuery("select isbn, title from Book", Object[].class)
        .getResultList();

for (var result : results) {
```

```

    var isbn = (String) result[0];
    var title = (String) result[1];
    ...
}

```

This is really a bit ugly. Java's record types now offer an interesting alternative:

```

record ISBNTitle(String isbn, String title) {}

var results =
    session.createQuery("select isbn, title from Book", ISBNTitle.class)
        .getResultList();

for (var result : results) {
    var isbn = result.isbn();
    var title = result.title();
    ...
}

```

Notice that we're able to declare the record right before the line which executes the query.

Now, this is only *superficially* more typesafe, since the query itself is not checked statically, and so we can't say it's objectively better. But perhaps you find it more aesthetically pleasing. And if we're going to be passing query results around the system, the use of a record type is *much* better.

The Criteria query API offers a much more satisfying solution to the problem. Consider the following code:

```

var builder = sessionFactory.getCriteriaBuilder();
var query = builder.createTupleQuery();
var book = query.from(Book.class);
var bookTitle = book.get(Book_.title);
var bookIsbn = book.get(Book_.isbn);
var bookPrice = book.get(Book_.price);
query.select(builder.tuple(bookTitle, bookIsbn, bookPrice));
var resultList = session.createQuery(query).getResultList();
for (var result: resultList) {
    String title = result.get(bookTitle);
    String isbn = result.get(bookIsbn);
    BigDecimal price = result.get(bookPrice);
    ...
}

```

This code is manifestly completely typesafe, and much better than we can hope to do with HQL.

5.15. Named queries

The `@NamedQuery` annotation lets us define a HQL query that is compiled and checked as part of the bootstrap process. This means we find out about errors in our queries earlier, instead of waiting until the query is actually executed. We can place the `@NamedQuery` annotation on any class, even on an entity class.

```

@NamedQuery(name="10BooksByTitle",
    query="from Book where title like :titlePattern order by title fetch first 10 rows only")
class BookQueries {}

```

We have to make sure that the class with the `@NamedQuery` annotation will be scanned by Hibernate, either:

- by adding `<class>org.hibernate.example.BookQueries</class>` to `persistence.xml`, or
- by calling `configuration.addClass(BookQueries.class)`.



Unfortunately, JPA's `@NamedQuery` annotation can't be placed on a package descriptor. Therefore, Hibernate provides a very similar annotation, `@org.hibernate.annotations.NamedQuery` which *can* be specified at the package level. If we declare a named query at the package level, we must call:

```
configuration.addPackage("org.hibernate.example")
```

so that Hibernate knows where to find it.

The `@NamedNativeQuery` annotation lets us do the same for native SQL queries. There's much less advantage to using

@NamedNativeQuery, because there is very little that Hibernate can do to validate the correctness of a query written in the native SQL dialect of your database.

Table 40. Executing named queries

Kind	Session method	EntityManager method	Query execution method
Selection	createNamedSelectionQuery(String,Class)	createNamedQuery(String,Class)	getResultList(), getSingleResult(), or getSingleResultOrNull()
Mutation	createNamedMutationQuery(String)	createNamedQuery(String)	executeUpdate()

We execute our named query like this:

```
List<Book> books =
    entityManager.createNamedQuery("10BooksByTitle")
        .setParameter("titlePattern", titlePattern)
        .getResultList()
```

Note that the code which executes the named query is not aware of whether the query was written in HQL or in native SQL, making it slightly easier to change and optimize the query later.



It's nice to have our queries checked at startup time. It's even better to have them checked at compile time. Back in [generated-query-methods], we mentioned that the [Query Validator](#) can do that for us. In fact, the Query Validator will even check HQL query strings that occur as arguments to `createQuery()` and friends. So if we use the Query Validator, there's not much advantage to the use of named queries.

5.16. Controlling lookup by id

We can do almost anything via HQL, criteria, or native SQL queries. But when we already know the identifier of the entity we need, a query can feel like overkill. And queries don't make efficient use of the [second level cache](#).

We met the `find()` method earlier. It's the most basic way to perform a *lookup* by id. But as we also [already saw](#), it can't quite do everything. Therefore, Hibernate has some APIs that streamline certain more complicated lookups:

Table 41. Operations for lookup by id

Method name	Purpose
<code>byId()</code>	Lets us specify association fetching via an EntityGraph, as we saw; also lets us specify some additional options, including how the lookup interacts with the second level cache , and whether the entity should be loaded in read-only mode
<code>byMultipleIds()</code>	Lets us load a <i>batch</i> of ids at the same time

Batch loading is very useful when we need to retrieve multiple instances of the same entity class by id:

```
var graph = session.createEntityGraph(Book.class);
graph.addSubgraph(Book_.publisher);

List<Book> books =
    session.byMultipleIds(Book.class)
        .withFetchGraph(graph) // control association fetching
        .withBatchSize(20)    // specify an explicit batch size
        .with(CacheMode.GET)  // control interaction with the cache
        .multiLoad(bookIds);
```

The given list of `bookIds` will be broken into batches, and each batch will be fetched from the database in a single `select`. If we don't specify the batch size explicitly, a batch size will be chosen automatically.

We also have some operations for working with lookups by [natural id](#):

Method name	Purpose
<code>bySimpleNaturalId()</code>	For an entity with just one attribute is annotated @NaturalId

Method name	Purpose
<code>byNaturalId()</code>	For an entity with multiple attributes are annotated <code>@NaturalId</code>
<code>byMultipleNaturalId()</code>	Lets us load a <i>batch</i> of natural ids at the same time

Here's how we can retrieve an entity by its composite natural id:

```
Book book =
    session.byNaturalId(Book.class)
        .using(Book_.isbn, isbn)
        .using(Book_.printing, printing)
        .load();
```

Notice that this code fragment is completely typesafe, again thanks to the [Metamodel Generator](#).

5.17. Interacting directly with JDBC

From time to time we run into the need to write some code that calls JDBC directly. Unfortunately, JPA offers no good way to do this, but the Hibernate Session does.

```
session.doWork(connection -> {
    try (var callable = connection.prepareCall("{call myproc(?)}") {
        callable.setLong(1, argument);
        callable.execute();
    }
});
```

The `Connection` passed to the work is the same connection being used by the session, and so any work performed using that connection occurs in the same transaction context.

If the work returns a value, use `doReturningWork()` instead of `doWork()`.



In a container environment where transactions and database connections are managed by the container, this might not be the easiest way to obtain the JDBC connection.

5.18. What to do when things go wrong

Object/relational mapping has been called the "Vietnam of computer science". The person who made this analogy is American, and so one supposes that he meant to imply some kind of unwinnable war. This is quite ironic, since at the very moment he made this comment, Hibernate was already on the brink of winning the war.

Today, Vietnam is a peaceful country with exploding per-capita GDP, and ORM is a solved problem. That said, Hibernate is complex, and ORM still presents many pitfalls for the inexperienced, even occasionally for the experienced. Sometimes things go wrong.

In this section we'll quickly sketch some general strategies for avoiding "quagmires".

- Understand SQL and the relational model. Know the capabilities of your RDBMS. Work closely with the DBA if you're lucky enough to have one. Hibernate is not about "transparent persistence" for Java objects. It's about making two excellent technologies work smoothly together.
- [Log the SQL](#) executed by Hibernate. You cannot know that your persistence logic is correct until you've actually inspected the SQL that's being executed. Even when everything seems to be "working", there might be a lurking [N+1 selects monster](#).
- Be careful when [modifying bidirectional associations](#). In principle, you should update *both ends* of the association. But Hibernate doesn't strictly enforce that, since there are some situations where such a rule would be too heavy-handed. Whatever the case, it's up to you to maintain consistency across your model.
- Never [leak a persistence context](#) across threads or concurrent transactions. Have a strategy or framework to guarantee this never happens.
- When running queries that return large result sets, take care to consider the size of the [session cache](#). Consider using a [stateless session](#).
- Think carefully about the semantics of the [second-level cache](#), and how the caching policies impact transaction isolation.
- Avoid fancy bells and whistles you don't need. Hibernate is incredibly feature-rich, and that's a good thing, because it serves the needs of a huge number of users, many of whom have one or two very specialized needs. But nobody has *all* those specialized needs. In all probability, you have none of them. Write your domain model in the simplest way that's reasonable, using the simplest mapping strategies that make sense.
- When something isn't behaving as you expect, *simplify*. Isolate the problem. Find the absolute minimum test case which

reproduces the behavior, *before* asking for help online. Most of the time, the mere act of isolating the problem will suggest an obvious solution.

- Avoid frameworks and libraries that "wrap" JPA. If there's any one criticism of Hibernate and ORM that sometimes *does* ring true, it's that it takes you too far from direct control over JDBC. An additional layer just takes you even further.
- Avoid copy/pasting code from random bloggers or stackoverflow reply guys. Many of the suggestions you'll find online just aren't the simplest solution, and many aren't correct for Hibernate 6. Instead, *understand* what you're doing; study the Javadoc of the APIs you're using; read the JPA specification; follow the advice we give in this document; go direct to the Hibernate team on Zulip. (Sure, we can be a bit cantankerous at times, but we *do* always want you to be successful.)
- Always consider other options. You don't have to use Hibernate for *everything*.

Chapter 6. Tuning and performance

Once you have a program up and running using Hibernate to access the database, it's inevitable that you'll find places where performance is disappointing or unacceptable.

Fortunately, most performance problems are relatively easy to solve with the tools that Hibernate makes available to you, as long as you keep a couple of simple principles in mind.

First and most important: the reason you're using Hibernate is that it makes things easier. If, for a certain problem, it's making things *harder*, stop using it. Solve this problem with a different tool instead.



Just because you're using Hibernate in your program doesn't mean you have to use it *everywhere*.

Second: there are two main potential sources of performance bottlenecks in a program that uses Hibernate:

- too many round trips to the database, and
- memory consumption associated with the first-level (session) cache.

So performance tuning primarily involves reducing the number of accesses to the database, and/or controlling the size of the session cache.

But before we get to those more advanced topics, we should start by tuning the connection pool.

6.1. Tuning the connection pool

The connection pool built in to Hibernate is suitable for testing, but isn't intended for use in production. Instead, Hibernate supports a range of different connection pools, including our favorite, Agroal.

To select and configure Agroal, you'll need to set some extra configuration properties, in addition to the settings we already saw in [Basic configuration settings](#). Properties with the prefix `hibernate.agroal` are passed through to Agroal:

```
# configure Agroal connection pool
hibernate.agroal.maxSize 20
hibernate.agroal.minSize 10
hibernate.agroal.acquisitionTimeout PT1s
hibernate.agroal.reapTimeout PT10s
```

As long as you set at least one property with the prefix `hibernate.agroal`, the `AgroalConnectionProvider` will be selected automatically. There's many to choose from:

Table 42. Settings for configuring Agroal

Configuration property name	Purpose
<code>hibernate.agroal.maxSize</code>	The maximum number of connections present on the pool
<code>hibernate.agroal.minSize</code>	The minimum number of connections present on the pool
<code>hibernate.agroal.initialSize</code>	The number of connections added to the pool when it is started
<code>hibernate.agroal.maxLifetime</code>	The maximum amount of time a connection can live, after which it is removed from the pool
<code>hibernate.agroal.acquisitionTimeout</code>	The maximum amount of time a thread can wait for a connection, after which an exception is thrown instead
<code>hibernate.agroal.reapTimeout</code>	The duration for eviction of idle connections
<code>hibernate.agroal.leakTimeout</code>	The duration of time a connection can be held without causing a leak to be reported
<code>hibernate.agroal.idleValidationTimeout</code>	A foreground validation is executed if a connection has been idle on the pool for longer than this duration
<code>hibernate.agroal.validationTimeout</code>	The interval between background validation checks
<code>hibernate.agroal.initialSql</code>	A SQL command to be executed when a connection is created

The following settings are common to all connection pools supported by Hibernate:

Table 43. Common settings for connection pools

<code>hibernate.connection.autocommit</code>	The default autocommit mode
<code>hibernate.connection.isolation</code>	The default transaction isolation level

Container-managed datasources

In a container environment, you usually don't need to configure a connection pool through Hibernate. Instead, you'll use a container-managed datasource, as we saw in [Basic configuration settings](#).

6.2. Enabling statement batching

An easy way to improve performance of some transactions, with almost no work at all, is to turn on automatic DML statement batching. Batching only helps in cases where a program executes many inserts, updates, or deletes against the same table in a single transaction.

All we need to do is set a single property:

Table 44. Enabling JDBC batching

Configuration property name	Purpose	Alternative
<code>hibernate.jdbc.batch_size</code>	Maximum batch size for SQL statement batching	<code>setJdbcBatchSize()</code>




Even better than DML statement batching is the use of HQL `update` or `delete` queries, or even native SQL that calls a stored procedure!

6.3. Association fetching

Achieving high performance in ORM means minimizing the number of round trips to the database. This goal should be uppermost in your mind whenever you're writing data access code with Hibernate. The most fundamental rule of thumb in ORM is:

- explicitly specify all the data you're going to need right at the start of a session/transaction, and fetch it immediately in one or two queries,
- and only then start navigating associations between persistent entities.



Without question, the most common cause of poorly-performing data access code in Java programs is the problem of *N+1 selects*. Here, a list of *N* rows is retrieved from the database in an initial query, and then associated instances of a related entity are fetched using *N* subsequent queries.



This isn't a bug or limitation of Hibernate; this problem even affects typical handwritten JDBC code behind DAOs. Only you, the developer, can solve this problem, because only you know ahead of time what data you're going to need in a given unit of work. But that's OK. Hibernate gives you all the tools you need.

In this section we're going to discuss different ways to avoid such "chatty" interaction with the database.

Hibernate provides several strategies for efficiently fetching associations and avoiding *N+1* selects:

- *outer join fetching*—where an association is fetched using a `left outer join`,
- *batch fetching*—where an association is fetched using a subsequent `select` with a batch of primary keys, and
- *subselect fetching*—where an association is fetched using a subsequent `select` with keys re-queried in a subselect.

Of these, you should almost always use outer join fetching. But let's consider the alternatives first.

6.4. Batch fetching and subselect fetching

Consider the following code:

```
List<Book> books =
    session.createQuery("from Book order by isbn", Book.class)
        .getResultList();
books.forEach(book -> book.getAuthors().forEach(author -> out.println(book.title + " by " + author.name)));
```

This code is *very* inefficient, resulting, by default, in the execution of *N+1* `select` statements, where *n* is the number of Books.

Let's see how we can improve on that.

SQL for batch fetching

With batch fetching enabled, Hibernate might execute the following SQL on PostgreSQL:

```
/* initial query for Books */
select b1_0.isbn,b1_0.price,b1_0.published,b1_0.publisher_id,b1_0.title
from Book b1_0
order by b1_0.isbn

/* first batch of associated Authors */
select a1_0.books_isbn,a1_1.id,a1_1.bio,a1_1.name
from Book_Author a1_0
join Author a1_1 on a1_1.id=a1_0.authors_id
where a1_0.books_isbn = any (?)

/* second batch of associated Authors */
select a1_0.books_isbn,a1_1.id,a1_1.bio,a1_1.name
from Book_Author a1_0
join Author a1_1 on a1_1.id=a1_0.authors_id
where a1_0.books_isbn = any (?)
```

The first `select` statement queries and retrieves Books. The second and third queries fetch the associated Authors in batches. The number of batches required depends on the configured *batch size*. Here, two batches were required, so two SQL statements were executed.



The SQL for batch fetching looks slightly different depending on the database. Here, on PostgreSQL, Hibernate passes a batch of primary key values as a SQL `ARRAY`.

SQL for subselect fetching

On the other hand, with subselect fetching, Hibernate would execute this SQL:

```
/* initial query for Books */
select b1_0.isbn,b1_0.price,b1_0.published,b1_0.publisher_id,b1_0.title
from Book b1_0
order by b1_0.isbn

/* fetch all associated Authors */
select a1_0.books_isbn,a1_1.id,a1_1.bio,a1_1.name
from Book_Author a1_0
join Author a1_1 on a1_1.id=a1_0.authors_id
where a1_0.books_isbn in (select b1_0.isbn from Book b1_0)
```

Notice that the first query is re-executed in a subselect in the second query. The execution of the subselect is likely to be relatively inexpensive, since the data should already be cached by the database. Clever, huh?

Enabling the use of batch or subselect fetching

Both batch fetching and subselect fetching are disabled by default, but we may enable one or the other globally using properties.

Table 45. Configuration settings to enable batch and subselect fetching

Configuration property name	Property value	Alternatives
hibernate.default_batch_fetch_size	A sensible batch size >1 to enable batch fetching	@BatchSize(), setFetchBatchSize()

Alternatively, we can enable one or the other in a given session:

```
session.setFetchBatchSize(5);
```



We may request subselect fetching more selectively by annotating a collection or many-valued association with the `@Fetch` annotation.

```
@ManyToMany @Fetch(SUBSELECT)
Set<Author> authors;
```

Note that `@Fetch(SUBSELECT)` has the same effect as `@Fetch(SELECT)`, except after execution of a HQL or criteria query. But after query execution, `@Fetch(SUBSELECT)` is able to much more efficiently fetch associations.

Later, we'll see how we can use [fetch profiles](#) to do this even more selectively.

That's all there is to it. Too easy, right?

Sadly, that's not the end of the story. While batch fetching might *mitigate* problems involving N+1 selects, it won't solve them. The truly correct solution is to fetch associations using joins. Batch fetching (or subselect fetching) can only be the *best* solution in rare cases where outer join fetching would result in a cartesian product and a huge result set.

But batch fetching and subselect fetching have one important characteristic in common: they can be performed *lazily*. This is, in principle, pretty convenient. When we query data, and then navigate an object graph, lazy fetching saves us the effort of planning ahead. It turns out that this is a convenience we're going to have to surrender.

6.5. Join fetching

Outer join fetching is usually the best way to fetch associations, and it's what we use most of the time. Unfortunately, by its very nature, join fetching simply can't be lazy. So to make use of join fetching, we must plan ahead. Our general advice is:



Avoid the use of lazy fetching, which is often the source of N+1 selects.

Now, we're not saying that associations should be mapped for eager fetching by default! That would be a terrible idea, resulting in simple session operations that fetch almost the entire database. Therefore:



Most associations should be mapped for lazy fetching by default.

It sounds as if this tip is in contradiction to the previous one, but it's not. It's saying that you must explicitly specify eager fetching for associations precisely when and where they are needed.

If we need eager join fetching in some particular transaction, we have four different ways to specify that.

Passing a JPA <code>EntityGraph</code>	We've already seen this in Entity graphs and eager fetching
Specifying a named <i>fetch profile</i>	We'll discuss this approach later in Named fetch profiles
Using <code>left join fetch</code> in HQL/JPQL	See <i>A guide to Hibernate Query Language 6</i> for details
Using <code>From.fetch()</code> in a criteria query	Same semantics as <code>join fetch</code> in HQL

Typically, a query is the most convenient option. Here's how we can ask for join fetching in HQL:

```
List<Book> booksWithJoinFetchedAuthors =
    session.createSelectionQuery("from Book join fetch authors order by isbn")
        .getResultList();
```

And this is the same query, written using the criteria API:

```
var builder = sessionFactory.getCriteriaBuilder();
var query = builder.createQuery(Book.class);
var book = query.from(Book.class);
book.fetch(Book_.authors);
query.select(book);
query.orderBy(builder.asc(book.get(Book_.isbn)));
List<Book> booksWithJoinFetchedAuthors =
    session.createSelectionQuery(query).getResultList();
```

Either way, a single SQL select statement is executed:

```
select b1_0.isbn,a1_0.books_isbn,a1_1.id,a1_1.bio,a1_1.name,b1_0.price,b1_0.published,b1_0.publisher_id,b1_0
.title
from Book b1_0
join (Book_Author a1_0 join Author a1_1 on a1_1.id=a1_0.authors_id)
on b1_0.isbn=a1_0.books_isbn
order by b1_0.isbn
```

Much better!

Join fetching, despite its non-lazy nature, is clearly more efficient than either batch or subselect fetching, and this is the source of our recommendation to avoid the use of lazy fetching.



There's one interesting case where join fetching becomes inefficient: when we fetch two many-values associations *in parallel*. Imagine we wanted to fetch both `Author.books` and `Author.royaltyStatements` in some unit of work. Joining both collections in a single query would result in a cartesian product of tables, and a large SQL result set. Subselect fetching comes to the rescue here, allowing us to fetch books using a join, and `royaltyStatements` using a single subselect select.

Of course, an alternative way to avoid many round trips to the database is to cache the data we need in the Java client. If we're expecting to find the associated data in a local cache, we probably don't need join fetching at all.



But what if we can't be *certain* that all associated data will be in the cache? In that case, we might be able to reduce the cost of cache misses by enabling batch fetching.

6.6. The second-level cache

A classic way to reduce the number of accesses to the database is to use a second-level cache, allowing data cached in memory to be shared between sessions.

By nature, a second-level cache tends to undermine the ACID properties of transaction processing in a relational database. We *don't* use a distributed transaction with two-phase commit to ensure that changes to the cache and database happen atomically. So a second-level cache is often by far the easiest way to improve the performance of a system, but only at the cost of making it much more difficult to reason about concurrency. And so the cache is a potential source of bugs which are difficult to isolate and reproduce.

Therefore, by default, an entity is not eligible for storage in the second-level cache. We must explicitly mark each entity that will be stored in the second-level cache with the `@Cache` annotation from `org.hibernate.annotations`.

But that's still not enough. Hibernate does not itself contain an implementation of a second-level cache, so it's necessary to configure an external *cache provider*.



Caching is disabled by default. To minimize the risk of data loss, we force you to stop and think before any entity goes into the cache.

Hibernate segments the second-level cache into named *regions*, one for each:

- mapped entity hierarchy or
- collection role.

For example, there might be separate cache regions for `Author`, `Book`, `Author.books`, and `Book.authors`.

Each region is permitted its own policies for expiry, persistence, and replication. These policies must be configured externally to Hibernate.

The appropriate policies depend on the kind of data an entity represents. For example, a program might have different caching

policies for "reference" data, for transactional data, and for data used for analytics. Ordinarily, the implementation of those policies is the responsibility of the underlying cache implementation.

6.7. Specifying which data is cached

By default, no data is eligible for storage in the second-level cache.

An entity hierarchy or collection role may be assigned a region using the `@Cache` annotation. If no region name is explicitly specified, the region name is just the name of the entity class or collection role.

```
@Entity
@Cache(usage=NONSTRICT_READ_WRITE, region="Publishers")
class Publisher {
    ...

    @Cache(usage=READ_WRITE, region="PublishedBooks")
    @OneToMany(mappedBy=Book_.PUBLISHER)
    Set<Book> books;

    ...
}
```

The cache defined by a `@Cache` annotation is automatically utilized by Hibernate to:

- retrieve an entity by id when `find()` is called, or
- to resolve an association by id.



The `@Cache` annotation must be specified on the *root class* of an entity inheritance hierarchy. It's an error to place it on a subclass entity.

The `@Cache` annotation always specifies a `CacheConcurrencyStrategy`, a policy governing access to the second-level cache by concurrent transactions.

Table 46. Cache concurrency

Concurrency policy	Interpretation	Explanation
READ_ONLY	• Immutable data • Read-only access	Indicates that the cached object is immutable, and is never updated. If an entity with this cache concurrency is updated, an exception is thrown. This is the simplest, safest, and best-performing cache concurrency strategy. It's particularly suitable for so-called "reference" data.
NONSTRICT_READ_WRITE	• Concurrent updates are extremely improbable • Read/write access with no locking	Indicates that the cached object is sometimes updated, but that it's extremely unlikely that two transactions will attempt to update the same item of data at the same time. This strategy does not use locks. When an item is updated, the cache is invalidated both before and after completion of the updating transaction. But without locking, it's impossible to completely rule out the possibility of a second transaction storing or retrieving stale data in or from the cache during the completion process of the first transaction.

Concurrency policy	Interpretation	Explanation
READ_WRITE	• Concurrent updates are possible but not common • Read/write access using soft locks	Indicates a non-vanishing likelihood that two concurrent transactions attempt to update the same item of data simultaneously. This strategy uses "soft" locks to prevent concurrent transactions from retrieving or storing a stale item from or in the cache during the transaction completion process. A soft lock is simply a marker entry placed in the cache while the updating transaction completes. • A second transaction may not read the item from the cache while the soft lock is present, and instead simply proceeds to read the item directly from the database, exactly as if a regular cache miss had occurred. • Similarly, the soft lock also prevents this second transaction from storing a stale item to the cache when it returns from its round trip to the database with something that might not quite be the latest version.
TRANSACTIONAL	• Concurrent updates are frequent • Transactional access	Indicates that concurrent writes are common, and the only way to maintain synchronization between the second-level cache and the database is via the use of a fully transactional cache provider. In this case, the cache and the database must cooperate via JTA or the XA protocol, and Hibernate itself takes on little responsibility for maintaining the integrity of the cache.

Which policies make sense may also depend on the underlying second-level cache implementation.



JPA has a similar annotation, named `@Cacheable`. Unfortunately, it's almost useless to us, since:

- it provides no way to specify any information about the nature of the cached entity and how its cache should be managed, and
- it may not be used to annotate associations, and so we can't even use it to mark collection roles as eligible for storage in the second-level cache.

6.8. Caching by natural id

If our entity has a [natural id](#), we can enable an additional cache, which holds cross-references from natural id to primary id, by annotating the entity `@NaturalIdCache`. By default, the natural id cache is stored in a dedicated region of the second-level cache, separate from the cached entity data.

```

@Entity
@Cache(usage=READ_WRITE, region="Book")
@NaturalIdCache(region="BookIsbn")
class Book {
    ...
    @NaturalId
    String isbn;

    @NaturalId
    int printing;
    ...
}

```

This cache is utilized when the entity is retrieved using one of the operations of `Session` which performs [lookup by natural id](#).



Since the natural id cache doesn't contain the actual state of the entity, it doesn't make sense to annotate an entity `@NaturalIdCache` unless it's already eligible for storage in the second-level cache, that is, unless it's also annotated `@Cache`.

It's worth noticing that, unlike the primary identifier of an entity, a natural id might be mutable.

We must now consider a subtlety that often arises when we have to deal with so-called "reference data", that is, data which fits easily in memory, and doesn't change much.

6.9. Caching and association fetching

Let's consider again our `Publisher` class:

```
@Cache(usage=NONSTRICT_READ_WRITE, region="Publishers")
@Entity
class Publisher { ... }
```

Data about publishers doesn't change very often, and there aren't so many of them. Suppose we've set everything up so that the publishers are almost *always* available in the second-level cache.

Then in this case we need to think carefully about associations of type `Publisher`.

```
@ManyToOne
Publisher publisher;
```

There's no need for this association to be lazily fetched, since we're expecting it to be available in memory, so we won't set it `fetch=LAZY`. But on the other hand, if we leave it marked for eager fetching then, by default, Hibernate will often fetch it using a join. This places completely unnecessary load on the database.

The solution is the `@Fetch` annotation:

```
@ManyToOne @Fetch(SELECT)
Publisher publisher;
```

By annotating the association `@Fetch(SELECT)`, we suppress join fetching, giving Hibernate a chance to find the associated `Publisher` in the cache.

Therefore, we arrive at this rule of thumb:



Many-to-one associations to "reference data", or to any other data that will almost always be available in the cache, should be mapped `EAGER,SELECT`.

Other associations, as we've [already made clear](#), should be `LAZY`.

Once we've marked an entity or collection as eligible for storage in the second-level cache, we still need to set up an actual cache.

6.10. Configuring the second-level cache provider

Configuring a second-level cache provider is a rather involved topic, and quite outside the scope of this document. But in case it helps, we often test Hibernate with the following configuration, which uses `EHCache` as the cache implementation, as above in [Optional dependencies](#):

Table 47. *EHCache configuration*

Configuration property name	Property value
<code>hibernate.cache.region.factory_class</code>	<code>jcache</code>
<code>hibernate.javax.cache.uri</code>	<code>/ehcache.xml</code>

If you're using `EHCache`, you'll also need to include an `ehcache.xml` file that explicitly configures the behavior of each cache region belonging to your entities and collections. You'll find more information about configuring `EHCache` [here](#).

We may use any other implementation of `JCache`, such as [Caffeine](#). `JCache` automatically selects whichever implementation it finds on the classpath. If there are multiple implementations on the classpath, we must disambiguate using:

Table 48. *Disambiguating the JCache implementation*

Configuration property name	Property value				
<code>hibernate.javax.cache.provider</code>	The implementation of <code>javax.cache.spi.CachingProvider</code> , for example: <table><tr><td><code>org.ehcache.jsr107.EhcacheCachingProvider</code></td><td>for <code>EHCache</code></td></tr><tr><td><code>com.github.benmanes.caffeine.jcache.spi.CaffeineCachingProvider</code></td><td>for <code>Caffeine</code></td></tr></table>	<code>org.ehcache.jsr107.EhcacheCachingProvider</code>	for <code>EHCache</code>	<code>com.github.benmanes.caffeine.jcache.spi.CaffeineCachingProvider</code>	for <code>Caffeine</code>
<code>org.ehcache.jsr107.EhcacheCachingProvider</code>	for <code>EHCache</code>				
<code>com.github.benmanes.caffeine.jcache.spi.CaffeineCachingProvider</code>	for <code>Caffeine</code>				

Alternatively, to use Infinispan as the cache implementation, the following settings are required:

Table 49. Infinispan provider configuration

Configuration property name	Property value				
hibernate.cache.region.factory_class	infinispan				
hibernate.cache.infinispan.cfg	Path to infinispan configuration file, for example: <table><tr><td>org/infinispan/hibernate/cache/commons/builder/infinispan-configs.xml</td><td>for a distributed cache</td></tr><tr><td>org/infinispan/hibernate/cache/commons/builder/infinispan-configs-local.xml</td><td>to test with local cache</td></tr></table>	org/infinispan/hibernate/cache/commons/builder/infinispan-configs.xml	for a distributed cache	org/infinispan/hibernate/cache/commons/builder/infinispan-configs-local.xml	to test with local cache
org/infinispan/hibernate/cache/commons/builder/infinispan-configs.xml	for a distributed cache				
org/infinispan/hibernate/cache/commons/builder/infinispan-configs-local.xml	to test with local cache				

Infinispan is usually used when distributed caching is required. There's more about using Infinispan with Hibernate [here](#).

Finally, there's a way to globally disable the second-level cache:

Table 50. Setting to disable caching

Configuration property name	Property value
hibernate.cache.use_second_level_cache	true to enable caching, or false to disable it

When `hibernate.cache.region.factory_class` is set, this property defaults to `true`.



This setting lets us easily disable the second-level cache completely when troubleshooting or profiling performance.

You can find much more information about the second-level cache in the [User Guide](#).

6.11. Caching query result sets

The caches we've described above are only used to optimize lookups by id or by natural id. Hibernate also has a way to cache the result sets of queries, though this is only rarely an efficient thing to do.

The query cache must be enabled explicitly:

Table 51. Setting to enable the query cache

Configuration property name	Property value
hibernate.cache.use_query_cache	true to enable the query cache

To cache the results of a query, call `SelectionQuery.setCacheable(true)`:

```
session.createQuery("from Product where discontinued = false")
    .setCacheable(true)
    .getResultList();
```

By default, the query result set is stored in a cache region named `default-query-results-region`. Since different queries should have different caching policies, it's common to explicitly specify a region name:

```
session.createQuery("from Product where discontinued = false")
    .setCacheable(true)
    .setCacheRegion("ProductCatalog")
    .getResultList();
```

A result set is cached together with a *logical timestamp*. By "logical", we mean that it doesn't actually increase linearly with time, and in particular it's not the system time.

When a `Product` is updated, Hibernate *does not* go through the query cache and invalidate every cached result set that's affected by the change. Instead, there's a special region of the cache which holds a logical timestamp of the most-recent update to each table. This is called the *update timestamps cache*, and it's kept in the region `default-update-timestamps-region`.



It's *your responsibility* to ensure that this cache region is configured with appropriate policies. In particular, update timestamps should never expire or be evicted.

When a query result set is read from the cache, Hibernate compares its timestamp with the timestamp of each of the tables that affect the results of the query, and *only* returns the result set if the result set isn't stale. If the result set *is* stale, Hibernate goes ahead and re-executes the query against the database and updates the cached result set.

As is generally the case with any second-level cache, the query cache can break the ACID properties of transactions.

6.12. Second-level cache management

For the most part, the second-level cache is transparent. Program logic which interacts with the Hibernate session is unaware of the cache, and is not impacted by changes to caching policies.

At worst, interaction with the cache may be controlled by specifying of an explicit `CacheMode`:

```
session.setCacheMode(CacheMode.IGNORE);
```

Or, using JPA-standard APIs:

```
entityManager.setCacheRetrieveMode(CacheRetrieveMode.BYPASS);  
entityManager.setCacheStoreMode(CacheStoreMode.BYPASS);
```

The JPA-defined cache modes come in two flavors: `CacheRetrieveMode` and `CacheStoreMode`.

Table 52. JPA-defined cache retrieval modes

Mode	Interpretation
<code>CacheRetrieveMode.USE</code>	Read data from the cache if available
<code>CacheRetrieveMode.BYPASS</code>	Don't read data from the cache; go direct to the database

We might select `CacheRetrieveMode.BYPASS` if we're concerned about the possibility of reading stale data from the cache.

Table 53. JPA-defined cache storage modes

Mode	Interpretation
<code>CacheStoreMode.USE</code>	Write data to the cache when read from the database or when modified; do not update already-cached items when reading
<code>CacheStoreMode.REFRESH</code>	Write data to the cache when read from the database or when modified; always update cached items when reading
<code>CacheStoreMode.BYPASS</code>	Don't write data to the cache

We should select `CacheStoreMode.BYPASS` if we're querying data that doesn't need to be cached.



It's a good idea to set the `CacheStoreMode` to `BYPASS` just before running a query which returns a large result set full of data that we don't expect to need again soon. This saves work, and prevents the newly-read data from pushing out the previously cached data.

In JPA we would use this idiom:

```
entityManager.setCacheStoreMode(CacheStoreMode.BYPASS);  
List<Publisher> allpubs =  
    entityManager.createQuery("from Publisher", Publisher.class)  
        .getResultList();  
entityManager.setCacheStoreMode(CacheStoreMode.USE);
```

But Hibernate has a better way:

```
List<Publisher> allpubs =  
    session.createSelectionQuery("from Publisher", Publisher.class)  
        .setCacheStoreMode(CacheStoreMode.BYPASS)
```



```
.getResultList();
```

A Hibernate `CacheMode` packages a `CacheRetrieveMode` with a `CacheStoreMode`.

Table 54. Hibernate cache modes and JPA equivalents

Hibernate <code>CacheMode</code>	Equivalent JPA modes
NORMAL	<code>CacheRetrieveMode.USE</code> , <code>CacheStoreMode.USE</code>
IGNORE	<code>CacheRetrieveMode.BYPASS</code> , <code>CacheStoreMode.BYPASS</code>
GET	<code>CacheRetrieveMode.USE</code> , <code>CacheStoreMode.BYPASS</code>
PUT	<code>CacheRetrieveMode.BYPASS</code> , <code>CacheStoreMode.USE</code>
REFRESH	<code>CacheRetrieveMode.REFRESH</code> , <code>CacheStoreMode.BYPASS</code>

There's no particular reason to prefer Hibernate's `CacheMode` over the JPA equivalents. This enumeration only exists because Hibernate had cache modes long before they were added to JPA.



For "reference" data, that is, for data which is expected to always be found in the second-level cache, it's a good idea to *prime* the cache at startup. There's a really easy way to do this: just execute a query immediately after obtaining the `EntityManager` or `SessionFactory`.

```
SessionFactory sessionFactory =
    setupHibernate(new Configuration())
        .buildSessionFactory();
// prime the second-level cache
sessionFactory.inSession(session -> {
    session.createQuery("from Countries")
        .setReadOnly(true)
        .getResultList();
    session.createQuery("from Product where discontinued = false")
        .setReadOnly(true)
        .getResultList();
});
```

Very occasionally, it's necessary or advantageous to control the cache explicitly, for example, to evict some data that we know to be stale. The `Cache` interface allows programmatic eviction of cached items.

```
sessionFactory.getCache().evictEntityData(Book.class, bookId);
```



Second-level cache management via the `Cache` interface is not transaction-aware. None of the operations of `Cache` respect any isolation or transactional semantics associated with the underlying caches. In particular, eviction via the methods of this interface causes an immediate "hard" removal outside any current transaction and/or locking scheme.

Ordinarily, however, Hibernate automatically evicts or updates cached data after modifications, and, in addition, cached data which is unused will eventually be expired according to the configured policies.

This is quite different to what happens with the first-level cache.

6.13. Session cache management

Entity instances aren't automatically evicted from the session cache when they're no longer needed. Instead, they stay pinned in memory until the session they belong to is discarded by your program.

The methods `detach()` and `clear()` allow you to remove entities from the session cache, making them available for garbage collection. Since most sessions are rather short-lived, you won't need these operations very often. And if you find yourself thinking you *do* need them in a certain situation, you should strongly consider an alternative solution: a *stateless session*.

6.14. Stateless sessions

An arguably-underappreciated feature of Hibernate is the `StatelessSession` interface, which provides a command-oriented, more bare-metal approach to interacting with the database.

You may obtain a reactive stateless session from the `SessionFactory`:

```
Stage.StatelessSession ss = getSessionFactory().openStatelessSession();
```

A stateless session:

- doesn't have a first-level cache (persistence context), nor does it interact with any second-level caches, and
- doesn't implement transactional write-behind or automatic dirty checking, so all operations are executed immediately when they're explicitly called.

For a stateless session, we're always working with detached objects. Thus, the programming model is a bit different:

Table 55. Important methods of the `StatelessSession`

Method name and parameters	Effect
<code>get(Class, Object)</code>	Obtain a detached object, given its type and its id, by executing a <code>select</code>
<code>fetch(Object)</code>	Fetch an association of a detached object
<code>refresh(Object)</code>	Refresh the state of a detached object by executing a <code>select</code>
<code>insert(Object)</code>	Immediately insert the state of the given transient object into the database
<code>update(Object)</code>	Immediately update the state of the given detached object in the database
<code>delete(Object)</code>	Immediately delete the state of the given detached object from the database
<code>`upsert(Object)1</code>	Immediately insert or update the state of the given detached object using a SQL <code>merge into</code> statement



There's no `flush()` operation, and so `update()` is always explicit.

In certain circumstances, this makes stateless sessions easier to work with, but with the caveat that a stateless session is much more vulnerable to data aliasing effects, since it's easy to get two non-identical Java objects which both represent the same row of a database table.



If we use `fetch()` in a stateless session, we can very easily obtain two objects representing the same database row!

In particular, the absence of a persistence context means that we can safely perform bulk-processing tasks without allocating huge quantities of memory. Use of a `StatelessSession` alleviates the need to call:

- `clear()` or `detach()` to perform first-level cache management, and
- `setCacheMode()` to bypass interaction with the second-level cache.



Stateless sessions can be useful, but for bulk operations on huge datasets, Hibernate can't possibly compete with stored procedures!

When using a stateless session, you should be aware of the following additional limitations:

- persistence operations never cascade to associated instances,
- changes to `@ManyToMany` associations and `@ElementCollections` cannot be made persistent, and
- operations performed via a stateless session bypass callbacks.

6.15. Optimistic and pessimistic locking

Finally, an aspect of behavior under load that we didn't mention above is row-level data contention. When many transactions try to read and update the same data, the program might become unresponsive with lock escalation, deadlocks, and lock acquisition timeout errors.

There's two basic approaches to data concurrency in Hibernate:

- optimistic locking using `@Version` columns, and
- database-level pessimistic locking using the SQL `for update` syntax (or equivalent).

In the Hibernate community it's *much* more common to use optimistic locking, and Hibernate makes that incredibly easy.



Where possible, in a multiuser system, avoid holding a pessimistic lock across a user interaction. Indeed, the usual practice is to avoid having transactions that span user interactions. For multiuser systems, optimistic locking is king.

That said, there *is* also a place for pessimistic locks, which can sometimes reduce the probability of transaction rollbacks.

Therefore, the `find()`, `lock()`, and `refresh()` methods of the reactive session accept an optional `LockMode`. We can also specify a `LockMode` for a query. The lock mode can be used to request a pessimistic lock, or to customize the behavior of optimistic locking:

Table 56. Optimistic and pessimistic lock modes

LockMode type	Meaning
READ	An optimistic lock obtained implicitly whenever an entity is read from the database using <code>select</code>
OPTIMISTIC	An optimistic lock obtained when an entity is read from the database, and verified using a <code>select</code> to check the version when the transaction completes
OPTIMISTIC_FORCE_INCREMENT	An optimistic lock obtained when an entity is read from the database, and enforced using an update to increment the version when the transaction completes
WRITE	A pessimistic lock obtained implicitly whenever an entity is written to the database using <code>update</code> or <code>insert</code>
PESSIMISTIC_READ	A pessimistic for <code>share</code> lock
PESSIMISTIC_WRITE	A pessimistic for <code>update</code> lock
PESSIMISTIC_FORCE_INCREMENT	A pessimistic lock enforced using an immediate update to increment the version

6.16. Collecting statistics

We may ask Hibernate to collect statistics about its activity by setting this configuration property:

Configuration property name	Property value
<code>hibernate.generate_statistics</code>	true to enable collection of statistics

The statistics are exposed by the `Statistics` object:

```
long failedVersionChecks =
    sessionFactory.getStatistics()
        .getOptimisticFailureCount();

long publisherCacheMissCount =
    sessionFactory.getStatistics()
        .getEntityStatistics(Publisher.class.getName())
        .getCacheMissCount()
```

Hibernate's statistics enable observability. Both [Micrometer](#) and [SmallRye Metrics](#) are capable of exposing these metrics.

6.17. Tracking down slow queries

When a poorly-performing SQL query is discovered in production, it can sometimes be hard to track down exactly where in the Java code the query originates. Hibernate offers two configuration properties that can make it easier to identify a slow query and find its source.

Table 57. Settings for tracking slow queries

Configuration property name	Purpose	Property value
<code>hibernate.log_slow_query</code>	Log slow queries at the INFO level	The minimum execution time, in milliseconds, which characterizes a "slow" query

Configuration property name	Purpose	Property value
<code>hibernate.use_sql_comments</code>	Prepend comments to the executed SQL	true or false

When `hibernate.use_sql_comments` is enabled, the text of the HQL query is prepended as a comment to the generated SQL, which usually makes it easy to find the HQL in the Java code.

The comment text may be customized:

- by calling `Query.setComment(comment)` or `Query.setHint(AvailableHints.HINT_COMMENT, comment)`, or
- via the `@NamedQuery` annotation.

6.18. Reactive programming with Hibernate

Finally, many systems which require high scalability now make use of reactive programming and reactive streams. [Hibernate Reactive](#) brings O/R mapping to the world of reactive programming. You can learn much more about Hibernate Reactive from its [Reference Documentation](#).



Hibernate Reactive may be used alongside vanilla Hibernate in the same program, and can reuse the same entity classes. This means you can use the reactive programming model exactly where you need it—perhaps only in one or two places in your system. You don't need to rewrite your whole program using reactive streams.

Chapter 7. Advanced Topics

In the last chapter of this Introduction, we turn to some topics that don't really belong in an introduction. Here we consider some problems, and solutions, that you're probably not going to run into immediately if you're new to Hibernate. But we do want you to know *about* them, so that when the time comes, you'll know what tool to reach for.

7.1. Filters

Filters are one of the nicest and under-usedest features of Hibernate, and we're quite proud of them. A filter is a named, globally-defined, parameterized restriction on the data that is visible in a given session.

Examples of well-defined filters might include:

- a filter that restricts the data visible to a given user according to row-level permissions,
- a filter which hides data which has been soft-deleted,
- in a versioned database, a filter that displays versions which were current at a given instant in the past, or
- a filter that restricts to data associated with a certain geographical region.

A filter must be declared somewhere. A package descriptor is as good a place as any:

```
@FilterDef(name = "ByRegion",
           parameters = @ParamDef(name = "region", type = String.class))
package org.hibernate.example;
```

This filter has one parameter. Fancier filters might in principle have multiple parameters, though we admit this must be quite rare.



If you add annotations to a package descriptor, and you're using `Configuration` to configure Hibernate, make sure you call `Configuration.addPackage()` to let Hibernate know that the package descriptor is annotated.

Typically, but not necessarily, a `@FilterDef` specifies a default restriction:

```
@FilterDef(name = "ByRegion",
           parameters = @ParamDef(name = "region", type = String.class),
           defaultCondition = "region = :region")
package org.hibernate.example;
```

The restriction must contain a reference to the parameter of the filter, specified using the usual syntax for named parameters.

Any entity or collection which is affected by a filter must be annotated `@Filter`:

```
@Entity
@Filter(name = example_.BY_REGION)
class User {

    @Id String username;

    String region;

    ...
}
```

If the `@Filter` annotation does not explicitly specify a restriction, the default restriction given by the `@FilterDef` will be applied to the entity. But an entity is free to override the default condition.

```
@Entity
@Filter(name = "ByRegion", condition = "name = :region")
class Region {

    @Id String name;

    ...
}
```

Note that the restriction specified by the `condition` or `defaultCondition` is a native SQL expression.

Table 58. Annotations for defining filters

Annotation	Purpose
@FilterDef	Defines a filter and declares its name (exactly one per filter)
@Filter	Specifies how a filter applies to a given entity or collection (many per filter)

By default, a new session comes with every filter disabled. A filter may be explicitly enabled in a given session by calling `enableFilter()` and assigning arguments to the parameters of the filter. You should do this right at the *start* of the session.

```
sessionFactory.inTransaction(session -> {
    session.enableFilter("ByRegion")
        .setParameter("region", "es")
        .validate();
    ...
});
```

Now, any queries executed within the session will have the filter restriction applied. Collections annotated @Filter will also have their members correctly filtered.



On the other hand, filters are not applied to @ManyToOne associations, nor to `find()`. This is completely by design and is not in any way a bug.

More than one filter may be enabled in a given session.

We've mentioned that a filter can be used to implement versioning, and to provide *historical* views of the data. Being such a general-purpose construct, filters provide a lot of flexibility here. But if you're after a more focused/opinionated solution to this problem, you should definitely check out [Envers](#).

Using Envers for auditing historical data

Envers is an add-on to Hibernate ORM which keeps a historical record of each versioned entity in a separate *audit table*, and allows past revisions of the data to be viewed and queried. A full introduction to Envers would require a whole chapter, so we'll just give you a quick taste here.

First, we must mark an entity as versioned, using the @Audited annotation:

```
@Audited @Entity
@Table(name="CurrentDocument")
@AuditTable("DocumentRevision")
class Document { ... }
```



The @AuditTable annotation is optional, and it's better to set either `org.hibernate.envers.audit_table_prefix` or `org.hibernate.envers.audit_table_suffix` and let the audit table name be inferred.

The `AuditReader` interface exposes operations for retrieving and querying historical revisions. It's really easy to get hold of one of these:

```
AuditReader reader = AuditReaderFactory.get(entityManager);
```

Envers tracks revisions of the data via a global *revision number*. We may easily find the revision number which was current at a given instant:

```
Number revision = reader.getRevisionNumberForDate(datetime);
```

We can use the revision number to ask for the version of our entity associated with the given revision number:

```
Document doc = reader.find(Document.class, id, revision);
```

Alternatively, we can directly ask for the version which was current at a given instant:

```
Document doc = reader.find(Document.class, id, datetime);
```

We can even execute queries to obtain lists of entities current at the given revision number:

```
List documents =  
    reader.createQuery()  
        .forEntitiesAtRevision(Document.class, revision)  
        .getResultList();
```

For much more information, see [the User Guide](#).

Another closely-related problem is multi-tenancy.

7.2. Multi-tenancy

A *multi-tenant* database is one where the data is segregated by *tenant*. We don't need to actually define what a "tenant" really represents here; all we care about at this level of abstraction is that each tenant may be distinguished by a unique identifier. And that there's a well-defined *current tenant* in each session.

We may specify the current tenant when we open a session:

```
var session =  
    sessionFactory.withOptions()  
        .tenantIdentifier(tenantId)  
        .openSession();
```

Or, when using JPA-standard APIs:

```
var entityManager =  
    entityManagerFactory.createEntityManager(Map.of(HibernateHints.HINT_TENANT_ID, tenantId));
```

However, since we often don't have this level of control over creation of the session, it's more common to supply an implementation of `CurrentTenantIdentifierResolver` to Hibernate.

There are three common ways to implement multi-tenancy:

1. each tenant has its own database,
2. each tenant has its own schema, or
3. tenants share tables in a single schema, and rows are tagged with the tenant id.

From the point of view of Hibernate, there's little difference between the first two options. Hibernate will need to obtain a JDBC connection with permissions on the database and schema owned by the current tenant.

Therefore, we must implement a `MultiTenantConnectionProvider` which takes on this responsibility:

- from time to time, Hibernate will ask for a connection, passing the id of the current tenant, and then we must create an appropriate connection or obtain one from a pool, and return it to Hibernate, and
- later, Hibernate will release the connection and ask us to destroy it or return it to the appropriate pool.



Check out `DataSourceBasedMultiTenantConnectionProviderImpl` for inspiration.

The third option is quite different. In this case we don't need a `MultiTenantConnectionProvider`, but we will need a dedicated column holding the tenant id mapped by each of our entities.

```
@Entity  
class Account {  
    @Id String id;  
    @TenantId String tenantId;  
  
    ...  
}
```

The `@TenantId` annotation is used to indicate an attribute of an entity which holds the tenant id. Within a given session, our data is automatically filtered so that only rows tagged with the tenant id of the current tenant are visible in that session.



Native SQL queries are *not* automatically filtered by tenant id; you'll have to do that part yourself.

Table 59. Multi-tenancy configuration

Configuration property name	Purpose
hibernate.tenant_identifier_resolver	Specifies the CurrentTenantIdentifierResolver
hibernate.multi_tenant_connection_provider	Specifies the MultiTenantConnectionProvider



If you only need to filter rows by a static condition with no parameters, `@Where` is a much simpler way to do that.

7.3. Using custom-written SQL

We've already discussed how to run [queries written in SQL](#), but occasionally that's not enough. Sometimes—but much less often than you might expect—we would like to customize the SQL used by Hibernate to perform basic CRUD operations for an entity or collection.

For this we can use `@SQLInsert` and friends:

```
@Entity
@SQLInsert(sql = "insert into person (name, id, valid) values (?, ?, true)", check = COUNT)
@SQLUpdate(sql = "update person set name = ? where id = ?")
@SQLDelete(sql = "update person set valid = false where id = ?")
@SQLSelect(sql = "select id, name from person where id = ? and valid = true")
public static class Person { ... }
```



If the custom SQL should be executed via a `CallableStatement`, just specify `callable=true`.

Any SQL statement specified by one of these annotations must have exactly the number of JDBC parameters that Hibernate expects, that is, one for each column mapped by the entity, in the exact order Hibernate expects. In particular, the primary key columns must come last.

However, the `@Column` annotation does lend some flexibility here:

- if a column should not be written as part of the custom insert statement, and has no corresponding JDBC parameter in the custom SQL, map it `@Column(insertable=false)`, or
- if a column should not be written as part of the custom update statement, and has no corresponding JDBC parameter in the custom SQL, map it `@Column(updatable=false)`.



If you need custom SQL, but are targeting multiple dialects of SQL, you can use the annotations defined in `DialectOverrides`. For example, this annotation lets us override the custom insert statement just for PostgreSQL:

```
@DialectOverride.SQLInsert(dialect = PostgreSQLDialect.class,
    override = @SQLInsert(sql="insert into person (name,id) values (?,gen_random_uuid())"))
```

It's even possible to override the custom SQL for specific *versions* of a database.

Sometimes a custom insert or update statement assigns a value to a mapped column which is calculated when the statement is executed on the database. For example, the value might be obtained by calling a SQL function:

```
@SQLInsert(sql = "insert into person (name, id) values (?, gen_random_uuid())")
```

But the entity instance which represents the row being inserted or updated won't be automatically populated with that value. And so our persistence context loses synchronization with the database. In situations like this, we may use the `@Generated` annotation to tell Hibernate to reread the state of the entity after each insert or update.

7.4. Handling database-generated columns

Sometimes, a column value is assigned or mutated by events that happen in the database, and aren't visible to Hibernate. For example:

- a table might have a column value populated by a trigger,

- a mapped column might have a default value defined in DDL, or
- a custom SQL insert or update statement might assign a value to a mapped column, as we saw in the previous subsection.

One way to deal with this situation is to explicitly call `refresh()` at appropriate moments, forcing the session to reread the state of the entity. But this is annoying.

The `@Generated` annotation relieves us of the burden of explicitly calling `refresh()`. It specifies that the value of the annotated entity attribute is generated by the database, and that the generated value should be automatically retrieved using a SQL `returning` clause, or separate `select` after it is generated.

A useful example is the following mapping:

```
@Entity
class Entity {
    @Generated @Id
    @ColumnDefault("gen_random_uuid()")
    UUID id;
}
```

The generated DDL is:

```
create table Entity (
    id uuid default gen_random_uuid() not null,
    primary key (uuid)
)
```

So here the value of `id` is defined by the column default clause, by calling the PostgreSQL function `gen_random_uuid()`.

When a column value is generated during updates, use `@Generated(event=UPDATE)`. When a value is generated by both inserts *and* updates, use `@Generated(event={INSERT, UPDATE})`.



For columns which should be generated using a SQL `generated always as` clause, prefer the `@GeneratedColumn` annotation, so that Hibernate automatically generates the correct DDL.

Actually, the `@Generated` and `@GeneratedColumn` annotations are defined in terms of a more generic and user-extensible framework for handling attribute values generated in Java, or by the database.

7.5. User-defined generators

JPA doesn't define a standard way to extend the set of id generation strategies, but Hibernate does:

- the `Generator` hierarchy of interfaces in the package `org.hibernate.generator` lets you define new generators, and
- the `@IdGeneratorType` meta-annotation from the package `org.hibernate.annotations` lets you write an annotation which associates a `Generator` type with identifier attributes.

Furthermore, the `@ValueGenerationType` meta-annotation lets you write an annotation which associates a `Generator` type with a non-`@Id` attribute.



These APIs are new in Hibernate 6, and supersede the classic `IdentifierGenerator` interface and `@GenericGenerator` annotation from older versions of Hibernate. However, the older APIs are still available and custom `IdentifierGenerators` written for older versions of Hibernate continue to work in Hibernate 6.

Hibernate has a range of built-in generators which are defined in terms of this new framework.

Table 60. Built-in generators

Annotation	Implementation	Purpose
<code>@Generated</code>	<code>GeneratedGeneration</code>	Generically handles database-generated values
<code>@GeneratedColumn</code>	<code>GeneratedAlwaysGeneration</code>	Handles values generated using <code>generated always</code>
<code>@CurrentTimestamp</code>	<code>CurrentTimestampGeneration</code>	Generic support for database or in-memory generation of creation or update timestamps
<code>@CreationTimestamp</code>	<code>CurrentTimestampGeneration</code>	A timestamp generated when an entity is first made persistent

Annotation	Implementation	Purpose
@UpdateTimestamp	CurrentTimestampGeneration	A timestamp generated when an entity is made persistent, and regenerated every time the entity is modified
@UuidGenerator	UuidGenerator	A more flexible generator for RFC 4122 UUIDs

Furthermore, support for JPA's standard id generation strategies is also defined in terms of this framework.

As an example, let's look at how @UuidGenerator is defined:

```
@IdGeneratorType(org.hibernate.id.uuid.UuidGenerator.class)
@ValueGenerationType(generatedBy = org.hibernate.id.uuid.UuidGenerator.class)
@Retention(RUNTIME)
@Target({ FIELD, METHOD })
public @interface UuidGenerator { ... }
```

@UuidGenerator is meta-annotated both @IdGeneratorType and @ValueGenerationType because it may be used to generate both ids and values of regular attributes. Either way, this Generator class does the hard work:

```
public class UuidGenerator
    // this generator produced values before SQL is executed
    implements BeforeExecutionGenerator {

    // constructors accept an instance of the @UuidGenerator
    // annotation, allowing the generator to be "configured"

    // called to create an id generator
    public UuidGenerator(
        org.hibernate.annotations.UuidGenerator config,
        Member idMember,
        CustomIdGeneratorCreationContext creationContext) {
        this(config, idMember);
    }

    // called to create a generator for a regular attribute
    public UuidGenerator(
        org.hibernate.annotations.UuidGenerator config,
        Member member,
        GeneratorCreationContext creationContext) {
        this(config, idMember);
    }

    ...

    @Override
    public EnumSet<EventType> getEventTypes() {
        // UUIDs are only assigned on insert, and never regenerated
        return INSERT_ONLY;
    }

    @Override
    public Object generate(SharedSessionContractImplementor session, Object owner, Object currentValue,
        EventType eventType) {
        // actually generate a UUID and transform it to the required type
        return valueTransformer.transform( generator.generateUuid( session ) );
    }
}
```

You can find out more about custom generators from the Javadoc for @IdGeneratorType and for org.hibernate.generator.

7.6. Naming strategies

When working with a pre-existing relational schema, it's usual to find that the column and table naming conventions used in the schema don't match Java's naming conventions.

Of course, the @Table and @Column annotations let us explicitly specify a mapped table or column name. But we would prefer to avoid scattering these annotations across our whole domain model.

Therefore, Hibernate lets us define a mapping between Java naming conventions, and the naming conventions of the relational

schema. Such a mapping is called a *naming strategy*.

First, we need to understand how Hibernate assigns and processes names.

- *Logical naming* is the process of applying naming rules to determine the *logical names* of objects which were not explicitly assigned names in the O/R mapping. That is, when there's no `@Table` or `@Column` annotation.
- *Physical naming* is the process of applying additional rules to transform a logical name into an actual "physical" name that will be used in the database. For example, the rules might include things like using standardized abbreviations, or trimming the length of identifiers.

Thus, there's two flavors of naming strategy, with slightly different responsibilities. Hibernate comes with default implementations of these interfaces:

Flavor	Default implementation
An <code>ImplicitNamingStrategy</code> is responsible for assigning a logical name when none is specified by an annotation	A default strategy which implements the rules defined by JPA
A <code>PhysicalNamingStrategy</code> is responsible for transforming a logical name and producing the name used in the database	A trivial implementation which does no processing



We happen to not much like the naming rules defined by JPA, which specify that mixed case and camel case identifiers should be concatenated using underscores. We bet you could easily come up with a much better `ImplicitNamingStrategy` than that! (Hint: it should always produce legit mixed case identifiers.)



A popular `PhysicalNamingStrategy` produces snake case identifiers.

Custom naming strategies may be enabled using the configuration properties we already mentioned without much explanation back in [Minimizing repetitive mapping information](#).

Table 61. Naming strategy configuration

Configuration property name	Purpose
<code>hibernate.implicit_naming_strategy</code>	Specifies the <code>ImplicitNamingStrategy</code>
<code>hibernate.physical_naming_strategy</code>	Specifies the <code>PhysicalNamingStrategy</code>

7.7. Spatial datatypes

Hibernate Spatial augments the [built-in basic types](#) with a set of Java mappings for OGC spatial types.

- `Geolatte-geom` defines a set of Java types implementing the OGC spatial types, and codecs for translating to and from database-native spatial datatypes.
- Hibernate Spatial itself supplies integration with Hibernate.

To use Hibernate Spatial, we must add it as a dependency, as described in [Optional dependencies](#).

Then we may immediately use Geolatte-geom and JTS types in our entities. No special annotations are needed:

```
import org.locationtech.jts.geom.Point;
import jakarta.persistence.*;

@Entity
class Event {
    Event() {}

    Event(String name, Point location) {
        this.name = name;
        this.location = location;
    }

    @Id @GeneratedValue
    Long id;

    String name;

    Point location;
}
```

```
}
```

The generated DDL uses `geometry` as the type of the column mapped by `location`:

```
create table Event (
  id bigint not null,
  location geometry,
  name varchar(255),
  primary key (id)
)
```

Hibernate Spatial lets us work with spatial types just as we would with any of the built-in basic attribute types.

```
var geometryFactory = new GeometryFactory();
...

Point point = geometryFactory.createPoint(new Coordinate(10, 5));
session.persist(new Event("Hibernate ORM presentation", point));
```

But what makes this powerful is that we may write some very fancy queries involving functions of spatial types:

```
Polygon triangle =
    geometryFactory.createPolygon(
        new Coordinate[] {
            new Coordinate(9, 4),
            new Coordinate(11, 4),
            new Coordinate(11, 20),
            new Coordinate(9, 4)
        }
    );
Point event =
    session.createQuery("select location from Event where within(location, :zone) = true", Point.class)
        .setParameter("zone", triangle)
        .getSingleResult();
```

Here, `within()` is one of the functions for testing spatial relations defined by the OpenGIS specification. Other such functions include `touches()`, `intersects()`, `distance()`, `boundary()`, etc. Not every spatial relation function is supported on every database. A matrix of support for spatial relation functions may be found in the [User Guide](#).



If you want to play with spatial functions on H2, run the following code first:

```
sessionFactory.inTransaction(session -> {
    session.doWork(connection -> {
        try (var statement = connection.createStatement()) {
            statement.execute("create alias if not exists h2gis_spatial for
\"org.h2gis.functions.factory.H2GISFunctions.load\"");
            statement.execute("call h2gis_spatial()");
        }
    });
} );
```

7.8. Ordered and sorted collections and map keys

Java lists and maps don't map very naturally to foreign key relationships between tables, and so we tend to avoid using them to represent associations between our entity classes. But if you feel like you *really* need a collection with a fancier structure than `Set`, Hibernate does have options.

The first two options let us map the index of a `List` or key of a `Map` to a column:

Table 62. Annotations for mapping lists and maps

Annotation	Purpose	JPA-standard
<code>@OrderColumn</code>	Specifies the column used to maintain the order of a list	✓
<code>@MapKeyColumn</code>	Specifies the column used to persist the keys of a map	✓

For an unowned association, the column must also be mapped on the owning side, usually by an attribute of the target entity.

Now, let's introduce a little distinction:

- an *ordered collection* is one with an ordering maintained in the database, and
- a *sorted collection* is one which is sorted in Java code.

These annotations allow us to specify how the elements of a collection should be ordered as they are read from the database:

Table 63. Annotations for ordered collections

Annotation	Purpose	JPA-standard
@OrderBy	Specifies a fragment of JPQL used to order the collection	✓

On the other hand, the following annotation specify how a collection should be sorted in memory, and are used for collections of type `SortedSet` or `SortedMap`:

Table 64. Annotations for sorted collections

Annotation	Purpose	JPA-standard
@SortNatural	Specifies that the elements of a collection are <code>Comparable</code>	✗
@SortComparator	Specifies a <code>Comparator</code> used to sort the collection	✗

Under the covers, Hibernate uses a `TreeSet` or `TreeMap` to maintain the collection in sorted order.

7.9. Any mappings

An `@Any` mapping is a sort of polymorphic many-to-one association where the target entity types are not related by the usual entity inheritance. The target type is distinguished using a discriminator value stored on the *referring* side of the relationship.

This is quite different to [discriminated inheritance](#) where the discriminator is held in the tables mapped by the referenced entity hierarchy.

For example, consider an `Order` entity containing `Payment` information, where a `Payment` might be a `CashPayment` or a `CreditCardPayment`:

```
interface Payment { ... }

@Entity
class CashPayment { ... }

@Entity
class CreditCardPayment { ... }
```

In this example, `Payment` is not declared as an entity type, and is not annotated `@Entity`. It might even be an interface, or at most just a mapped superclass, of `CashPayment` and `CreditCardPayment`. So in terms of the object/relational mappings, `CashPayment` and `CreditCardPayment` would not be considered to participate in the same entity inheritance hierarchy.

On the other hand, `CashPayment` and `CreditCardPayment` do have the same identifier type. This is important.

An `@Any` mapping would store the discriminator value identifying the concrete type of `Payment` along with the state of the associated `Order`, instead of storing it in the table mapped by `Payment`.

```
@Entity
class Order {
    ...

    @Any
    @AnyKeyJavaClass(UUID.class) //the foreign key type
    @JoinColumn(name="payment_id") // the foreign key column
    @Column(name="payment_type") // the discriminator column
    // map from discriminator values to target entity types
    @AnyDiscriminatorValue(discriminator="CASH", entity=CashPayment.class)
    @AnyDiscriminatorValue(discriminator="CREDIT", entity=CreditCardPayment.class)
    Payment payment;

    ...
}
```

}

It's reasonable to think of the "foreign key" in an `@Any` mapping as a composite value made up of the foreign key and discriminator taken together. Note, however, that this composite foreign key is only conceptual and cannot be declared as a physical constraint on the relational database table.

There are a number of annotations which are useful to express this sort of complicated and unnatural mapping:

Table 65. Annotations for `@Any` mappings

Annotations	Purpose
<code>@Any</code>	Declares that an attribute is a discriminated polymorphic association mapping
<code>@AnyDiscriminator</code>	Specify the Java type of the discriminator
<code>@JdbcType</code> or <code>@JdbcTypeCode</code>	Specify the JDBC type of the discriminator
<code>@AnyDiscriminatorValue</code>	Specifies how discriminator values map to entity types
<code>@Column</code> or <code>@Formula</code>	Specify the column or formula in which the discriminator value is stored
<code>@AnyKeyJavaType</code> or <code>@AnyKeyJavaClass</code>	Specify the Java type of the foreign key (that is, of the ids of the target entities)
<code>@AnyKeyJdbcType</code> or <code>@AnyKeyJdbcTypeCode</code>	Specify the JDBC type of the foreign key
<code>@JoinColumn</code>	Specifies the foreign key column

Of course, `@Any` mappings are disfavored, except in extremely special cases, since it's much more difficult to enforce referential integrity at the database level.

There's also currently some limitations around querying `@Any` associations in HQL. This is allowed:

```
from Order ord
join CashPayment cash
on id(ord.payment) = cash.id
```



Polymorphic association joins for `@Any` mappings are not currently implemented.

7.10. Selective column lists in inserts and updates

By default, Hibernate generates `insert` and `update` statements for each entity during bootstrap, and reuses the same `insert` statement every time an instance of the entity is made persistent, and the same `update` statement every time an instance of the entity is modified.

This means that:

- if an attribute is `null` when the entity is made persistent, its mapped column is redundantly included in the SQL `insert`, and
- worse, if a certain attribute is unmodified when other attributes are changed, the column mapped by that attribute is redundantly included in the SQL `update`.

Most of the time this just isn't an issue worth worrying about. The cost of interacting with the database is *usually* dominated by the cost of a round trip, not by the number of columns in the `insert` or `update`. But in cases where it does become important, there are two ways to be more selective about which columns are included in the SQL.

The JPA-standard way is to indicate statically which columns are eligible for inclusion via the `@Column` annotation. For example, if an entity is always created with an immutable `creationDate`, and with no `completionDate`, then we would write:

```
@Column(updatable=false) LocalDate creationDate;
@Column(insertable=false) LocalDate completionDate;
```

This approach works quite well in many cases, but often breaks down for entities with more than a handful of updatable columns.

An alternative solution is to ask Hibernate to generate SQL dynamically each time an `insert` or `update` is executed. We do this by

annotating the entity class.

Table 66. Annotations for dynamic SQL generation

Annotation	Purpose
@DynamicInsert	Specifies that an insert statement should be generated each time an entity is made persistent
@DynamicUpdate	Specifies that an update statement should be generated each time an entity is modified

It's important to realize that, while @DynamicInsert has no impact on semantics, the more useful @DynamicUpdate annotation *does* have a subtle side effect.



The wrinkle is that if an entity has no version property, @DynamicUpdate opens the possibility of two optimistic transactions concurrently reading and selectively updating a given instance of the entity. In principle, this might lead to a row with inconsistent column values after both optimistic transactions commit successfully.

Of course, this consideration doesn't arise for entities with a @Version attribute.



But there's a solution! Well-designed relational schemas should have *constraints* to ensure data integrity. That's true no matter what measures we take to preserve integrity in our program logic. We may ask Hibernate to add a check constraint to our table using the @Check annotation. Check constraints and foreign key constraints can help ensure that a row never contains inconsistent column values.

7.11. Using the bytecode enhancer

Hibernate's [bytecode enhancer](#) enables the following features:

- *attribute-level lazy fetching* for basic attributes annotated @Basic(fetch=LAZY) and for lazy non-polymorphic associations,
- *interception-based*—instead of the usual *snapshot-based*—detection of modifications.

To use the bytecode enhancer, we must add the Hibernate plugin to our gradle build:

```
plugins {  
    id "org.hibernate.orm" version "6.2.2.Final"  
}  
  
hibernate { enhancement }
```

Consider this field:

```
@Entity  
class Book {  
    ...  
  
    @Basic(optional = false, fetch = LAZY)  
    @Column(length = LONG32)  
    String fullText;  
    ...  
}
```

The fullText field maps to a clob or text column, depending on the SQL dialect. Since it's expensive to retrieve the full book-length text, we've mapped the field fetch=LAZY, telling Hibernate not to read the field until it's actually used.

- *Without* the bytecode enhancer, this instruction is ignored, and the field is always fetched immediately, as part of the initial select that retrieves the Book entity.
- *With* bytecode enhancement, Hibernate is able to detect access to the field, and lazy fetching is possible.



By default, Hibernate fetches all lazy fields of a given entity at once, in a single select, when any one of them is accessed. Using the @LazyGroup annotation, it's possible to assign fields to distinct "fetch groups", so that different lazy fields may be fetched independently.

Similarly, interception lets us implement lazy fetching for non-polymorphic associations without the need for a separate proxy object. However, if an association is polymorphic, that is, if the target entity type has subclasses, then a proxy is still required.

Interception-based change detection is a nice performance optimization with a slight cost in terms of correctness.

- *Without* the bytecode enhancer, Hibernate keeps a snapshot of the state of each entity after reading from or writing to the database. When the session flushes, the snapshot state is compared to the current state of the entity to determine if the entity has been modified. Maintaining these snapshots does have an impact on performance.
- *With* bytecode enhancement, we may avoid this cost by intercepting writes to the field and recording these modifications as they happen.

This optimization isn't *completely* transparent, however.



Interception-based change detection is less accurate than snapshot-based dirty checking. For example, consider this attribute:

```
byte[] image;
```

Interception is able to detect writes to the `image` field, that is, replacement of the whole array. It's not able to detect modifications made directly to the *elements* of the array, and so such modifications may be lost.

7.12. Named fetch profiles

We've already seen two different ways to override the default [fetching strategy](#) for an association:

- JPA entity graphs, and
- the `join fetch` clause in [HQL](#), or, equivalently, the method `From.fetch()` in the criteria query API.

A third way is to define a named fetch profile. First, we must declare the profile, by annotating a class or package:

```
@FetchProfile(name = "EagerBook")
@Entity
class Book { ... }
```

Note that even though we've placed this annotation on the `Book` entity, a fetch profile—unlike an entity graph—*isn't "rooted"* at any particular entity.

We may specify association fetching strategies using the `fetchOverrides` member of the `@FetchProfile` annotation.



Similarly, a JPA [entity graph](#) may be defined using `@NamedEntityGraph`. But the format of this annotation is *even worse* than `@FetchProfile(fetchOverrides=...)`, so we can't recommend it. 🤖

A better way is to annotate an association with the fetch profiles it should be fetched in:

```
@FetchProfile(name = "EagerBook", fetchOverrides = {
    @FetchProfile.FetchOverride(entity = Book.class, association = "publisher", mode = JOIN),
    @FetchProfile.FetchOverride(entity = Book.class, association = "authors", mode = JOIN),
    @FetchProfile.FetchOverride(entity = Author.class, association = "person", mode = JOIN)
})
@Entity
class Book {
    ...

    @ManyToOne(fetch = LAZY)
    Publisher publisher;

    @ManyToMany
    Set<Author> authors;

    ...
}

@Entity
class Author {
    ...

    @OneToOne
    Person person;

    ...
}
```


For collections, we may even request subselect fetching:

```
@FetchProfile(name = "EagerBook", fetchOverrides = {
    @FetchProfile.FetchOverride(entity = Book.class, association = "person", mode = JOIN),
    @FetchProfile.FetchOverride(entity = Book.class, association = "authors", mode = JOIN)
})
@FetchProfile(name = "BookWithAuthorsBySubselect", fetchOverrides = {
    @FetchProfile.FetchOverride(entity = Book.class, association = "authors", mode = SUBSELECT)
})
@Entity
class Book {
    ...

    @OneToOne
    Person person;

    @ManyToMany
    Set<Author> authors;

    ...
}
```

We may define as many different fetch profiles as we like.

Table 67. Annotations for defining fetch profiles

Annotation	Purpose
@FetchProfile	Declares a named fetch profile, optionally including a list of @FetchOverrides
@FetchProfile.FetchOverride	Declares a fetch strategy override as part of the @FetchProfile declaration

A fetch profile must be explicitly enabled for a given session:

```
session.enableFetchProfile("EagerBook");
Book eagerBook = session.find(Book.class, bookId);
```

So why or when might we prefer named fetch profiles to entity graphs? Well, it's really hard to say. It's nice that this feature *exists*, and if you love it, that's great. But Hibernate offers alternatives that we think are more compelling most of the time.

The one and only advantage unique to fetch profiles is that they let us very selectively request subselect fetching. We can't do that with entity graphs, and we can't do it with HQL. Unresolved directive in Hibernate_Introduction.adoc - include::Credits.adoc[]